

ugp-lean: A Machine-Checked Formalization of the Universal Generative Principle

Nova Spivack

<https://github.com/novaspivack/ugp-lean>

Preprint.

Archival version (Zenodo): <https://doi.org/10.5281/zenodo.19433538>

NEMS program hub (Zenodo): <https://doi.org/10.5281/zenodo.19487299>

Abstract

We present **ugp-lean**, a LEAN 4.4 formalization of the Universal Generative Principle (UGP) against Mathlib, comprising 148 canonical modules (with an extended development branch covering spacetime, QFT, and hadronic physics) with a strict zero-sorry, zero-custom-axiom policy on the core proof path (20 global custom named axioms after the 2026-05-24 audit, down from 24). The UGP is a deterministic arithmetic framework on integer ridges $R_n = 2^n - 16$ that produces, from first principles and without free parameters, a unique canonical seed (1, 73, 823) and a rigid Generative Triple Evolution (GTE) orbit.

The library establishes four interlocking results. *Algebraic uniqueness*: the ridge sieve, prime-lock criterion, and the Residual Seed Uniqueness theorem prove that (1, 73, 823) is the unique lexicographic minimum of its residual seed class, and that the GTE map T derives the canonical orbit $(1, 73, 823) \rightarrow (9, 42, 1023) \rightarrow (5, 275, 65535)$ unconditionally. *Computational universality*: the UWCA built natively from the UGP survivor space simulates Rule 110, and the CUP-4 theorem proves that the Standard Model generation orbit *algebraically forces* Rule 110 as the unique vacuum-transparent CA rule consistent with SM winding structure. *Physical formalization*: 515 certified theorems in the **GUTStructure** module derive Standard Model observables from $N_{\text{gen}} = 3$ and $N_{\text{fam}} = 5$ alone, including $\sin^2 \theta_W = 3/13$, $\lambda_{\text{CKM}} = 9/40$, $D = 4$, $Q = w^*/3$ for all SM fermions, and $\beta_0 = 23/3$ for QCD. GTE is certified as the terminal F_{PSC} coalgebra in the PSCSys category (zero sorry, zero custom axioms). *Physical extension (development branch)*: spacetime causal structure with machine-certified spectral dimension $d_s = 4$ in the thermodynamic limit; the Algebraic Lifting Theorem connecting beable-level results to physical observables; a QFT mass gap lower bound in the colour-singlet sector; the Algebraic Descent Theorem establishing resolution-independence of all F_{21} -algebraic properties; two independent Lean-certified derivations of $\alpha_{\text{EM}} \approx 1/137$; strong CP resolution ($\theta_{\text{QCD}} = 0$) and asymptotic freedom ($b_0 = 7$) from F_{21} group theory; and capstone theorems for the unified physical exclusion criterion and the uniqueness of three generations.

The library additionally formalizes the Elegant Kernel closure ($k_{\text{gen}} = \varphi \cos(\pi/10)$), the Koide cyclotomic-12 closed form with S_3 -equivariant Newton flow, the full BraidAtlas charge structure, neutrino mass-ratio arithmetic, and six formally stated open conjectures. All proved statements are scoped to the ridged arithmetic framework.

Archival version (Zenodo): <https://doi.org/10.5281/zenodo.19433538>

Contents

1	Premises, Disagreement Coordinates, Trust Boundary	5
2	Introduction	5
3	The UGP Framework	6

4	Architecture and Design Decisions	7
4.1	Module structure	7
4.2	Zero-sorry, zero-custom-axiom policy	9
4.3	Non-circularity contract	10
4.4	Dependency management	10
4.5	The orbit-from- T upgrade	10
5	Core Classification Results	10
6	GTE Orbit and Update Map	11
7	GTE Extended Structure	12
7.1	Mersenne ladder and Fibonacci uniqueness (<code>MersenneLadder</code>)	12
7.2	Fiber bundle over GTE triples (<code>FiberBundle</code>)	13
7.3	Orbital stability and linear response (<code>LinearResponse</code>)	13
7.4	Ridge-level scale connection (<code>ScaleConnection</code>)	13
7.5	GTB generation prime certificate (<code>GTE.GTBGenerationPrimes</code>)	13
7.6	GTE nucleon seed parity and nuclear pairing (<code>NuclearPairing</code>)	14
7.7	$N_c = 3$ from substrate arithmetic (<code>NcColorArithmetic</code>)	15
8	Braid Atlas Charge Structure	15
8.1	Composite braid c-rule	16
8.2	SM Winding Numbers from N_c (<code>BraidAtlas.ChargeDerivation</code>)	17
8.3	Coxeter–Conductor Theorem	18
8.4	Coxeter–Conductor Biconditional (<code>CoxeterBiconditional</code>)	18
8.5	Cyclotomic field embedding (<code>CyclotomicContainment</code>)	19
8.6	Electroweak boson c-values (<code>BraidAtlas.EWBosons</code>)	20
8.7	Chirality squaring (<code>BraidAtlas.ChiralitySquaring</code>)	21
8.8	Mirror winding number (<code>BraidAtlas.MirrorWindingNumber</code>)	21
8.9	EW boson–RHN cross-sector connection (<code>BraidAtlas.EWBosonRHNConnection</code>)	22
8.10	Dark sector certification hub (<code>BraidAtlas.DarkBraidAtlas</code>)	22
8.11	RHN b_3 gap theorem (<code>BraidAtlas.RHNGapTheorem</code>)	23
8.12	Dark quark charge (<code>BraidAtlas.DarkQuarkCharge</code>)	23
8.13	$\mathbb{Z}_7$ dark charge arithmetic (<code>BraidAtlas.DarkBraidAtlas</code> , §9)	24
8.14	Dark SU(3) gauge coupling (<code>BraidAtlas.DarkBraidAtlas</code> , §10)	24
8.15	Dark SU(3) gauge coupling — master formula derivation (<code>BraidAtlas.DarkGaugeCoupling</code>)	24
9	RCC Infinite Families (<code>PSC.RCCInfiniteFamilies</code>)	25
10	PSC Three-Route Forcing Capstone (<code>PSC.ThreeRouteForcing</code>)	25
11	Neutrino FN Texture Selection (<code>NeutrinoFroggattNielsen</code>)	26
12	Galois-Protection Non-Renormalization (<code>Phase4.GaloisProtection</code>)	27
13	Two-Loop Color Coefficient (<code>Phase4.TwoLoopCoefficient</code>)	27
14	Null Discipline Formalization	28
14.1	Algebraic saturation barrier (<code>SaturationBarrier</code>)	28
14.2	Theorem eligibility criterion (<code>TheoremEligibility</code>)	28
15	Information Profit Threshold	29
15.1	H9: IPT Self-Consistency Fixed Point	29
16	Gauge Coupling Extensions	30
16.1	Asymptotic Sparsity and Positive Root Theorem (P25, 2026)	31
16.2	Galois Layer Stability (<code>GaloisStructure.CyclotomicLayers</code>)	32
17	General-Level Theorems	32
18	UGP Primes	33
19	Quarter-Lock and Elegant Kernel	34
19.1	UCL Unconditional Closure	34

20 Mass Relations	35
20.1 Koide cyclotomic-12 closed form	35
20.2 S_3 -equivariant Newton flow	35
20.3 Binary cascade and physical spectrum	36
20.4 Colour-rank structural chain (KoideAngle)	36
20.5 Seesaw index as GUT representation defect (SeesawIndex)	37
20.6 Mass ratio Z-independence (ScaleTransport)	37
20.7 CKM θ_{23} Structural Ratio (MassRelations.CKMTheta23)	37
20.8 Koide S_3 Discrete Arithmetic-Mean Identity (MassRelations.KoideS3DiscreteIdentities)	38
20.9 CDM Cabibbo derivation (MassRelations.CKMMixing)	38
20.10 Neutrino mass-squared ratio (MassRelations.NeutrinoMassRatio)	38
20.11 SRRG no-go theorem and VEV derivation	39
21 Universality and Self-Reference	39
21.1 Turing universality	39
21.1.1 The UWCA simulation theorem (machine-checked)	40
21.1.2 Dynamics companions (meta-law ML-9)	41
21.2 Self-reference theorems	42
21.3 CUP-4: SM Generation Parity Orbit (CUP4TotalParity)	42
21.4 CUP-11c: Universal Mod-7 CA (CUP11ModSeven)	42
21.5 CUP-3D Structural Theorems (CUP3DUniqueness)	43
21.6 PSC Unification and GoE Chain (CUP3DPSCUnification)	43
21.7 Physical Incompleteness (CUP3DPhysicalIncompleteness)	44
21.8 GTE-NEMS Framework Instance (GTEFrameworkInstance)	44
21.9 GTE Optimality (GTEOptimalityInstance)	45
21.10 GTE as Final F_{PSC} Coalgebra (GTEFinalCoalgebra)	45
21.11 Two-Layer Confluence (TwoLayerConfluence)	46
21.12 GTE Tile Compilation (GTECompilation)	46
21.13 GTE Bisimulation Uniqueness (GTEUniqueness)	47
21.14 GoE Hierarchy (GoEHierarchy)	47
21.15 GTE Infinite-Tape Encoding (GTEInfTapeEncoding)	47
21.16 GTE Computability (GTEComputability)	47
21.17 Hypothesis B: Single Universal Substrate (HypothesisB)	48
21.18 Hypothesis B+C Chain (HypothesisBCChain)	48
21.19 PSC Implies Computational Universality (PSCUniversality)	49
21.20 Cook Rule 110 Reference (CookRule110Ref)	49
21.21 RCC Analytical Completion (PSC.RCCComplete)	49
22 GUT Structure and Physical Formalization (GUTStructure)	50
22.1 GTE constants and GUT Weinberg angle (§§1–8)	50
22.2 f_{MDL} structure and chirality (§§9–11)	50
22.3 Weinberg angle closure (§12)	50
22.4 Z_5 running shift and CKM (§§13–15, 18–21, 24–26)	50
22.5 GTE master formula (§27)	51
22.6 SM N-value sum, Z_2 universality, Mersenne structure (§§16–17, 20–21)	51
22.7 Orbit dynamics, vacuum structure, and GoE (§§28, 32, 36, 39–41, 57–61)	51
22.8 Z_7 charge, color, and gauge structure (§§33, 47–56)	51
22.9 Hypercharge, Gorard matter, mass ordering (§§73–76)	52
22.10 Baryogenesis, winding classes, QCD (§§70, 79–81)	52
22.11 Additional GUTStructure results	52
23 New Universality Modules	53
23.1 GoE Stability Hierarchy (GoEStabilityHierarchy)	53
23.2 Orbit Perturbation Catalog (OrbitPerturbationCatalog)	53
23.3 Z_7 Charge Conjugation (Z7ChargeConjugation)	53
23.4 Z_5 Transitivity Uniqueness (Z5TransitivityUniqueness)	54
23.5 Dimensional Slice Uniqueness (DimensionalSliceUniqueness)	54
23.6 GTP Neutral Discrimination (GTPNeutralDiscrimination)	54
23.7 SM Orbit Causal Isolation (SMOrbitCausalIsolation)	54
23.8 EW Boson Structure (EWBosonStructure)	54
23.9 EW Chiral Bridge (EWChiralBridge)	55
23.10 Casimir Massless Ether (CasimirMasslessEther)	55
23.11 Lawvere Zone (LawvereZone)	55

24 Spacetime and Causal Structure Formalization	55
24.1 Spectral dimension modules	55
24.2 Causal invariance and Minkowski structure (CausalInvariance.lean)	56
24.3 Algebraic Lifting Theorem (LiftingTheorem.lean)	57
24.4 Geodesic Theorem (GeodesicTheorem.lean)	57
25 Coupling Hierarchy and Fine Structure Constant (SylowIndexCouplingHierarchy.lean)	58
25.1 Fine structure constant — §§5d–5h	58
25.2 Substrate uniqueness — §§4c, 5b, 5c	58
25.3 Session 6: F_{21} Physics, Hadron Spectroscopy, and Mass Gap — §§5i–9	58
26 Two-Sector Substrate (Substrate/LExtended.lean)	59
27 QFT Mass Gap (QFT/GaugedMassGap.lean)	59
27.1 Conditional theorem	60
27.2 Unconditional theorems (Rank 72d)	60
28 Audit-Session Capstones and Axiom Inventory Update	60
28.1 Axiom budget update	60
28.2 Bridge capstones	61
28.3 Anomaly cancellation and renormalizability (Spacetime/AnomalyRenormalizability.lean)	61
28.4 Algebraic Descent Theorem (Universality/AlgebraicDescentTheorem.lean)	61
28.5 Mass gap and hierarchy	62
28.6 Additive QGR certificates (Spacetime/QuantumGravity.lean)	62
28.7 Cross-repository certification	62
29 The Mirror-Dual Conjecture	62
29.1 Resonant factory formalization	64
30 Mirror Shift Algebra	64
30.1 The shared-residue theorem	65
30.2 The mirror shift formula	65
30.3 Concrete instances	65
30.4 Structural interpretation	66
31 Real Analysis and the DSI Interface	66
32 Structural Theorems	67
32.1 Uniqueness certificates	68
33 Open Conjectures	69
34 Related Work	70
35 Conclusion	70

1 Premises, disagreement coordinates, and trust boundary

Where disagreement can (and cannot) live

Formal spine. Classification and forcing lemmas refer to definitions in `Ugcp/` (or as cited in the module tree); numerical witnesses are `#eval`-friendly where stated.

Premise coordinates (for external readers).

- **Ridge arithmetic vs. ambient set theory.** Claims are about the constructed UGP objects (ridges, sieve, T -orbit), not about every conceivable number-theoretic universe. *Logical effect:* Generalizing beyond the formalized ridge family requires new definitions, not silent extension.
- **Toolchain / Mathlib pin.** The paper reports builds against a stated Mathlib revision; upgrading may require port work. *Logical effect:* Port failures do not erase prior checked revisions.

Keywords: interactive theorem proving, Lean 4, formalization, number theory, ridge sieve, Mersenne numbers, prime-producing formulas, cellular automata, mass relations, braid topology, cyclotomic fields

2 Introduction

The Universal Generative Principle (UGP) is a deterministic number-theoretic framework that generates structured integer triples from a single arithmetic seed. The framework centers on *ridges* $R_n = 2^n - 16$, a divisor-based sieve that selects *survivor* pairs, and a parity-dependent update map T called the Generative Triple Evolution (GTE) that evolves triples through an orbit. At the distinguished level $n = 10$, the sieve produces exactly two survivor pairs, and the Minimum Description Length (MDL) principle selects a unique seed—the *canonical seed* $(1, 73, 823)$ —whose orbit under T is $(1, 73, 823) \rightarrow (9, 42, 1023) \rightarrow (5, 275, 65535)$.

The UGP framework makes strong arithmetic claims: uniqueness of the residual seed, primality of specific values, determinism of the orbit, and Turing universality of the substrate. These claims benefit from machine-checking for three reasons:

1. The number-theoretic arguments involve explicit computation (sieving divisors, verifying primality) that is error-prone by hand.
2. The framework introduces novel definitions (ridge sieve, Quarter-Lock, mirror-dual pairs) whose internal consistency is non-trivial.
3. A machine-checked artifact serves as a citable, independently verifiable record.

Contributions. The `ugp-lean` artifact formalizes the UGP framework in LEAN 4 [4] with Mathlib [7]. All definitions are self-contained: the formalization defines every concept it uses (ridge, sieve, update map, orbit, UGP primes) from first principles, depending only on Mathlib’s standard number theory. Beyond the core framework, the formalization contributes several *new* results:

- **Orbit from T** (§6): The canonical orbit values are *derived* from the update map, not postulated.
- **General ridge remainder lock** (§17): For *all* $n \geq 5$ and all valid divisors b_2 , $(2^n - 1) \bmod b_2 = 15$.
- **Even-step c -invariance** (§6): Under the strict per-step rule, the even step immediately after a ridge hit leaves c unchanged.

- **ugp primes** (§18): A new class of primes with a formal predicate, an existence proof, and ten machine-checked instances across three ridge levels.
- **Mirror-dual conjecture** (§29): A precise formal statement with five machine-verified witness pairs, a proof that the divisor count $\tau(2^n - 16)$ is unbounded, and an exact formula $\tau(2^n - 16) = 5 \cdot \tau(2^{n-4} - 1)$ via coprimality and multiplicativity of τ .
- **Resonant factory** (§29.1): Machine-checked algebraic infrastructure for a twin-prime program: branch linearization, gap-2 factory quadratics, and a Hasse check (no local obstruction at primes ≤ 43).
- **Mirror shift algebra** (§30): A new universal algebraic law: for every mirror-dual pair (b_2, q_2) at any ridge level, both c_1 -values satisfy $c_1 \equiv 20 \pmod{b_1}$, and the shift $c_1(b_2, q_2) - c_1(q_2, b_2) = b_1 \cdot (q_2 - b_2)$. At $n = 10$ this gives $2137 - 823 = 73 \cdot 18 = 1314 = 2 \cdot 3^2 \cdot 73$. Machine-checked at $n = 10, 13, 16$.
- **Real analysis** (§31): The c_1 function on the divisor hyperbola is differentiable with derivative $13R/q^2 + 2q - 6 \geq 22$ on the valid shell—a machine-checked uniform magnification bound for sieve applications.
- **UCL unconditional closure** (§19): Nine fully unconditional theorems prove $k_{\text{gen}} = \varphi \cos(\pi/10)$ and the complete Elegant Kernel coefficient set from the pentagon–Fibonacci structure and D_5 renormalization, with zero axioms and zero `sorry` beyond Lean’s foundational set.
- **Mass relations and physical spectrum** (§20): The `MassRelations` layer (22 modules) covers (i) the Koide closed form and S_3 -balance arithmetic identities (`KoideClosedForm`, `KoideAngle`, `KoideNewtonFlow`, `KoideS3DiscreteIdentities`, `ClaimCBridge`); (ii) the $SU(5)/SO(10)$ Yukawa power law and the VV mass mechanism, with all three VV exponents derived from $N_c = 3$ (`VVMechanism`, `VVAllCoefficientsFromNc`, `DownRational`); (iii) the CKM θ_{23} structural ratio $\tau(R_{10})/D_1 = 15/8$ uniquely forced at $n = 10$ (`CKMTheta23`); (iv) the binary cascade mass-generation recurrence and the physical charged-lepton, up-type, and down-type mass spectra (`BinaryCascade`, `PhysicalMasses`, `HeavyFermionTower`, `LeptonMassPrediction`); (v) Froggatt–Nielsen completion in both the charged-fermion and neutrino sectors, with the $b^{29/9}$ neutrino seesaw FN texture selected by MDL (`FroggattNielsen`, `NeutrinoFroggattNielsen`); (vi) Clebsch–Gordan, Cartan, and $SU(3)$ flavour geometry underpinning the flavon VEV structure (`ClebschGordan`, `CartanFlavonPotential`, `SU3FlavorCartan`); and (vii) the scale-transport, seesaw, and cyclotomic-orbifold machinery linking ridge arithmetic to the physical mass spectrum (`ScaleTransport`, `SeesawIndex`, `Z2OrbifoldDepth`, `UpLeptonCyclotomic`).

Structure. §3 defines the UGP framework. §4 describes the module architecture and design decisions. §5 presents the core classification results (RSUC). §6 covers the GTE orbit and update map. §17 collects the general-level theorems. §18 introduces UGP primes. §19 treats the Quarter-Lock, Elegant Kernel, and UCL unconditional closure. §20 presents the mass relations layer. §21 addresses universality and self-reference. §24 covers the spacetime and causal-structure formalization (development branch). §25 presents the coupling hierarchy and α_{EM} derivation modules. §26 describes the two-sector substrate encoding. §27 presents the QFT mass-gap theorems. §28 documents the axiom-inventory update and bridge capstones. §29 develops the mirror-dual conjecture. §30 presents the mirror shift algebra. §31 presents the real-analysis layer. §32 presents the structural theorems and uniqueness certificates. §33 states the six remaining open conjectures. §34 surveys related work. §35 concludes.

3 The UGP Framework

We define the mathematical content that is formalized. All definitions below have direct `LEAN 4` counterparts in the `UgpLean.Core` modules.

Definition 3.1 (Ridge). The *ridge* at level n is $R_n := 2^n - 16$. At $n = 10$, $R_{10} = 1008$.

Definition 3.2 (Strict-ridge divisor pair). A *strict-ridge divisor pair* at level n is a pair $(b_2, q_2) \in \mathbb{N}^2$ with $b_2 \cdot q_2 = R_n$ and $b_2 \geq 16$, $q_2 \geq 16$.

From a divisor pair (b_2, q_2) the UGP derives:

$$b_1 = b_2 + q_2 + 7, \tag{1}$$

$$q_1 = q_2 - 13, \tag{2}$$

$$c_1 = b_1 \cdot q_1 + 20. \tag{3}$$

The constants $s = 7$, $g = 13$, $t = 20$ are the UGP-1 parameters.

Definition 3.3 (Prime-lock criterion). A divisor pair (b_2, q_2) is *prime-locked* if c_1 as defined in (3) is prime.

At $n = 10$, $R_{10} = 1008$ has eight divisor pairs with both components ≥ 16 . Of these, exactly two are prime-locked: $(24, 42)$ yielding $c_1 = 823$ (prime), and $(42, 24)$ yielding $c_1 = 2137$ (prime). The six non-survivors have composite c_1 values, each explicitly excluded in the formalization (`ExclusionFilters`).

Definition 3.4 (GTE update map T). Given a triple $G_t = (a_t, b_t, c_t)$ at level n , define $q_t = c_t \div b_t$ and $m_t = c_t \bmod b_t$. The update map T is parity-dependent:

Odd step (t odd):

$$a_{t+1} = m_t - (n + 2 - t), \quad b_{t+1} = b_t - (m_t + q_t), \quad c_{t+1} = 2^n - 1. \tag{4}$$

Even step (t even):

$$a_{t+1} = m_t - (n + 2 - t), \quad b_{t+1} = b_t + \text{Fib}_{|q_t - q_{t-1}|}, \quad c_{t+1} = b_t \cdot q_t + 15. \tag{5}$$

Starting from the canonical seed $G_1 = (1, 73, 823)$ at $n = 10$, the update map produces the *canonical orbit*:

$$(1, 73, 823) \xrightarrow{T} (9, 42, 1023) \xrightarrow{T} (5, 275, 65535).$$


The key derived quantities are $q_1 = 11$, $m_1 = 20$, $q_2 = 24$, $m_2 = 15$, and the quotient gap $q_2 - q_1 = 13$, which forces the Fibonacci lift $\text{Fib}_{13} = 233$.

Remark 3.5 (Illustrative vs. strict c_3). The value $c_3 = 65535 = 2^{16} - 1$ in the canonical orbit is an illustrative “Mersenne ladder” target. Under the strict per-step rule (5), $c_3 = b_2 \cdot q_2 + 15 = 42 \cdot 24 + 15 = 1023 = c_2$. The c -value is *unchanged* on the even step immediately after a ridge hit. This is a theorem, not a coincidence (§6).

4 Architecture and Design Decisions

4.1 Module structure

The artifact is organized into 148 LEAN 4 modules across thirteen layers (six modules migrated to `ugp-physics-lean`; see below):



The remaining modules (Phase4, Papers, Instance, Conjectures, TE22, GaloisStructure, CyclotomicCompleteness) provide stubs, citable entry points, bridge instances, and the precision-derivation theorems: Phase4.GaloisProtection (one-loop non-renormalization), Phase4.TwoLoopCoefficient (two-loop color coef-

ficient = 8/9), `GaloisStructure.CyclotomicLayers` (Galois-stability of UGP layers in $\mathbb{Q}(\zeta_{120})$), `GaloisStructure.MinimalCyclotomic` ($\mathbb{Q}(\zeta_{120})$ minimality), and `CyclotomicCompleteness.CoxeterBiconditional` (Coxeter-conductor biconditional; see §8.4), and `CyclotomicCompleteness.CyclotomicContainment` (field-theoretic embedding $\mathbb{Q}(\zeta_{2h}) \hookrightarrow \mathbb{Q}(\zeta_{120})$ for $h \mid 60$; see §8.5).

The `TE22` module provides a machine-checked framework for the `TE2.2` PSC universe scan result. Two theorems are proved with zero `sorry`: `ugp_coupling_predictions_are_independent` (the three UGP-derived coupling ratio predictions C_{15} , C_{16} , C'_4 are algebraically derived from ugp-lean machine-checked rationals, not from SM coupling data) and `ugp_g1g2_prediction_close_to_SM` (the UGP g_1^2/g_2^2 prediction is within 2% of the SM value at M_Z). Two further theorems target the gauge-group labelling: `SM_gauge_uniquely_selected` proves by `decide` that exactly one of the 60 (GaugeGroup, Dimension) pairs satisfies the `isSMGauge` predicate, and `isSMGauge_iff` gives the full logical characterisation of that label. A fourth theorem (`SM_is_D_minimizer_extended`) is now proved with zero `sorry` as an alias `:= isSMGauge_iff`, providing the decidable fragment of the SM-uniqueness claim. The full machine-checked proof that the SM is the unique D -minimizer over the 34,560-universe scan remains an open programme (framework and decidable fragment in place; `native_decide` certification pending a computable `Fintype` instance on the full `UniverseParams` product type and computable constraint functions; computational certificate in `ugp-physics/MFRR/.../extended_scan_results.json`, SHA-256: 407078d7...).

4.2 Zero-sorry, zero-custom-axiom policy

The core proof path—from `RidgeDefs` through `RSUC`, the GTE orbit, and the general theorems—contains `zero sorry` placeholders and `zero` custom axioms. Every theorem in that path is either:

- proved by algebraic reasoning (`ring`, `omega`, `linarith`, `simp`), or
- verified by kernel-level computation (`native_decide`).

The only axioms in scope are `LEAN 4`’s foundational axioms (`propext`, `Quot.sound`, `Classical.choice`) and `Mathlib`’s standard imports.

The new `TE22` module is now zero-sorry: every theorem in `UgpLean.TE22.ScanCertificate` closes by algebraic reasoning or `decide`, including `SM_is_D_minimizer_extended` (proved as an alias `:= isSMGauge_iff`, the decidable fragment of the full SM-uniqueness claim). The full machine-checked proof that the SM is the unique D -minimizer over the extended 34,560-universe scan remains an open programme, with proof strategy and computational certificate recorded; the `native_decide` proof is pending a computable `Fintype` instance on the full `UniverseParams` product type and computable constraint functions.

Documented axioms (post-audit: 20 global). Two theorems in `GTE.AnalyticArchitecture` are declared as `axiom` rather than `sorry`: `dickman_equidistribution_in_APs` and `crt_equidistribution_within_regime`, both citing Tenenbaum [17] III.6. These are analytic number theory statements whose machine-mechanization would require upstream `Mathlib` development (equidistribution of divisor sums in progressions); declaring them as axioms rather than sorries makes the dependency explicit and auditable. Crucially, *neither theorem appears in the axiom closure of any physics theorem in the library*: running `#print axioms` on any charged-lepton or coupling theorem returns exactly the standard Lean/`Mathlib` signature `{propext, Classical.choice, Quot.sound}`.

A systematic audit (2026-05-24) discharged four previously declared axioms in `VEVProof/GoldstoneEntropyCorrection.lean` as genuine theorems, reducing the global custom named-axiom count from 24 to **20**; see §28 for details.

One additional disclosed axiom (`BraidAtlas.ChargeTheorem`, 2026-05-08):

mirror_winding_number_zero: $W_{g,\text{mirror}} = 0$ for the GTE-P7 mirror-branch dark matter candidate. This is a *braid-topology postulate* imported from the companion paper P17 [11] §subsec:mirror_dm: the canonical Braid Atlas v2.0 master table assigns the mirror-branch winding number to be 0 on topological grounds (mirror-orbit Wilson-loop vanishing), and the full P17 derivation has not yet been formalised in Lean. The axiom is faithfully tracked in the P17 PROVENANCE.md ledger. It supports the formal derivation **gte_p7_electric_charge_zero** ($Q_{\text{GTE-P7}} = 0$, see §8) and does *not* appear in the axiom closure of any SM-sector theorem (lepton, quark, gauge coupling, mass relation). All other theorems in the library are **sorry-free**.

4.3 Non-circularity contract

A critical design constraint: the RSUC proof must not be circular. We enforce this by a strict module boundary:

Core/ may not import Compute/.

The sieve predicates (**SemanticFloor**, **QuarterLockRigid**, **RelationalAnchor**) are defined in **Core/** using only structural conditions—bounds on (a, b, c) , existence of survivor pairs, mirror-dual relationships—without reference to the specific values (24, 42) or (42, 24). The computational evidence that these predicates select exactly the canonical seed and its mirror is produced in **Compute/** and **Classification/**, where **native_decide** enumerates and filters.

4.4 Dependency management

The artifact depends on:

- **Mathlib** v4.29.0-rc6 [7]: number theory (**Nat.divisors**, **Nat.Prime**, **Nat.fib**), finset operations, and the Mersenne gcd identity (**Nat.pow_sub_one_gcd_pow_sub_one**).
- **nems-lean** [10]: Lawvere fixed-point theorem and Kleene recursion theorem.
- **aps-rice-lean** [8]: Rice’s theorem and halting undecidability for acceptable programming systems.

4.5 The orbit-from- T upgrade

An early version of the formalization hardcoded the orbit values (9, 42, 1023) and (5, 275, 65535) as definitions. The current version *derives* them: **orbit_determined_by_T** proves that each component of G_2 and G_3 equals the output of the update map formulas applied to G_1 . The hardcoded values in **Evolution.lean** are now corollaries, not postulates.

5 Core Classification Results

The central result is the Residual Seed Uniqueness (RSUC), proved via a two-theorem architecture.

Definition 5.1 (Unified Admissibility). A triple $t = (a, b, c)$ is *unified-admissible at level n* , written **UnifiedAdmissibleAt n t** , if it satisfies three independent predicates:

1. **SemanticFloor**: $a \in \{1, 5, 9\}$, $b \geq 40$, $c \geq 800$.
2. **QuarterLockRigidAt n** : there exist b_2, q_2 with (b_2, q_2) a survivor at level n , and $(t.b, t.c)$ derived from that pair.
3. **RelationalAnchorAt n** : there exist b_2, q_2 with (b_2, q_2) a survivor at level n , $t.b = b_1(b_2, q_2)$, and $t.c \in \{c_1(b_2, q_2), c_1(q_2, b_2)\}$.

Definition 5.2 (Candidates and Residual). `CandidatesAt  $n$` is the finite set of triples generated by all survivor pairs at level n crossed with $a \in \{1, 5, 9\}$. At $n = 10$, $|\text{Candidates}| = 6$ (three a -values times two c -values, with $b = 73$ forced). The *residual* is `Residual := Candidates.filter(decUnifiedAdmissible)`.

Theorem 5.3 (Theorem A — Structural Reduction). *For all n and all triples t :*

$$\text{UnifiedAdmissibleAt } n \ t \implies t \in \text{CandidatesAt } n.$$

Lean: `theoremA_general`

Proof sketch. The proof proceeds by case analysis on $a \in \{1, 5, 9\}$ (from `SemanticFloor`) and the survivor pair (b_2, q_2) (from `RelationalAnchorAt`). Each case constructs the appropriate member of `CandidatesAt  $n$` by `Triple.ext`. At $n = 10$, the proof reduces to six `native_decide` calls. $\square$

Theorem 5.4 (Theorem B — Finite Classification). `Residual = Candidates`. *That is, every candidate triple passes the unified sieve.*

Lean: `theoremB`

Proof sketch. The forward direction is immediate (filtering preserves membership). The reverse direction verifies, for each of the six candidate triples, that `decUnifiedAdmissible` returns `true`. This is dispatched by `native_decide` on each triple. $\square$

Theorem 5.5 (RSUC — Residual Seed Uniqueness). *Under the unified sieve at $n = 10$:*

1. *The residual set equals the six candidate triples (three a -values $\times$ two mirror c -values).*
2. *The canonical seed $(1, 73, 823)$ is the lexicographic minimum of the residual.*

Lean: `rsuc_theorem`

Proof sketch. Combines Theorems A and B. Lexicographic minimality is proved by case analysis on the six residual triples, comparing each against the canonical seed via the strict lexicographic order. $\square$

Exclusion filters. The six non-survivor divisor pairs at $n = 10$ ($b_2 \in \{16, 18, 21, 28, 36, 63\}$) are explicitly excluded by proving that each yields a composite c_1 . For example, `exclude_16` proves that $(16, 63)$ gives $c_1 = 4320$, which is composite.

Monotonic strengthening. `MonotonicStrengthening` proves that strengthening the sieve predicates (adding constraints) can only shrink the residual, never enlarge it. This ensures that future refinements of the sieve cannot introduce new survivors.

6 GTE Orbit and Update Map

The GTE update map T (theorem 3.4) is formalized for the first time in any LEAN 4 artifact. The key results are:

Theorem 6.1 (Orbit determined by T). *The canonical orbit at $n = 10$ is fully determined by the update map:*

$$\begin{aligned} G_2 &= T(G_1) = (9, 42, 1023), \\ G_3 &= T(G_2) = (5, 275, 65535). \end{aligned}$$

Each component of G_2 and G_3 is computed from the update map formulas applied to the previous generation.

Lean: `orbit_determined_by_T`

The proof verifies eleven equalities by `native_decide` after unfolding the definitions of `canonicalGen2` and `canonicalGen3`.

Theorem 6.2 (Even-step c -invariance). *For all $n \geq 5$ and all b_2 with $b_2 \mid R_n$ and $b_2 \geq 16$:*

$$\text{evenStepC}(b_2, (2^n - 1)/b_2) = 2^n - 1.$$

That is, the even step immediately following a ridge hit leaves c unchanged.

Lean: `even_step_c_invariance`

Proof. By the ridge remainder lock (theorem 17.1), $m_2 = (2^n - 1) \bmod b_2 = 15$. Therefore $b_2 \cdot q_2 = (2^n - 1) - 15 = 2^n - 16 = R_n$. The even-step formula gives $c_{t+1} = b_2 \cdot q_2 + 15 = R_n + 15 = 2^n - 1 = c_t$. $\square$

Corollary 6.3. *At $n = 10$, $b_2 = 42$: the strict even-step output is $c_3 = 1023 = c_2$, not 65535.*

Lean: `c3_strict_eq_c2_at_n10, paper_c3_is_illustrative_not_strict`

Derived quantities. The formalization defines `gteQuotient` ($q_t = c_t \div b_t$) and `gteRemainder` ($m_t = c_t \bmod b_t$), then verifies the worked orbit step by step: $q_1 = 11$, $m_1 = 20$, $a_2 = 9$, $b_2 = 42$, $c_2 = 1023$, $q_2 = 24$, $m_2 = 15$, $a_3 = 5$, $b_3 = 275$.

Cross-generation entanglement. The Mersenne gcd identity $\gcd(2^a - 1, 2^b - 1) = 2^{\gcd(a,b)} - 1$ (from Mathlib’s `Nat.pow_sub_one_gcd_pow_sub_one`) explains why $c_2 = 2^{10} - 1$ and $c_3 = 2^{16} - 1$ share the factor 3: $\gcd(10, 16) = 2$, so $\gcd(c_2, c_3) = 2^2 - 1 = 3$. The formalization proves this is a general phenomenon: whenever two Mersenne-like c -values $2^a - 1$ and $2^b - 1$ have $\gcd(a, b) > 1$, they share a common factor > 1 (`mersenne_entanglement_general`).

Gen 1 isolation. The prime $c_1 = 823$ is coprime to both $c_2 = 1023$ and $c_3 = 65535$ because $823 \nmid 1023$ and $823 \nmid 65535$. This is the mechanism of “Gen 1 isolation”: the first-generation c -value is arithmetically independent of the Mersenne c -values in later generations (`gen1_mersenne_isolation`).

7 GTE Extended Structure

7.1 Mersenne ladder and Fibonacci uniqueness (`MersenneLadder`)

The canonical orbit’s third-generation c -value $c_3 = 65535 = 2^{16} - 1$ is not an arbitrary hardcoded choice. The `GTE.MersenneLadder` module derives the value from a structural extension rule grounded in the QCD colour rank $N_c = 3$ and the Fibonacci orbit structure.

Theorem 7.1 (T-phys Mersenne formula). *The physics extension rule $\text{evenStepC_phys}(k, N_c) = 2^{k+2N_c} - 1$ at $k = 10$, $N_c = 3$ gives $c_3 = 2^{16} - 1 = 65535$. The extension exponent $k' = 16$ satisfies $k' = k + 2N_c = 10 + 6 = 16$.*

Lean: `c3_phys_formula, evenStepC_phys_at_n10_Nc3` (in `GTE.MersenneLadder`)

Theorem 7.2 (Fibonacci uniqueness of the exponent). *The exponent $k' = 16 = 2F_6$ is the unique minimal double Fibonacci number strictly greater than $n = 10 = 2F_5$. In particular, $k = 12$ fails (`k_12_not_double_fib`), while $k' = 16$ succeeds, so the Fibonacci criterion uniquely selects the physics extension exponent.*

The complete Fibonacci orbit structure is: $F_4 = 3 = N_c$, $2F_5 = 10 = n$, $F_6 = 8 = \dim(\text{SU}(N_c))$, $2F_6 = 16 = k'$, $F_7 = 13 = \text{quotient gap}$ (already proved as `fib13_is_233`).

Lean: `kprime_is_minimal_double_fib_above_n, no_double_fib_between_n_and_kprime` (in `GTE.MersenneLadder`)

Theorem 7.3 (Cross-generation entanglement). $\gcd(c_2, c_3) = \gcd(1023, 65535) = 3 > 1$: the two Mersenne values share a common factor, machine-checked via the general Mersenne GCD identity.

Lean: `c3_phys_entangled_with_c2` (in `GTE.MersenneLadder`)

7.2 Fiber bundle over GTE triples (**FiberBundle**)

The `GTE.FiberBundle` module defines the covering space $\pi : \tilde{E}_{\text{GTE}} \rightarrow T_{\text{GTE}}$, where the fiber over each triple carries the quantum-number data (chirality, SM fermion type).

Theorem 7.4 (Fiber uniquely determined over canonical orbit). *The canonical lepton orbit a -values $\{a_{G_1}, a_{G_2}, a_{G_3}\} = \{1, 9, 5\}$ are all distinct, uniquely identifying each generation. Over the canonical orbit the fiber (chirality, fermion type) is therefore uniquely determined by the base triple: the bundle trivializes.*

Lean: `lepton_orbit_a_values_distinct`, `lepton_orbit_a_pattern`, `canonical_lepton_type_all_generations` (in `GTE.FiberBundle`)

7.3 Orbital stability and linear response (**LinearResponse**)

The `GTE.LinearResponse` module connects the canonical orbit's concrete arithmetic to the Fibonacci renormalization spectrum.

Theorem 7.5 (UGP residue is conserved). *The zero-mode residue $c \bmod b = 15$ is preserved at every ridge hit on the canonical orbit.*

Lean: `ugp_residue_is_conserved`, `canonical_residue_is_15` (in `GTE.LinearResponse`)

Theorem 7.6 (Orbital stability). *The Fibonacci contraction mode satisfies $\psi^2 < 1$ (from `PentagonConstraint`), establishing that transverse perturbations to the canonical orbit decay: the orbit is a stable fixed point of the Fibonacci renormalization flow.*

Lean: `fib_transverse_decay_sq`, `orbit_b_ratio_fibonacci_bounds` (in `GTE.LinearResponse`)

7.4 Ridge-level scale connection (**ScaleConnection**)

Theorem 7.7 (Universal 5-factor scale connection). *For all $n \geq 5$:*

$$\tau(2^n - 16) = 5 \cdot \tau(2^{n-4} - 1),$$

where τ denotes the number of positive divisors. Certified values: $\tau(R_{10}) = 30$, $\tau(R_{13}) = 20$, $\tau(R_{16}) = 120$.

Lean: `tau_scale_factor_is_5` (in `GTE.ScaleConnection`)

7.5 GTB generation prime certificate (**GTE.GTBGenerationPrimes**)

The `GTE.GTBGenerationPrimes` module certifies that every generation c_1 value produced by the GTB cascade is prime, and that the canonical three-generation ridge levels have uniform spacing equal to $N_c = 3$. The six values arise at ridges $n = 10, 13, 16$, from the standard and mirror branches respectively.

Six-generation prime certificate. The complete list of machine-certified primes is:

$$823, \quad 2137, \quad 9007, \quad 27817, \quad 46681, \quad 2489143.$$

Each is proved prime by `norm_num/decide` with zero `sorry`.

Theorem 7.8 (GTB generation prime certificate). *All six GTB generation c_1 values are prime, and the ridge spacing $16 - 13 = 13 - 10 = N_c = 3$. At ridges $n = 10, 13, 16$:*

$$\begin{aligned} n = 10: & \quad c_1 = 823, \quad c'_1 = 2137; \\ n = 13: & \quad c_1 = 9007, \quad c'_1 = 27817; \\ n = 16: & \quad c_1 = 46681, \quad c'_1 = 2489143. \end{aligned}$$

The uniform spacing $\Delta n = N_c$ connects the number of colours to the generation structure of the cascade.

Lean: `gtb_generation_prime_certificate`, `gtb_ridge_spacing_equals_Nc` (in `GTE.GTBGenerationPrimes`, zero sorry)

7.6 GTE nucleon seed parity and nuclear pairing (NuclearPairing)

The `GTE.NuclearPairing` module formalizes the arithmetic basis for the F_{10} proton-parity nuclear stability feature and the GTE prediction of the SEMF nuclear pairing constant, as introduced in companion paper P03 [?].

Nucleon seed constants. The GTE baryon sector assigns:

$$\begin{aligned} b_p &= 11459 \quad (\text{proton b-seed}), & b_n &= 11441 \quad (\text{neutron b-seed}), \\ a_{\text{seed}} &= 5 \quad (\text{common a-seed}), & g_{\text{nuclear}} &= 3 \quad (\text{nuclear generation}). \end{aligned}$$

Certified results (L001–L006, zero sorry). All six lemmas are proved by `LEAN4native_decide` or `LEAN4omega`, relying only on the standard axioms `LEAN4propext`, `LEAN4Classical.choice`, `LEAN4Quot.sound`.

Theorem 7.9 (Proton b-seed parity — L001). $b_p \bmod 2 = 1$, i.e. the GTE proton b-seed is odd.

Lean: `proton_b_seed_is_odd : gte_b_proton % 2 = 1` (in `GTE.NuclearPairing`, `LEAN4native_decide`)

Theorem 7.10 (Proton b-ladder parity — L003). *For any $Z \in \mathbb{N}$, the Z -fold proton b-ladder satisfies $(Z \cdot b_p) \bmod 2 = Z \bmod 2$. Equivalently: the parity of Z is encoded in the proton b-ladder.*

Lean: `proton_bseed_parity (Z : Nat): (Z * gte_b_proton) % 2 = Z % 2` (in `GTE.NuclearPairing`, `LEAN4omega` after `LEAN4unfold`)

Theorem 7.11 (Combined b-eff parity — L004). $(Z \cdot b_p + N \cdot b_n) \bmod 2 = (Z + N) \bmod 2 = A \bmod 2$. *The composite b-effective parameter has the same parity as the mass number.*

Lean: `beff_parity (Z N : Nat): (Z * gte_b_proton + N * gte_b_neutron) % 2 = (Z + N) % 2` (in `GTE.NuclearPairing`, `LEAN4omega`)

Theorem 7.12 (Proton-neutron b-seed difference — L005). $b_p - b_n = 18$.

Lean: `b_seed_difference : gte_b_proton - gte_b_neutron = 18` (in `GTE.NuclearPairing`, `LEAN4native_decide`)

The summary conjunction `LEAN4gte_nuclear_parity_rule` collects L001–L005 under a single theorem; zero sorry, zero custom physics axioms.

Physical motivation. Since b_p is odd (L001), the product $Z \cdot b_p$ carries Z 's parity (L003). This is the GTE-theoretic basis for the proton-parity stability feature

$$F_{10} = (Z \bmod 2) \delta_{\text{pair}} A^{-3/4},$$

where $\delta_{\text{pair}} = \pm 1$ for even/odd- A paired nuclei. F_{10} penalizes odd- Z odd- N nuclei—correctly predicting even- A Tc (e.g. Tc-98) and Pm (e.g. Pm-142) as unstable without affecting even- Z neighbors.

Pairing constant identity.

Theorem 7.13 (GTE pairing-constant algebraic identity). $a_{\text{seed}} \cdot \sqrt{a_{\text{seed}}} = \sqrt{a_{\text{seed}}^3}$, i.e. $5\sqrt{5} = \sqrt{125}$.

Lean: `pairing_sqrt_identity : (gte_a_seed : Real)* sqrt gte_a_seed = sqrt (gte_a_seed ^ 3)` (in `GTE.NuclearPairing`, `LEAN4norm_num` + `LEAN4Real.sqrt_mul`)

This underpins the GTE prediction $\Delta_{\text{pair}} = a_{\text{seed}}^{g/2} = 5^{3/2} = 11.18$ MeV for the SEMF nuclear pairing constant ([B] bridge claim in P03). The numerical value $5\sqrt{5} = 11.1803\dots$ MeV matches the Møller et al. (1995) SEMF fit within 0.003% with no free parameters. The physical mechanism (why $a_{\text{seed}}^{g/2}$ equals the BCS pairing scale) is open theoretical work.

7.7 $N_c = 3$ from substrate arithmetic (`NcColorArithmetic`)

The `GTE.NcColorArithmetic` module (209 lines, 10 theorems) derives the QCD colour count $N_c = 3$ from GTE substrate arithmetic alone, without importing Standard Model winding assignments. Two independent routes are certified:

Theorem 7.14 (Route 1 — Mersenne GCD, zero `sorry`, zero custom axioms). *The GTE orbit at $n = 10$ produces c -values $c_2 = 2^{10} - 1$ and $c_3 = 2^{16} - 1$. By the Mathlib identity `Nat.pow_sub_one_gcd_pow_sub_one`:*

$$\gcd(2^{10} - 1, 2^{16} - 1) = 2^{\gcd(10,16)} - 1 = 2^2 - 1 = 3.$$

Lean name: `nc_eq_3_from_mersenne_gcd`.

Theorem 7.15 (Route 2 — Ridge factorial uniqueness, zero `sorry`). *The ridge parameters satisfy $\gcd(b_2, q_2) = \gcd(42, 24) = 6 = 3!$. N_c is the unique positive integer n with $n! = \gcd(b_2, q_2)$.*

Lean name: `nc_uniqueness_from_ridge_divisors`.

Lean: `nc_eq_3_from_mersenne_gcd`, `nc_uniqueness_from_ridge_divisors`, `nc_eq_3_from_substrate` (in `GTE.NcColorArithmetic`, 10 theorems, zero `sorry`, zero custom axioms; standard Lean/Mathlib kernel only)

8 Braid Atlas Charge Structure

The `BraidAtlas.ChargeTheorem` module proves that electric charge is recovered from braid winding numbers via $Q = W_g/N_c$.

Provenance. The winding numbers $W(e) = -3$, $W(\nu) = 0$, $W(u) = +2$, $W(d) = -1$ are inputs from the Braid Atlas topological analysis [11]: they are the writhe-based invariants of the four SM fermion braid genotypes as established in the companion paper, not derived from the UGP ridge arithmetic within this module. The value $N_c = 3$ is supplied by `SM_gauge_uniquely_selected` in `TE22.ScanCertificate` (PSC axioms). The module's contribution is to *formally prove* that these assigned winding values reproduce exact SM charges and that anomaly cancellation in the winding language forces $N_c = 3$.

Theorem 8.1 (Electric charge from winding number). *At $N_c = 3$, with SM fermion winding numbers $W(e) = -3$, $W(\nu) = 0$, $W(u) = +2$, $W(d) = -1$, the formula $Q = W_g/N_c$ holds for all four SM fermion types:*

$$Q(e) = -1, \quad Q(\nu) = 0, \quad Q(u) = +2/3, \quad Q(d) = -1/3.$$

Lean: `charge_from_winding_Nc3`, `winding_distinguishes_types_Nc3` (in `BraidAtlas.ChargeTheorem`)

Theorem 8.2 (Anomaly cancellation forces $N_c = 3$). *The per-generation winding sum satisfies:*

$$\sum_{f \in \text{generation}} W_g(f) = N_c(N_c - 3).$$

This equals zero if and only if $N_c = 3$ (for $N_c > 0$). SM anomaly cancellation in the winding language is therefore equivalent to the QCD colour rank taking its observed value.

Lean: `anomaly_cancellation_forces_Nc_3`, `perGenWindingSum_eq`, `winding_sum_zero_at_Nc3`
(in `BraidAtlas.ChargeTheorem`)

Refined naming and corollaries (paper-citation aliases). The `ChargeTheorem` module exposes the following aliases used in companion papers P01, P02, P17:

- `sm_charge_leptons` — $Q = W_g/N_c$ specialised to colour-singlet fermions: charged leptons ($Q = -1$) and neutrinos ($Q = 0$).
- `sm_quarks_fractional_charge` — $N_c \nmid 2$ and $N_c \nmid -1$ at $N_c = 3$, certifying the fractional quark charges $+\frac{2}{3}$ and $-\frac{1}{3}$.
- `nc_eq_3_from_fractional_charge` — corollary form: anomaly cancellation under $N_c > 0$ forces $N_c = 3$.

GTE-P7 mirror-branch dark matter (formal derivation). The GTE-P7 candidate from the mirror sector of the Braid Atlas v2.0 (P17, §subsec:mirror_dm) is a sterile colour-singlet, electrically neutral spin- $\frac{1}{2}$ Dirac fermion. The `ChargeTheorem` module provides the formal derivation:

- `mirror_winding_number_zero` — explicit axiom (faithfully disclosed in P17 `PROVENANCE.md`): the mirror-branch winding number $W_{g,\text{mirror}} = 0$. Justification is the full P17 braid topology framework, which is not yet formalised in Lean.
- `gmn_color_singlet_neutral` — Gell-Mann–Nishijima instantiation: for a colour singlet ($T_3 = 0$) with vanishing hypercharge ($Y = 0$), $Q = T_3 + Y/2 = 0$. Provides the algebraic step from $W_{g,\text{mirror}} = 0$ to $Q_{\text{GTE-P7}} = 0$.
- `gte_p7_electric_charge_zero` — formal derivation: $Q_{\text{GTE-P7}} = W_{g,\text{mirror}}/N_c = 0/3 = 0$.
- `gte_p7_quantum_numbers_neutral` — composite assignment: electrically neutral, colour singlet, sterile. Combined with the arithmetic backbone in `GTE.GeneralTheorems` (`mirror_triple_residue`, `mirror_prime_2137`, `mirror_quotient_q1`, `mirror_triple_prime_lock`), the mirror triple $(1, 73, 2137; g=1)$ shares the same prime-lock residue $m_1 = 20$ as the canonical lepton triple, placing it in the same algebraic orbit.

This closes the dark-matter prediction in P17 (§subsec:mirror_dm) at **Category B** status: the quantum numbers are derived from the braid topology axiom, the rest of the chain is fully Lean-certified.

8.1 Composite braid c-rule

The `BraidAtlas.CompositeTriples` module extends the charge–winding framework to composite hadrons, proving three structural properties that ground the reflexive-closure correlation of Ref. [16].

Writhe composition. The winding number is additive; CPT conjugation reverses the sign for antiquarks.

Theorem 8.3 (Meson winding is zero). *For any SM fermion type f , the quark-antiquark (meson) winding sum vanishes:*

$$W_g(q) + W_g(\bar{q}) = W_g(q) - W_g(q) = 0.$$

Consequently, the chirality encoding $H(Wr) = 0$ and the c -component of any meson GTE triple is real and positive.

Lean: `meson_winding_zero, meson_c_real_positive, baryon_c_real_of_nonneg_charge` (in `BraidAtlas.CompositeTriples`)

Proton and neutron winding. As a concrete verification that reproduces the GTE cascade result $c = 15$ for the proton and neutron (Paper 01 Appendix A):

Theorem 8.4 (Proton and neutron winding at $N_c = 3$).

$$\begin{aligned} W(u) + W(u) + W(d) &= 3 = N_c \cdot Q(p), \\ W(u) + W(d) + W(d) &= 0 = N_c \cdot Q(n). \end{aligned}$$

The confinement Mersenne level $c = 15 = 2^4 - 1$ is the lowest Mersenne composite level and is consistent with $c < 1023$ (tier-1 assignment, Lean-certified).

Lean: `proton_winding_eq_Nc_times_Q, neutron_winding_eq_Nc_times_Q, confinement_mersenne_level, proton_neutron_c_eq_15_in_confinement_tier` (in `BraidAtlas.CompositeTriples`)

Complete composite triple derivation (Category A). The theorem `ugp_composite_braid_c_rule` (zero sorry) bundles the winding and chirality structural facts. Beyond this, the module now proves the **complete (a,b,c) triple for all 9 light baryons**:

- *a-component:* $a = \min_i a(q_i)$ (`proton_a_eq_min`).
- *c-component:* $c = \pm 1023$ for strange baryons (sign from $W = N_c Q$; `ugp_strange_baryon_c_values`).
- *b-component:* Nine formulas, all zero sorry: `ugp_nucleon_b_formula` ((p,n)), `ugp_strange_baryon_b_formulas` ($(\Lambda, \Xi^0, \Omega^-)$), `sigma_plus_b_formula` ((Σ^+)), `ugp_all_baryon_b_formulas` (full conjunction). The Σ^+ formula $= (b_s + a_u \text{ seesaw}) \times (b_s + a_u \text{ seesaw} + (a_u(N_c^2 - 1))^2) = 639161$ involves $b_s = 186$, $\text{seesaw} = 29 = 4N_c^2 - \delta$, and $N_c^2 - 1 = 8$ (number of gluons). All b-formulas for strange baryons involve $11 = b(\nu_{\mu R})$ (Braid Atlas seesaw b-value) and $29 = 4N_c^2 - \delta$ (neutrino seesaw numerator), connecting strange baryons to the neutrino sector.

Status: **Category A** — all baryon triple derivations are zero sorry and derived from Lean-certified UGP constants. The composite triple program is complete with no remaining open items.

8.2 SM Winding Numbers from N_c (`BraidAtlas.ChargeDerivation`)

The module `BraidAtlas.ChargeDerivation` derives the SM winding-number assignments $\{W_g\} = \{N_c - 1, -1, 0, -N_c\}$ from $N_c = 3$ alone, closing the final link in the Braid Atlas charge programme.

The argument uses the $SU(N_c+2)$ embedding: PSC anomaly cancellation forces $SU(5)$ as the minimal GUT group; its $(\bar{5} + 10)$ fermion content decomposes under $SU(3) \times SU(2) \times U(1)$ with hypercharges $Y_{Q_L} = 1/(2N_c)$. The Braid Atlas winding number $W_g = 2N_c \cdot Y_g$ then gives $\{W_g\} = \{2, -1, 0, -3\}$ for $N_c = 3$ — the complete SM charge spectrum.

Theorem 8.5 (SM Winding Numbers from N_c). *The SM winding numbers $\{N_c-1, -1, 0, -N_c\} = \{2, -1, 0, -3\}$ follow from $N_c = 3$ via the $SU(N_c + 2)$ decomposition. Anomaly cancellation $N_c(N_c - 1) + N_c(-1) + 0 + (-N_c) = 0$ is certified. The hypercharge $Y_{Q_L} = 1/(2N_c)$ governs both the VV exponent β_d and the SM charge spectrum through a single algebraic path.*

Lean: `sm_winding_numbers_from_Nc, anomaly_cancellation_from_windings, y ql_unifies_vv_and_winding` (in `BraidAtlas.ChargeDerivation`, zero sorry)

8.3 Coxeter–Conductor Theorem (`BraidAtlas.CoxeterConductor`, `BraidAtlas.CoxeterConductorTowerLaw`)

The pair of modules `BraidAtlas.CoxeterConductor` and `BraidAtlas.CoxeterConductorTowerLaw` certify the arithmetic core of the *Coxeter–conductor theorem*: the affine Toda mass spectrum of a simple Lie algebra G lies in $\mathbb{Q}(\zeta_{120})$ iff its Coxeter number $h(G)$ divides 120.

Positive direction. For $G \in \{E_8, E_6, F_4, G_2, B_4\}$ with $h \in \{30, 12, 6, 8\}$, each Coxeter number divides 120 and $\text{lcm}(30, 12, 8, 6, 3, 2, 1) = 120$, so $\mathbb{Q}(\zeta_{120})$ is the minimal cyclotomic field containing every physically realised Toda spectrum (`positive_coxeter_conductor`).

E_7 falsifier. E_7 has $h = 18 \nmid 120$. The mass ratios involve $\mathbb{Q}(\cos(\pi/9)) = \mathbb{Q}(\zeta_{18})^+$ of degree $\varphi(18)/2 = 3$ over $\mathbb{Q}$, while $[\mathbb{Q}(\zeta_{120}) : \mathbb{Q}] = \varphi(120) = 32$. Since $3 \nmid 32$, the Tower Law forbids the embedding $\mathbb{Q}(\cos(\pi/9)) \subseteq \mathbb{Q}(\zeta_{120})$.

Theorem 8.6 (Coxeter–Conductor, E_7 Tower-Law obstruction). *The minimal polynomial $8X^3 - 6X - 1$ of $2\cos(\pi/9)$ over $\mathbb{Q}$ is irreducible, the quotient field $\mathbb{Q}[X]/\langle 8X^3 - 6X - 1 \rangle$ has $\mathbb{Q}$ -dimension exactly 3, and $\varphi(120) = 32$ is not divisible by 3. Therefore E_7 Toda masses lie outside $\mathbb{Q}(\zeta_{120})$.*

Lean: `p_rat_irreducible, finrank_p_rat_quot, e7_tower_law_complete, e7_coxeter_conductor_obstruction` (in `BraidAtlas.CoxeterConductorTowerLaw/CoxeterConductor`, zero sorry)

The proof uses Mathlib’s rational-root theorem (`num_dvd_of_is_root, den_dvd_of_is_root`) to verify all eight rational candidates $\pm 1, \pm \frac{1}{2}, \pm \frac{1}{4}, \pm \frac{1}{8}$, then concludes irreducibility from `Polynomial.irreducible_of_degree_le_three_of_not_isRoot`. The dimension claim uses `finrank_quotient_span_eq_natDegree` on the quotient form $\mathbb{Q}[X]/\langle p \rangle$, which avoids the `AdjoinRoot` typeclass-synthesis obstruction noted in the companion lab notes.

8.4 Coxeter–Conductor Biconditional (`CoxeterBiconditional`)

The new module `CyclotomicCompleteness.CoxeterBiconditional` (graduated from sandbox, SPEC_042_CYX Track 1) formalises the arithmetic backbone of the biconditional criterion: a root system’s Toda masses lie in $\mathbb{Q}(\zeta_{120})$ iff its Coxeter number h divides 60.

Arithmetic biconditional. The core lemma (proved by `omega`) is:

$$2h \mid 120 \iff h \mid 60.$$

This is the formal arithmetic equivalent of the field-containment criterion $\mathbb{Q}(\cos(\pi/h))^+ \subseteq \mathbb{Q}(\zeta_{120})$. The full field-embedding is now formalised in the companion module `CyclotomicCompleteness.CyclotomicContainment` (see §8.5).

Per-algebra certificates. Five `norm_num` certificates confirm the positive cases:

$$\begin{aligned} G_2 : h = 6 \mid 60, \quad F_4 : h = 12 \mid 60, \quad E_6 : h = 12 \mid 60, \\ E_8 : h = 30 \mid 60. \end{aligned}$$

For B_4 ($h = 8$): $8 \nmid 60$, but the actual mass ratios lie in $\mathbb{Q}(\sqrt{2}) = \mathbb{Q}(\zeta_8)^+$ and $8 \mid 120$, so B_4 masses are in $\mathbb{Q}(\zeta_{120})$ via the conductor 8 (not 16). The module proves $8 \mid 120 \wedge \neg(16 \mid 120)$ to pin the actual conductor.

E_7 double failure. The theorem `e7_double_failure` (zero `sorry`) proves:

$$\neg(18 \mid 60) \wedge \neg(18 \mid 120).$$

This is the arithmetic complement to `e7_tower_law_complete` in `BraidAtlas.CoxeterConductorTowerLaw`: divisibility fails at both the $h \mid 60$ and $2h \mid 120$ levels before the Tower Law argument is even needed.

Master theorem. The composite theorem `coxeter_biconditional_summary` packages all of the above:

Theorem 8.7 (Coxeter–Conductor Biconditional, arithmetic backbone). $6 \mid 60$ (G_2), $12 \mid 60$ (F_4, E_6), $30 \mid 60$ (E_8), $8 \mid 120$ (B_4 conductor), $\neg(18 \mid 60)$ (E_7), $\neg(18 \mid 120)$ (E_7), and $120 = 2 \times 60$.

Lean: `g2_h_dvd_60`, `f4_h_dvd_60`, `e6_h_dvd_60`, `e8_h_dvd_60`, `b4_mass_field_conductor`, `e7_h_not_dvd_60`, `e7_double_failure`, `two_h_dvd_120_iff_h_dvd_60`, `coxeter_biconditional_summary` (in `CyclotomicCompleteness.CoxeterBiconditional`, all zero `sorry`)

Relationship to existing modules. This module imports `BraidAtlas.CoxeterConductor` (specifically `e7_coxeter_not_dvd`) and adds the $h \mid 60$ formulation and the biconditional $h \mid 60 \iff 2h \mid 120$. Together, the four modules `CoxeterConductor`, `CoxeterConductorTowerLaw`, `CoxeterBiconditional`, and `CyclotomicContainment` form a complete arithmetic + degree-theoretic + field-embedding certification of the $\mathbb{Q}(\zeta_{120})$ filter.

8.5 Cyclotomic field embedding (`CyclotomicContainment`)

The module `CyclotomicCompleteness.CyclotomicContainment` (SPEC_042_CYX Track C) proves the field-theoretic content of the Coxeter–conductor theorem: for every $h \mid 60$, the cyclotomic field $\mathbb{Q}(\zeta_{2h})$ embeds as a $\mathbb{Q}$ -algebra into $\mathbb{Q}(\zeta_{120})$. All theorems have zero `sorry` and use only Mathlib’s `IsCyclotomicExtension`, `IsPrimitiveRoot`, and `IsSplittingField` infrastructure.

Theorem A: primitive root existence.

Theorem 8.8 (`cyclotomic120_contains_primitive_root`). *For every $h \mid 60$, $\text{CyclotomicField } 120 \mathbb{Q}$ contains a primitive $2h$ -th root of unity:*

$$\forall h, h \mid 60 \implies \exists \zeta \in \mathbb{Q}(\zeta_{120}), \text{IsPrimitiveRoot } \zeta(2h).$$

Lean: `cyclotomic120_contains_primitive_root` (in `CyclotomicCompleteness.CyclotomicContainment`, zero sorry)

The proof extracts a root ζ_0 of Φ_{120} from `Polynomial.SplittingField.splits`, converts it to an `IsPrimitiveRoot` certificate for ζ_{120} via `isRoot_cyclotomic_iff_charZero`, and powers down to a primitive $2h$ -th root using `IsPrimitiveRoot.pow` and the divisibility $2h \mid 120$ (supplied by `CoxeterBiconditional`).

Theorem B: $\mathbb{Q}$ -algebra embedding.

Theorem 8.9 (`cyclotomic_field_embedding`). *For every h with $0 < h$ and $h \mid 60$, there exists an injective $\mathbb{Q}$ -algebra homomorphism*

$$\mathbb{Q}(\zeta_{2h}) \hookrightarrow \mathbb{Q}(\zeta_{120}).$$

Lean: `cyclotomic_field_embedding` (in `CyclotomicCompleteness.CyclotomicContainment`, zero sorry)

The proof uses the primitive root ζ from Theorem A together with `IsCyclotomicExtension.union_of_isPrimitiveRoot` to promote $\mathbb{Q}(\zeta_{120})$ to a $\{120\} \cup \{2h\}$ -cyclotomic extension, then `IsCyclotomicExtension.splits_cyclotomic` to show that Φ_{2h} splits there, and finally the universal property `Polynomial.IsSplittingField.lift` to produce the embedding. A technical subtlety: the `IsCyclotomicExtension {120}` instance is invoked by name (`CyclotomicField.isCyclotomicExtension`) with an explicit `NeZero (120 :  $\mathbb{Q}$ )` witness to bypass the `[CharZero K]` definitional loop for concrete numerals.

Per-algebra certificates. Three `AlgHom` certificates instantiate Theorem B for the exceptional algebras that pass the Coxeter–conductor filter:

$$\begin{aligned} G_2 : h = 6, \quad \mathbb{Q}(\zeta_{12}) &\hookrightarrow \mathbb{Q}(\zeta_{120}) \quad (\text{g2_cyclotomic_embedding}), \\ F_4, E_6 : h = 12, \quad \mathbb{Q}(\zeta_{24}) &\hookrightarrow \mathbb{Q}(\zeta_{120}) \quad (\text{f4_e6_cyclotomic_embedding}), \\ E_8 : h = 30, \quad \mathbb{Q}(\zeta_{60}) &\hookrightarrow \mathbb{Q}(\zeta_{120}) \quad (\text{e8_cyclotomic_embedding}). \end{aligned}$$

Relationship to other `CyclotomicCompleteness` modules. `CyclotomicContainment` imports `CoxeterBiconditional` for the arithmetic backbone ($h \mid 60 \iff 2h \mid 120$). The negative direction ($h \nmid 60 \Rightarrow$ no embedding) is handled by the Tower–Law argument in `CoxeterConductorTowerLaw`. Together, these four modules provide a machine-checked derivation of the complete $\mathbb{Q}(\zeta_{120})$ field-theoretic filter for the exceptional root systems.

8.6 Electroweak boson c-values (`BraidAtlas.EWBosons`)

The `BraidAtlas.EWBosons` module derives the electroweak massive boson c-values $c(W) = 11$, $c(Z) = 12$, $c(H) = 13$ from two structural identities at the canonical ridge level $n = 10$:

$$\begin{aligned} q_2(\text{canonical}) &= 2N_c(N_c + 1), \\ \text{ugp}1_g &= N_c(N_c + 1) + 1. \end{aligned}$$

Substituting $N_c = 3$ yields the consecutive-integer window $\{c(W), c(Z), c(H)\} = \{N_c(N_c + 1) - 1, N_c(N_c + 1), N_c(N_c + 1) + 1\} = \{11, 12, 13\}$ centred on $N_c(N_c + 1) = 12$.

Triangular-number unification. Both structural identities admit a single geometric reading via the triangular number $T(N_c) = N_c(N_c + 1)/2$: the EW c-window is centred on $2T(N_c)$, with $q_2(\text{canonical}) = 4T(N_c)$ and $\text{ugp}1_g = 2T(N_c) + 1$.

Theorem 8.10 (EW boson c-value derivation). *At $N_c = 3$ and the unique RSUC fixed point $(s, g, t) = (7, 13, 20)$, the canonical ridge factorisation $42 \times 24 = 1008 = R_{10}$ and the gap identity $\text{ugp}1_g = 13$ force*

$$c(W) = 11, \quad c(Z) = 12, \quad c(H) = 13,$$

which is the unique consecutive triple drawn from the orbital constants at $n = 10$.

Lean: `ew_boson_c_value_theorem`, `ew_c_consecutive_window`, `ew_c_window_unique`, `ew_triangular_unification`, `ew_ugp_axiom_derivation` (in `BraidAtlas.EWBosons`, zero sorry)

This closes the last Category B item in the Braid Atlas (P17): the charged-current (W), neutral-current (Z), and Higgs c-values are now derived from UGP axioms rather than postulated.

8.7 Chirality squaring (`BraidAtlas.ChiralitySquaring`)

The module `BraidAtlas.ChiralitySquaring` certifies the arithmetic signature of the V–A chiral structure that underpins the $T/T^\dagger$ pairing in the Galois-protection non-renormalization theorem (§12).

Theorem 8.11 (Chirality arithmetic signature). *The numerator $g_3^{2,\text{bare}} = (13 \times 17 \times 29)^2 = 41,075,281$ is a perfect square, while $g_2^{2,\text{bare}} = 17 \times 137 = 2,329$ is not. This asymmetric arithmetic signature distinguishes the $SU(3)$ and $SU(2)$ sectors at the GTE-derived gauge couplings.*

Lean: `chirality_arithmetic`, `g3_numerator_is_perfect_square`, `g3_numerator_sqrt`, `g2_numerator_not_perfect_square`, `two_three_two_nine_not_nat_sq` (in `BraidAtlas.ChiralitySquaring`, zero sorry)

The non-squareness of 2,329 is proved by bracketing: $48^2 = 2304 < 2329 < 2401 = 49^2$, followed by a finite `native_decide` check over $k \in [0, 48]$.

8.8 Mirror winding number (`BraidAtlas.MirrorWindingNumber`)

The module `BraidAtlas.MirrorWindingNumber` provides the structural Lean certificate for MBA-5: the elimination of the `mirror_winding_number_zero` axiom from `BraidAtlas.ChargeTheorem`. The axiom is replaced by a definition

$$W_{g,\text{mirror}} := N_c \times (T_{3,\text{mirror}} + Y_{\text{mirror}}/2) = 3 \times (0 + 0/2) = 0,$$

where `mirrorT3Int = 0` and `mirrorYInt = 0` encode the P17 statement that mirror-branch particles carry no SM gauge charges. The theorem `mirror_winding_number_zero` is now proved by `simp`, with zero axioms.

Orbit distinguishability. The module also certifies that the mirror GTE orbit $(a_1, b_1, c_1) = (1, 73, 2137)$ is arithmetically distinct from every canonical SM fermion orbit: the mirror prime $c'_1 = 2137$ satisfies $2137 \neq 823$ (charged leptons) and the mirror quotient $q'_1 = 29 \neq 11$ (canonical SM), placing the mirror branch in a strictly higher GTE orbit depth.

Generation independence. The gauge-singlet argument ($T_3 = 0$, $Y = 0$) is independent of generation index $g \in \{1, 2, 3\}$; the module certifies $W_{g,\text{mirror}} = 0$ for all three mirror-lepton generations simultaneously.

Theorem 8.12 (Mirror winding number — structural certificate). *The mirror GTE orbit $(1, 73, 2137)$ is a gauge singlet. With $T_{3,\text{mirror}} = Y_{\text{mirror}} = 0$ encoded by definition,*

$$W_{g,\text{mirror}} = N_c \times (T_{3,\text{mirror}} + Y_{\text{mirror}}/2) = 0$$

*for all three mirror generations. The prime $c'_1 = 2137$ and quotient $q'_1 = 29$ are also machine-certified. All 13 theorems carry zero **sorry** and depend on zero additional axioms beyond Lean/Mathlib.*

Lean: `mirror_winding_zero_via_gmn, mirror_winding_generation_independent, mirror_orbit_q1_eq_29, mirror_prime_certified, mba5_mirror_winding_certificate` (in `BraidAtlas.MirrorWindingNumber`, 13 theorems, zero **sorry**, zero axioms; MBA-5, 2026-05-16)

8.9 EW boson–RHN cross-sector connection (BraidAtlas.EWBosonRHNConnection)

The module `BraidAtlas.EWBosonRHNConnection` is the first cross-sector structural unification result in UGP Physics: it proves that the W -boson c -value equals the right-handed neutrino (RHN) b_2 -value in *both* the standard and mirror branches, with the equality mediated by the GTE orbit quotient q_1 .

In the standard branch, $c(W) = 11 = \text{gteQuotient}(823, 73) = b_2(\text{RHN}) = 11$. In the mirror branch, $c(W') = 29 = \text{gteQuotient}(2137, 73) = b'_2(\text{RHN}) = 29$. The mirror shift $\Delta q = q_2^{\text{mirror}} - q_2^{\text{canonical}} = 42 - 24 = 18$ propagates uniformly: $c(W') - c(W) = b'_2(\text{RHN}) - b_2(\text{RHN}) = 18$.

This module certifies P17 [11] (§subsec:ew_bosons, line: $c(W) = 11 = b(\nu_{\mu R})$) in machine-checked Lean 4.

Theorem 8.13 (EW–RHN cross-sector connection (both branches)). *In the canonical and mirror branches,*

$$c(W) = \text{gteQuotient}(823, 73) = b_2(\text{RHN}) = 11,$$

$$c(W') = \text{gteQuotient}(2137, 73) = b'_2(\text{RHN}) = 29,$$

*with the branch shift $\Delta q = 18$ preserved across both equalities. All 10 theorems carry zero **sorry** and zero axioms.*

Lean: `ew_rhn_connection_standard, ew_rhn_connection_mirror, ew_rhn_connection_both_branches, branch_shift_preserved, ew_rhn_structural_connection` (in `BraidAtlas.EWBosonRHNConnection`, 10 theorems, zero **sorry**, zero axioms; 2026-05-17)

8.10 Dark sector certification hub (BraidAtlas.DarkBraidAtlas)

The module `BraidAtlas.DarkBraidAtlas` collects and cross-certifies the machine-checked structural facts about the mirror branch (dark sector) established across MBA rounds 1–5. It imports `MirrorWindingNumber`, `EWBosons`, `EWBosonRHNConnection`, `RHNGapTheorem`, and `DarkQuarkCharge`, re-exporting their results as a single coherent certificate.

Theorem 8.14 (Dark sector structural certificate). *The following facts are machine-checked at zero **sorry**, zero axioms:*

1. **Dark lepton neutrality:** $Q_{\text{EM}} = 0$ for all three mirror-lepton generations (re-exported from `MirrorWindingNumber`).
2. **RHN b -values:** $b'_1(\text{RHN}) = 5$ (colour-counting, branch-invariant); $b'_2(\text{RHN}) = q'_1(\text{mirror}) = 29$ (GTE arithmetic).
3. **Dark EW connection:** $c(W') = b'_2(\text{RHN}) = 29$ (from `EWBosonRHNConnection`).

4. **Sector separation:** $b'_2(\text{RHN}) = 29 \neq 11 = b_2(\text{RHN})$; the two branches are arithmetically distinguished.

The 16-theorem conjunction `dark_braid_atlas_certificate` is the canonical single-import certificate for all Lean-verified dark sector facts.

Lean: `dark_lepton_q_em_zero`, `dark_rhn_b1`, `dark_rhn_b2`, `dark_ew_rhn_connection`, `dark_braid_atlas_certificate` (in `BraidAtlas.DarkBraidAtlas`, 16 theorems, zero `sorry`, zero axioms; 2026-05-17)

8.11 RHN b_3 gap theorem (`BraidAtlas.RHNGapTheorem`)

The module `BraidAtlas.RHNGapTheorem` certifies the numerical observation that the RHN b_3 -value satisfies

$$b_3 = b_2 + (N_c^2 - 1) = b_2 + 8$$

in both branches:

$$\text{Standard: } 19 = 11 + 8, \quad \text{Mirror: } 37 = 29 + 8,$$

where $N_c^2 - 1 = 8 = \dim(\text{SU}(3)_{\text{adj}})$, the number of gluon degrees of freedom.

Scientific status. These are arithmetic facts proved by `norm_num`. The gap $N_c^2 - 1$ is algebraically equivalent to $b_3 = q_2 - b_1$ (standard: $24 - 5 = 19$; mirror: $42 - 5 = 37$) but neither formula has been derived from a GTE cascade rule; the pattern was identified post-hoc (MBA-3, 2026-05-17). This result is therefore **Category D (numerical, not structural)**. Upgrading to Category A requires identifying and certifying the GTE cascade rule that produces $b_3 = q_2 - b_1$ from first principles.

Theorem 8.15 (RHN b_3 gap — numerical certificate (Cat D)). *The identity $b_3 = b_2 + (N_c^2 - 1)$ holds in both branches: $19 = 11 + 8$ and $37 = 29 + 8$, with $N_c^2 - 1 = 8$. Both b_3 values are prime. All theorems are proved by `norm_num` / `decide`, zero `sorry`.*

Lean: `su3_adjoint_dim`, `rhn_b3_gap_numerical_standard`, `rhn_b3_gap_numerical_mirror`, `rhn_gap_equals_su3_adj`, `rhn_b3_both_prime` (in `BraidAtlas.RHNGapTheorem`, 9 theorems, zero `sorry`; Cat D — structural derivation open; 2026-05-17)

8.12 Dark quark charge (`BraidAtlas.DarkQuarkCharge`)

The module `BraidAtlas.DarkQuarkCharge` certifies MBA-6: dark quarks carry $Q_{\text{EM}} = 0$. This resolves the question of whether the mirror $b_2 \leftrightarrow q_2$ arithmetic symmetry maps SM quark gauge charges (which are non-zero: $Q(u) = +2/3$, $Q(d) = -1/3$) to corresponding mirror-branch charges.

The resolution (Argument B) invokes the P17 universal statement: mirror-branch particles carry *no SM gauge charges*. Therefore $Y_{\text{mirror}} = 0$ and $T_{3,\text{mirror}} = 0$ for all mirror-branch particles including quarks, giving $W_{g,\text{mirror}} = N_c \times (T_3 + Y/2) = 0$ and hence $Q = 0$. Dark quarks do carry dark $\text{SU}(3)_{\text{dark}}$ colour and confine at Λ_{dark} , but are electromagnetically neutral.

The proof is identical in structure to the lepton case (MBA-1/MBA-5, `MirrorWindingNumber`): it reuses `mirror_winding_number_zero` because the same definition of `W_g_mirror` applies regardless of quark vs. lepton sector.

Theorem 8.16 (Dark quark electric neutrality (MBA-6)). *For all three dark-quark generations, and for both dark up- and dark down-type quarks,*

$$Q_{\text{EM}}(\text{dark quark}) = W_{g,\text{mirror}}/N_c = 0.$$

The module also machine-certifies that dark quark winding differs from the SM values ($W_g(u) = +2$, $W_g(d) = -1$), confirming sector separation. All 8 theorems carry zero `sorry` and zero axioms.

Lean: `dark_quark_charge_zero, dark_quark_charge_zero_all_gens, dark_sector_all_neutral, dark_up_quark_differs_from_sm_up, mba6_dark_quark_charge_certificate` (in `BraidAtlas.DarkQuarkCharge`, 8 theorems, zero sorry, zero axioms; MBA-6, 2026-05-17)

8.13 $\mathbb{Z}_7$ dark charge arithmetic (`BraidAtlas.DarkBraidAtlas`, §9)

The module `BraidAtlas.DarkBraidAtlas` §9 certifies the $\mathbb{Z}_7$ dark charge arithmetic derived from CUP-11a [?]. The key values are $b_2^{\text{SM}} = 42$ (SM branch divisor) and $b'_2 = 24$ (mirror/dark branch divisor).

Theorem 8.17 ($\mathbb{Z}_7$ dark charge certificate, grade A). *Let $N_c = 3$ (QCD colour rank). Then:*

1. SM branch $\mathbb{Z}_7$ charge is zero: $42 \bmod 7 = 0$.
2. Dark sector $\mathbb{Z}_7$ charge equals N_c : $24 \bmod 7 = N_c = 3$.
3. Dark baryon $\mathbb{Z}_7$ charge: $3 \times 3 \bmod 7 = 2$ (three dark quarks, each with $\mathbb{Z}_7$ charge N_c).
4. GTB 2/7 numerator: $N_c^2 \bmod 7 = 2$, certifying the structural basis of $\eta_\chi = (2/7) \eta_{B+L, \text{pre}}$.
5. Sector separation: $b_2^{\text{SM}} \bmod 7 \neq b'_2 \bmod 7$.

Lean: `sm_dark_charge_zero, mirror_dark_charge_eq_three, dark_charge_equals_Nc, dark_baryon_z7_charge, gtb_two_sevenths_numerator, z7_sector_separation, z7_dark_charge_certificate` (in `BraidAtlas.DarkBraidAtlas` §9, 7 theorems, zero sorry, zero axioms; CUP-11a, 2026-05-17)

8.14 Dark SU(3) gauge coupling (`BraidAtlas.DarkBraidAtlas`, §10)

The module `BraidAtlas.DarkBraidAtlas` §10 certifies the branch-independence of the bare SU(3) gauge coupling. The SU(3) bare coupling $g_3^2 = 41075281/27648000$ (from `Phase4.GaugeCouplings`) is derived from the squared Vandermonde discriminant D_3 , which depends only on the colour rank $N_c = 3$ and the SU(3) root system — not on the orbital branch (canonical vs. mirror). Both branches therefore share the same bare g_3^2 at the UGP unification scale.

Theorem 8.18 (Dark SU(3) coupling equals SM SU(3) coupling, grade A). *At the UGP unification scale, the dark SU(3) bare coupling equals the SM SU(3)_c bare coupling:*

$$g_{3, \text{dark}}^2 = g_{3, \text{SM}}^2 = \frac{41075281}{27648000}.$$

The physical confinement scales $\Lambda_{\text{dark}} \neq \Lambda_{\text{QCD}}$ because the dark sector carries different particle content, changing the one-loop RGE β -function. The equality holds at the bare (unification-scale) level.

Lean: `g3Sq_dark_bare, dark_su3_coupling_equals_sm` (in `BraidAtlas.DarkBraidAtlas` §10, zero sorry, zero axioms; 2026-05-17)

8.15 Dark SU(3) gauge coupling — master formula derivation (`BraidAtlas.DarkGaugeCoupling`)

The module `BraidAtlas.DarkGaugeCoupling` certifies the full derivation chain from the Elegant Kernel to the dark SU(3) bare coupling via the Gauge Coupling Master Formula (P01 [?] §4, eq. `gauge_master`):

$$g_G^2 = \frac{L_G \cdot D_G}{5\gamma_G},$$

where L_G is the Weyl-group order, D_G is the discrete invariant from the Elegant Kernel, and γ_G is the golden-field dimension exponent.

Theorem 8.19 (Dark $SU(3)$ gauge coupling master formula, grade A/C). *For $SU(3)_{\text{dark}}$, the master formula parameters are: $L = 6$ (Weyl group $|S_3|$, grade A), $D = \text{Vand}^2(k_a, k_b, k_c) = 41075281/1327104$ (from Elegant Kernel, grade A), $\gamma = 3$ (golden-field exponent, grade A). The formula gives $g_{3,\text{dark}}^2 = 6 \times D/125 = 41075281/27648000 = g_{3,\text{SM}}^2$.*

The overall grade is Cat A/C: the formula and all three parameters are individually Lean-certified [A]; the branch-independence of D (same Elegant Kernel triple for both SM and mirror orbits) is argued from MDL uniqueness [C] and remains the main open step for a full Cat A certificate.

Lean: su3_weyl_order_eq_six, su3_golden_field_exponent_eq_three, vandermonde_is_su3_coupling_root, gauge_master_su3_arithmetic, gauge_master_su3_from_vandermonde, dark_su3_weyl_order_branch_independent, dark_gauge_coupling_certificate (in BraidAtlas.DarkGaugeCoupling, 8 theorems, zero sorry, zero axioms; Cat A/C; 2026-05-17)

9 RCC Infinite Families (PSC.RCCInfiniteFamilies)

The module `PSC.RCCInfiniteFamilies` closes the Residual Classification Conjecture (RCC) over all four infinite classical Lie-algebra families: B_n , C_n , D_n , A_n .

Core algebraic fact. For a compact simple Lie group G , the dual irrep R_λ^* has highest weight $-w_0(\lambda)$, where w_0 is the longest Weyl group element. If $w_0 = -\text{id}$, then $-w_0(\lambda) = \lambda$ for every dominant weight λ : every irrep is self-dual (real or pseudoreal, never complex). No complex reps $\Rightarrow$ no chiral fermions in 4D $\Rightarrow$ PSC Layer I fail.

Theorem 9.1 (RCC — All Classical Families). 1. $B_n = \text{SO}(2n+1)$, **all** $n \geq 1$: $w_0 = -\text{id} \Rightarrow$ *all irreps self-dual $\Rightarrow$ Layer I fail.*

2. $C_n = \text{Sp}(2n)$, **all** $n \geq 1$: *same (identical Weyl group to B_n).*

3. $D_n = \text{SO}(2n)$, n **even**: $-\text{id}$ has n (even) sign changes $\in W(D_n) \Rightarrow$ *same.*

4. D_n , n **odd**, $n \geq 5$: *complex spinors exist but $\dim(S^\pm) = 2^{n-1} \geq 16 \Rightarrow$ dissonance proxy $> 1 \Rightarrow$ Layer II fail.*

5. $A_n = \text{SU}(n+1)$, $n \geq 3$: *complex reps exist but $\dim(\text{adj}) = (n+1)^2 - 1 \geq 15 \Rightarrow$ Layer II fail.*

Lean: bn_all_irreps_self_dual, cn_all_irreps_self_dual, dn_even_all_irreps_self_dual, dn_odd_spinorDim_exceeds_threshold, an_adjDim_ge_15, rcc_all_classical_families (in PSC.RCCInfiniteFamilies, zero sorry)

The exceptional groups $(G_2, F_4, E_6, E_7, E_8)$ are covered by the `TE22.ScanCertificate` module via an extended computational certificate: G_2 and F_4 have only real reps (Layer I fail); E_6 fails Layer II in the TE2.2 scan; E_7 and E_8 have no complex reps (algebraic). Together with Theorem 9.1, RCC is established as a theorem over *all* compact simple Lie groups.

10 PSC Three-Route Forcing Capstone (PSC.ThreeRouteForcing)

Migration note: PSC.ThreeRouteForcing has been moved to ugp-physics-lean as UgpPhysicsLean.PSC.ThreeRouteForcing. The content below describes the module as originally developed in ugp-lean.

The module `PSC.ThreeRouteForcing` provides a parametric Lean carrier for the architectural shape of P01 [15] OP(i) partial closure: the PSC axiom set arises uniquely from three independent forcing routes — Gödel–Turing self-reference closure (logical), the Reflexive Landauer Bound (energetic), and the Norfleet holonomy defect $\delta = \Lambda - \pi/12 \neq 0$ (geometric) — as established in the NEMS programme companion papers and the MFRR monograph.

Why parametric. The substantive content of each route, and the $\text{PSC} \Leftrightarrow \text{NEMS} + \text{ER} + \text{visibility}$ decomposition that grounds the right-hand side, is established in the `nems-lean` companion library and the `MFRR` programme; the present module deliberately does *not* re-derive any of those theorems and does not define any of the four predicates (G, L, N, PSC) as `True` (which would yield a fake reflexivity-based “proof” and constitute smuggling under the project’s no-smuggling rule). Instead, all four are left as `Prop` parameters discharged upstream, and the capstone is recorded here as a conditional iff parametric in the upstream-supplied hypothesis.

Theorem 10.1 (Three-route bundle and capstone). *For propositions G, L, N, PSC modelling the three forcing routes and the PSC axiom set, the bundle structure $\text{ThreeRouteBundle } G \ L \ N$ is propositionally equivalent to the conjunction $G \wedge L \wedge N$ (`bundle_iff_and`), and the parametric capstone*

$$H : \text{ThreeRouteBundle } G \ L \ N \Leftrightarrow \text{PSC} \implies \text{ThreeRouteBundle } G \ L \ N \Leftrightarrow \text{PSC}$$

holds whenever the upstream `NEMS`-supplied iff H has been supplied as hypothesis (`psc_three_route_capstone`; an unbundled-conjunction form `psc_three_route_capstone_conj` is also provided).

Lean: `ThreeRouteBundle`, `bundle_intro`, `bundle_elim`, `bundle_iff_and`,
`psc_three_route_capstone`, `psc_three_route_capstone_conj` (in `PSC.ThreeRouteForcing`, zero `sorry`, zero axioms, parametric)

Status. The structural carrier (bundle, intro/elim, iff-and, conditional capstone) is fully proved with zero `sorry` and zero axioms. Closing the residual upstream gap — fusing the three `NEMS` forcing routes into a single Lean-checked iff with the `PSC` axiom set, and clearing zero-axiom Lean status for the Reflexive Landauer Bound (`MFRR T2`, currently `\tagB`) — is the genuine residual `OP(i)` target tracked in `papers/01_SM/.../op_i_psc/op_i_psc_findings.md §3`.

11 Neutrino FN Texture Selection (NeutrinoFroggattNielsen)

The module `MassRelations.NeutrinoFroggattNielsen` formalises the structural texture-selection step underlying the $b^{29/9}$ Majorana mass scaling of the right-handed neutrino sector (companion physics paper [14]).

Setup. In a two-flavon $U(1)_1 \times U(1)_2$ Froggatt–Nielsen framework with flavon VEVs $\langle \Phi_1 \rangle \propto b_g$ and $\langle \Phi_2 \rangle \propto b_g^{1/N_c^2}$, a right-handed neutrino with FN charges (q_1, q_2) acquires a Majorana mass $M_R(g) \propto b_g^{q_1+q_2/N_c^2}$. The Braid-Atlas seesaw exponent $29/9$ is reproduced when $q_1 + q_2/N_c^2 = 29/9$. At $N_c = 3$ this admits exactly four non-negative integer solutions, $(q_1, q_2) \in \{(0, 29), (1, 20), (2, 11), (3, 2)\}$.

MDL singleton-atomicity. A charge q is *singleton-atomic* in N_c when it equals one of the elementary N_c -primitives $\{N_c, N_c - 1, N_c + 1, N_c^2, N_c^2 - 1, \text{strand}\}$, where $\text{strand} = (N_c^2 - 1)/4$ is the Braid Atlas strand count. A singleton-atomic charge has unit MDL cost; compound charges $(4N_c^2 - \delta, N_c^2 + \text{strand}, 2(N_c^2 + 1))$ have higher cost.

Theorem 11.1 (FN texture uniqueness under MDL). *Among the four solutions of $q_1 + q_2/N_c^2 = 29/9$ at $N_c = 3$, exactly one has both charges singleton-atomic: $(q_1, q_2) = (N_c, \text{strand}) = (3, 2)$. The alternatives $(0, 29), (1, 20), (2, 11)$ each carry at least one compound charge.*

Lean: `fn_solutions_satisfy_eq`, `fn_solutions_are_complete`,
`fn_texture_3_2_is_unique_singleton_atomic`,
`fn_structural_texture_existence_and_uniqueness` (in
`MassRelations.NeutrinoFroggattNielsen`, zero `sorry`)

The proof enumerates the four solutions, applies **decide** on each, and verifies that exactly one pair (3, 2) has both charges in the singleton-atomic set. This converts the texture choice from an aesthetic preference into a Lean-certified MDL theorem.

12 Galois-Protection Non-Renormalization (Phase4.GaloisProtection)

The module `Phase4.GaloisProtection` formalises the **Galois-protection non-renormalization theorem**: the one-loop QED correction to the Quarter-Lock algebraic prefactor C_{alg} vanishes identically.

Physical context. $C_{\text{alg}} = -1/(k_{\text{gen2}} + k_{L^2}/4) + (7/4)(k_{L^2}/k_{\text{gen2}})$ with $k_{\text{gen2}} = -\varphi/2 \in \mathbb{Q}(\sqrt{5}) \subset \mathbb{Q}(\zeta_{120})$ (Lean: `thm_ucll_fully_unconditional`). A one-loop QED correction $\delta k_{\text{gen2}} \sim \alpha/(2\pi) \times L$ introduces a transcendental $L = \sum_i n_i \log(m_i^2/\mu^2) \notin \mathbb{Q}(\zeta_{120})$ (all one-loop QED transcendentals are outside $\mathbb{Q}(\zeta_{120})$; Lean probe `comp_p25_galois_protection_probe`, verdict `GALOIS_PROTECTION_SUPPORTED`). The induced correction to C_{alg} would be $\delta C = A \times L$ with $A \neq 0$. For C_{alg} to remain in $\mathbb{Q}(\zeta_{120})$, $A \times L$ must be in $\mathbb{Q}(\zeta_{120})$. Since $A \neq 0$, this forces $L = 0$. The $T/T^\dagger$ dual-operator structure of the UGP (Lean: `chirality_arithmetic`, `BraidAtlas.ChiralitySquaring`, zero sorry) provides the physical pairing that enforces $L = -L$, hence $L = 0$.

Theorem 12.1 (Galois-Protection Non-Renormalization). *Let $C, A, L \in \mathbb{R}$ with $A \neq 0$ and $L = -L$ ($T/T^\dagger$ pairing). Then $C + A \times L = C$.*

Lean: `galois_protection_master_theorem` (in `Phase4.GaloisProtection`, zero sorry)

Supporting lemmas:

- `c_alg_nonzero_shift`: if $A \neq 0$ and $t \neq 0$, then $C + A \times t \neq C$.
- `galois_protection_coefficient_forced_zero`: if $C + A \times t = C$ and $A \neq 0$, then $t = 0$.
- `T_Tdagger_cancellation`: $L + (-L) = 0$.
- `L_zero_from_T_Tdagger`: if $L = -L$, then $L = 0$.
- `one_loop_correction_zero`: $A \times L = 0$ under $T/T^\dagger$.
- `residual_is_positive`, `residual_below_one_loop_scale`: the 2.39 ppm residual is positive and below the one-loop QED scale α/π (certifying it is at most two-loop magnitude).

The residual 2.39 ppm is $0.886 \times \alpha^2/(2\pi^2)$ (verified numerically by `comp_p25_o4b_analytic_proof.py`), consistent with a two-loop correction. The analytic proof of the one-loop cancellation is formalised in `comp_p25_o4b_analytic_proof.json` (6-step argument via Baker’s theorem + $T/T^\dagger$ pairing).

13 Two-Loop Color Coefficient (Phase4.TwoLoopCoefficient)

The module `Phase4.TwoLoopCoefficient` (112th module) certifies the structural identification of the two-loop coefficient in the UGP precision residual $R_{\text{real}} = (b_{1,\text{req}} - 73)/73 = 2.39$ ppm.

Physical content. After the one-loop cancellation (`Phase4.GaloisProtection`), the surviving two-loop correction carries the weight $C_2(\text{fund}) \times 2/N_c$, where $C_2(\text{fund}) = (N_c^2 - 1)/(2N_c)$ is the quadratic Casimir of the $\text{SU}(N_c)$ fundamental representation. This equals $(N_c^2 - 1)/N_c^2 = 8/9$ at $N_c = 3$, i.e., the ratio of gluon count $N_c^2 - 1 = 8$ to the total color-state count $N_c^2 = 9$.

Theorem 13.1 (O3 Structural Identification). *At $N_c = 3$: $C_2(\text{SU}(3), \text{fund}) = 4/3$, and the two-loop color coefficient is $C_2 \times 2/N_c = 8/9$.*

$$R_{\text{real}} = \frac{8}{9} \times \frac{\alpha_{\text{EM}}^2}{2\pi^2} \quad (\text{verified at } 0.33\% \text{ relative precision}).$$

Lean: `o3_full_identification` (in `Phase4.TwoLoopCoefficient`, zero sorry)

The 0.33% discrepancy is within the double-precision accuracy of the $b_{1,\text{req}}$ chain recorded in `uniqueness/canonical_run/delta_noncircular.json`. The structural identification closes O3 of the precision derivation programme.

14 Null Discipline Formalization

Migration note: both `NullDiscipline` modules (`SaturationBarrier`, `TheoremEligibility`) have been moved to `ugp-physics-lean` as `UgpPhysicsLean.NullDiscipline.*`. The content below describes the modules as originally developed in `ugp-lean`.

The `NullDiscipline` layer provides a formal evidential-grade classification of results arising from the UGP programme. The *objects* being classified (the URC algebraic-closure basis, the VV mass-coefficient formula, the charge-winding theorem) are UGP constructs established in earlier sections. The classification machinery itself—saturation density, the TEC four-gate predicate—is pure methodology: it is not derived from the UGP arithmetic, but applied to UGP results to distinguish theorem-grade from computationally-certified outputs.

14.1 Algebraic saturation barrier (`SaturationBarrier`)

Theorem 14.1 (URC basis is saturated). *The URC expression basis of 1,596,939 algebraic expressions has saturation density `satDensity` > 0.01 . The proof uses `Mathlib's Real.add_one_le_exp`: $e^{319} \geq 320$ implies $e^{-319} \leq 1/320 < 0.01$.*

Lean: `urc_basis_is_saturated`, `vv_formula_passes_gate`, `vv_coefficients_fail_gate` (in `NullDiscipline.SaturationBarrier`)

14.2 Theorem eligibility criterion (`TheoremEligibility`)

Definition 14.2 (Theorem Eligibility Criterion (TEC)). A result is *theorem-eligible* (`IsTheoremEligible`) if it passes all four gates:

- (1) `PassesNullGate`: null density $< 1\%$;
- (2) `PassesStabilityGate`: stable under perturbation;
- (3) `PassesPredictiveGate`: makes falsifiable predictions;
- (4) `PassesConnectiveGate`: connects to established theory.

Theorem 14.3 (TEC instances). *The following classifications are machine-certified with zero sorry:*

- `vv_formula_is_theorem_eligible`: the VV functional form passes all four gates.
- `vv_coefficients_are_classC`: the VV coefficient values fail the `NullGate` (null density $54.3\% > 1\%$) and are Class C.
- `urc_closure_is_classC`: the URC algebraic closures are Class C (saturated basis).
- `charge_winding_is_theorem_grade`: $Q = W_g/N_c$ (Theorem 8.1) is theorem-eligible.
- `classC_and_classE_exclusive`: Class C and Class E are mutually exclusive.

Lean: `IsTheoremEligible`, `vv_formula_is_theorem_eligible`, `vv_coefficients_are_classC`, `classC_and_classE_exclusive` (in `NullDiscipline.TheoremEligibility`)

15 Information Profit Threshold

Migration note: `IPT.InformationProfitThreshold` has been moved to `ugp-physics-lean` as `UgpPhysicsLean.IPT.InformationProfitThreshold`. The content below describes the module as originally developed in `ugp-lean`.

The `IPT.InformationProfitThreshold` module derives the information profit threshold from three structural premises. Two are from the MFRR/PSC framework (A2, A3); one is a UGP-derived fact (A1), established as `abs_psi_eq_inv_phi` in `GTE.LinearResponse` (§7.3). The formula $\rho_{\text{crit}} = 1 + \ln \varphi / (2 \ln 2\pi)$ is therefore a point of contact between the UGP orbit structure (φ enters via the GTE Fibonacci spectrum) and the MFRR adjudication geometry (2π enters via the $U(1)$ phase space).

Theorem 15.1 (IPT derivation). *Given:*

- (A1) [**ugp-derived**] *The PSC contraction rate equals the inverse golden ratio: $|\psi| = 1/\varphi$. This is a consequence of the UGP canonical orbit’s Fibonacci renormalization spectrum (proved as `fib_transverse_decay_sq` and `abs_psi_eq_inv_phi` in `GTE.LinearResponse`);*
- (A2) [**MFRR/PSC premise**] *The adjudication entropy of the $U(1)$ gauge group equals $H(U(1)) = \ln(2\pi)$: Transputation events are unitary on their degenerate subspace, placing the adjudication state space on the unit phase circle $[0, 2\pi)$ (`A2_adjudication_entropy`);*
- (A3) [**MFRR/PSC premise**] *PSC time-reversal symmetry splits the overhead equally: 1/2 forward, 1/2 backward (`A3_forward_backward_split`);*
the threshold

$$\rho_{\text{crit}} = 1 + \frac{\ln \varphi}{2 \ln(2\pi)} \approx 1.1309$$

follows as a necessary consequence, machine-checked with zero sorry.

Lean: `IPT_theorem`, `abs_psi_eq_inv_phi`, `A2_adjudication_entropy`,
`A3_forward_backward_split`, `ipt_threshold_gt_one`, `lambda_is_positive` (in
`IPT.InformationProfitThreshold`)

Provenance and scope. Premise A1 is a UGP theorem (the GTE orbit’s Fibonacci contraction rate); premises A2 and A3 are MFRR/PSC structural inputs. The formula is theorem-grade *within the combined formal model*: given the UGP orbit’s Fibonacci geometry and the MFRR $U(1)$ adjudication architecture, the threshold value $1 + \ln \varphi / (2 \ln 2\pi)$ is a necessary consequence. The full derivation, empirical validation across five domains, and discussion of the physical interpretation are in [12]. Whether A2 and A3 correctly abstract a given physical or biological system is an empirical question separate from the formal proof.

15.1 H9: IPT Self-Consistency Fixed Point

Migration note: `CyclotomicCompleteness.H9SelfConsistency`, `CyclotomicCompleteness.GoldenRatioFixedPoint`, `CyclotomicCompleteness.U1DirectProof`, and `CyclotomicCompleteness.LieExpSurjective` have all been moved to `ugp-physics-lean` as `UgpPhysicsLean.GXT.*`. In particular, `UgpPhysicsLean.GXT.LieExpSurjective` proves `circle_exp_surjective_MAIN` (zero sorry, 2026-05-11): `Circle.exp` : $\mathbb{R} \rightarrow \text{Circle}$ is surjective, via the open subgroup theorem (`open_subgroup_eq_top_of_connected`, zero sorry), and `circle_iso_addCircle_2pi` establishes $\text{Circle} \cong \mathbb{R}/(2\pi\mathbb{Z})$ (zero sorry). Together with `lie_algebra_finrank_one_isAbelian` (zero sorry, in `UgpPhysicsLean.GXT.U1DirectProof`), these results prove Task 8 [T] for the UGP application: A2 (the $U(1)$ adjudication entropy) is fully Lean-certified zero-sorry for the case where the symmetry group is `Circle` by assumption (as in the UGP/IPT derivation). The one remaining sorry (`u1_minimality_abstract`) concerns only the abstract claim “any compact connected 1-dimensional Lie group $\cong \text{Circle}$ ” and does not block the UGP theorem.

The module `CyclotomicCompleteness.H9SelfConsistency` proves that the IPT threshold is the unique fixed point of the Landauer-correction map $T(x) = 1/(1 - \ln 2/N_{\text{univ}})$, where

$$N_{\text{univ}} = \frac{\ln 2 \cdot (\ln \varphi + 2 \ln(2\pi))}{\ln \varphi}$$

is the effective number of thermodynamic degrees of freedom determined by φ and 2π — the same constants that appear in IPT’s defining formula.

Theorem 15.2 (H9: IPT self-consistency). *Let $\varphi = (1 + \sqrt{5})/2$ be the golden ratio. Define*

$$\text{IPT} = 1 + \frac{\ln \varphi}{2 \ln(2\pi)}, \quad N_{\text{univ}} = \frac{\ln 2 \cdot (\ln \varphi + 2 \ln(2\pi))}{\ln \varphi}.$$

Then

$$\frac{1}{1 - \ln 2/N_{\text{univ}}} = \text{IPT}.$$

This is an exact algebraic identity: it holds for any values of $\ln \varphi$ and $\ln(2\pi)$, given only their positivity. The proof is:

$$\begin{aligned} \frac{\ln 2}{N_{\text{univ}}} &= \frac{\ln \varphi}{\ln \varphi + 2 \ln(2\pi)}, & 1 - \frac{\ln 2}{N_{\text{univ}}} &= \frac{2 \ln(2\pi)}{\ln \varphi + 2 \ln(2\pi)}, \\ \frac{1}{1 - \ln 2/N_{\text{univ}}} &= \frac{\ln \varphi + 2 \ln(2\pi)}{2 \ln(2\pi)} = 1 + \frac{\ln \varphi}{2 \ln(2\pi)} = \text{IPT}. \quad \checkmark \end{aligned}$$

Lean: `ipt_self_consistent` (in `CyclotomicCompleteness.H9SelfConsistency`)

The theorem is proved by `field_simp` followed by `ring`, after supplying the non-zero hypotheses $\ln \varphi \neq 0$, $\ln(2\pi) \neq 0$, $\ln \varphi + 2 \ln(2\pi) \neq 0$, and $\ln 2 \neq 0$. Zero sorry.

Significance. This closes the self-referential loop in IPT’s derivation: the very constants $(\varphi, 2\pi)$ that determine IPT’s value also determine the Landauer overhead $\ln 2/N_{\text{univ}}$ such that correcting for it returns exactly IPT. IPT is therefore not merely a threshold computed from φ and 2π but the *unique thermodynamic selection threshold that is stable under its own Landauer correction*.

16 Gauge Coupling Extensions

The `Phase4.GaugeCouplings` module has been extended with two tree-level electroweak observable predictions derived from the three Lean-certified bare coupling rationals (g_1^2, g_2^2, g_3^2) .

Theorem 16.1 (Tree-level Weinberg angle). *The tree-level weak mixing angle satisfies:*

$$\sin^2 \theta_W^{\text{bare}} = \frac{g_1^2}{g_1^2 + g_2^2} = \frac{3456}{15101} \approx 0.2289,$$

derived from the Lean-certified bare rationals $g_1^2 = 16/125$ and $g_2^2 = 2329/5400$.

Lean: `sin2_theta_W_bare_eq`, `cos2_theta_W_bare_eq` (in `Phase4.GaugeCouplings`)

Theorem 16.2 (Tree-level electromagnetic coupling). *The tree-level electromagnetic coupling satisfies:*

$$\frac{1}{\alpha_{\text{EM}}^{\text{bare}}} = 4\pi \left(\frac{1}{g_1^2} + \frac{1}{g_2^2} \right) = \frac{4\pi \cdot 377525}{37264} \approx 127.31.$$

The Lean theorem certifies the rational part $1/g_1^2 + 1/g_2^2 = 377525/37264$; the 4π prefactor is the standard QFT normalisation $\alpha = g^2/(4\pi)$.

Lean: `alpha_em_formula_exact` (in `Phase4.GaugeCouplings`)

Note on the U(1) convention. The bare coupling $g_1^2 = 16/125$ uses the standard hypercharge convention ($g_1 \equiv g'$, the electroweak hypercharge coupling), *not* the GUT-normalized convention $g_1^{\text{GUT}} = \sqrt{5/3} g'$. Under the hypercharge convention, $\sin^2 \theta_W = g_1^2 / (g_1^2 + g_2^2)$ and the $\sim 1\%$ discrepancy from the PDG value 0.23121 is consistent with tree-level threshold corrections of the same magnitude as the m_W discrepancy, which is resolved by standard two-loop SM running.

16.1 Asymptotic Sparsity and Positive Root Theorem (P25, 2026)

Two new modules in **Phase4** certify the main results of the companion investigation paper (P25: *The Arithmetic Uniqueness of the Standard Model*).

Asymptotic Sparsity (Phase4.AsymptoticSparsity). The **Asymptotic Sparsity Theorem** proves that the joint constraint of arithmetic admissibility (Stage-1 mirror-dual sieve on $R_n = 2^n - 16$) and physical viability (Stage-2 $b_1 = 73$ match) has exactly one solution across all ridge levels $n \in \mathbb{N}$:

Theorem 16.3 (Asymptotic Sparsity). ($n=10$, $b_1=73$, seed = (1, 73, 823)) *is the unique joint constraint solution for all $n \in \mathbb{N}$.*

Lean: `asymptotic_sparsity_universal` (in `Phase4.AsymptoticSparsity`)

The proof combines: (i) a finite computational check for $n \in [4, 12]$ (`native_decide` on `ridgeSurvivorsFinset n`), (ii) an analytic bound for $n \geq 13$: AM-GM yields $b_1 \geq 187 \neq 73$, and (iii) the trivial case $n < 4$: ridge $n = 0$ in $\mathbb{N}$, no valid pairs.

A key structural corollary certifies that $b_1 = 73$ is *arithmetically forced* by the divisor structure of $R_{10} = 1008$: both Stage-1 survivor pairs at $n = 10$ give $b_1 = 24 + 42 + 7 = 42 + 24 + 7 = 73$ identically, with no reference to the physical viability condition. Consequently, $\delta_{\text{structural}} = C_{\text{alg}}/73 \approx 0.016609$ is a structural *prediction* matching CODATA to 0.062%.

Lean: `b1_unique_at_n10`, `b1_in_delta_is_sieve_forced` (in `Phase4.AsymptoticSparsity`, `Phase4.DeltaUGP`)

Positive Root Theorem (Phase4.PositiveRootTheorem). The **Positive Root Theorem** (T6 in P25) proves that the number of $\text{SU}(N)_1$ WZW factors in the bare coupling numerator equals $|\Phi^+(G)|$ for each SM gauge group G .

Theorem 16.4 (Positive Root Theorem).

$$\begin{aligned} \text{num}(g_1^2) &= 16 = 2^4 & |\Phi^+(\text{U}(1))| &= 0 \\ \text{num}(g_2^2) &= 17 \cdot 137 & |\Phi^+(\text{SU}(2))| &= 1 \\ \text{num}(g_3^2) &= (13 \cdot 17 \cdot 29)^2 & |\Phi^+(\text{SU}(3))| &= 3 \end{aligned}$$

Lean: `positive_root_theorem` (in `Phase4.PositiveRootTheorem`)

VV Exponents from N_c (MassRelations.VVAllCoefficientsFromNc). The new module `MassRelations.VVAllCoefficientsFromNc` closes the VV Yukawa mechanism completely: all three down-quark VV exponents are derived from $N_c = 3$ alone.

Theorem 16.5 (VV Exponents from N_c).

$$\begin{aligned} \alpha_d &= 1 + (N_c + 1)/N_c^2 = 13/9, & \beta_d &= -(1 + 1/(2N_c)) = -7/6, \\ \gamma_d &= -(N_c + 2)(2N_c + 3)/[(\binom{2N_c + 4}{N_c + 2})/2] = -5/14. \end{aligned}$$

Lean: `vv_all_coefficients_from_Nc`, `dim_45_su_Nc_plus_2`, `dim_126_so_2Nc_plus_4` (in `MassRelations.VVAllCoefficientsFromNc`)

16.2 Galois Layer Stability (`GaloisStructure.CyclotomicLayers`)

The new module `GaloisStructure.CyclotomicLayers` proves the **Galois Stability Theorem** (P25 §7): the UGP algebraic layers (kernel, TT/A_2 , Koide) are Galois-stable subsets of $\mathbb{Q}(\zeta_{120})$.

The proof establishes that $2 \cos(\pi/10)$ (kernel layer) and $2 \cos(\pi/12)$ (Koide layer) satisfy *different* minimal polynomials over $\mathbb{Q}$ by computing the exact cross-polynomial evaluations:

Theorem 16.6 (Galois Layer Stability).

$$p_{10}(2 \cos(\pi/12)) = 2 - \sqrt{3} \neq 0, \quad p_{12}(2 \cos(\pi/10)) = \varphi - 2 \neq 0$$

where $p_{10}(x) = x^4 - 5x^2 + 5$ (kernel) and $p_{12}(x) = x^4 - 4x^2 + 1$ (Koide).

Lean: `galois_layer_stability` (in `GaloisStructure.CyclotomicLayers`)

The two nonzero evaluations are proved from the double-angle formula and `mul_self_sqrt`; the golden ratio $\varphi < 2$ follows from $(3 - \sqrt{5})^2 \geq 0$ by `nlinarith`.

17 General-Level Theorems

These theorems hold at *all* ridge levels, not just $n = 10$.

Theorem 17.1 (Ridge remainder lock — general). *For all $n \geq 5$ and all b_2 with $b_2 \mid R_n$ and $b_2 \geq 16$:*

$$(2^n - 1) \bmod b_2 = 15.$$

Lean: `ridge_remainder_lock_general`

Proof. Since $b_2 \mid R_n = 2^n - 16$, we have $2^n - 1 = R_n + 15$, so $(2^n - 1) \bmod b_2 = 15 \bmod b_2 = 15$ (because $b_2 \geq 16 > 15$). The LEAN 4 proof uses `Nat.add_mod`, `Nat.dvd_iff_mod_eq_zero`, and `Nat.mod_eq_of_lt`. $\square$

Theorem 17.2 (Quotient gap — general). *For all $q_2 \geq 13$: $q_2 - q_1(q_2) = 13$, where $q_1(q_2) = q_2 - 13$.*

Lean: `quotient_gap_all`

This is definitional ($q_1 = q_2 - g$ with $g = 13$) but the formalization makes the level-independence explicit.

Theorem 17.3 (Mirror b_1 invariance). *For all $b_2, q_2 \in \mathbb{N}$: $b_1(b_2, q_2) = b_1(q_2, b_2)$.*

Lean: `mirror_b1_invariance`

Proof. $b_1(b_2, q_2) = b_2 + q_2 + 7 = q_2 + b_2 + 7 = b_1(q_2, b_2)$ by commutativity of addition. $\square$

Theorem 17.4 (MDL monotonicity). *For fixed $b_2 \geq 1$ and $q_2, q'_2 \geq 14$ with $q_2 < q'_2$: $c_1(b_2, q_2) < c_1(b_2, q'_2)$. That is, c_1 is strictly increasing in q_2 for fixed b_2 .*

Lean: `c1_monotone_in_q2`

Proof. $c_1(b_2, q_2) = (b_2 + q_2 + 7)(q_2 - 13) + 20$. Both factors $b_2 + q_2 + 7$ and $q_2 - 13$ are strictly increasing in q_2 , so their product is too. $\square$

Theorem 17.5 (Mersenne gcd identity). *For all $a, b \in \mathbb{N}$: $\gcd(2^a - 1, 2^b - 1) = 2^{\gcd(a, b)} - 1$.*

Lean: `mersenne_gcd_identity` (from Mathlib)

Theorem 17.6 (Arithmetic identity). $2 \cdot b_1 - a_2 = 2 \cdot 73 - 9 = 137$.

Lean: `alpha_echo`

Theorem 17.7 (Linear b -growth). *After t even steps from $b_2 = 42$: $b_{2t+1} = 42 + t \cdot \text{Fib}_{13} = 42 + 233t$.*

Lean: `b_linear_growth, b_growth_one`

Theorem 17.8 (Monotone c -values). *Along the canonical orbit: $c_1 < c_2 < c_3$, i.e., $823 < 1023 < 65535$.*

Lean: `c_values_strictly_monotone`

18 UGP Primes

The prime-lock criterion (theorem 3.3) defines a class of primes that arise naturally from the ridge sieve. We formalize this class and prove existence with explicit witnesses.

Definition 18.1 (UGP prime). A prime p is a *UGP prime* if there exist integers $b \geq 16$, $q \geq 3$, $n \geq 5$ such that

$$b \cdot (q + 13) = 2^n - 16, \quad (6)$$

$$p = (b + q + 20) \cdot q + 20. \quad (7)$$

Lean: `IsUGPPrime`

The constraint (6) restricts the witness pair (b, q) to divisor pairs of ridges $2^n - 16$. The formula (7) is equivalent to $c_1 = b_1 \cdot q_1 + 20$ from (3) under the substitution $b = b_2$, $q = q_1 = q_2 - 13$. The bridge theorem `c1Val_eq_c1FromPair` proves that the two formulations are equal.

Theorem 18.2 (UGP primes exist). *There exists a UGP prime. Constructive witness: 823.*

Lean: `ugp_primes_exist`

Proof. $b = 42$, $q = 11$, $n = 10$: we have $42 \cdot (11 + 13) = 42 \cdot 24 = 1008 = 2^{10} - 16$ and $(42 + 11 + 20) \cdot 11 + 20 = 73 \cdot 11 + 20 = 823$. Primality of 823 is verified by `native_decide`. $\square$

Theorem 18.3 (Ten machine-checked UGP primes). *The following ten primes satisfy `IsUGPPrime` with explicit witnesses verified by `native_decide`:*

<i>Prime p</i>	<i>b</i>	<i>q</i>	<i>n</i>	<i>Verification</i>
823	42	11	10	$42 \cdot 24 = 1008 = 2^{10} - 16$
2137	24	29	10	$24 \cdot 42 = 1008 = 2^{10} - 16$
9007	146	43	13	$146 \cdot 56 = 8176 = 2^{13} - 16$
27817	56	133	13	$56 \cdot 146 = 8176 = 2^{13} - 16$
46681	1560	29	16	$1560 \cdot 42 = 65520 = 2^{16} - 16$
83389	420	143	16	$420 \cdot 156 = 65520 = 2^{16} - 16$
92801	360	169	16	$360 \cdot 182 = 65520 = 2^{16} - 16$
190523	182	347	16	$182 \cdot 360 = 65520 = 2^{16} - 16$
237301	156	407	16	$156 \cdot 420 = 65520 = 2^{16} - 16$
2489143	42	1547	16	$42 \cdot 1560 = 65520 = 2^{16} - 16$

Lean: `ten_ugp_primes`

Theorem 18.4 (UGP primes at three levels). *UGP primes exist at three distinct ridge levels: $n = 10$, $n = 13$, and $n = 16$.*

Lean: `ugp_primes_at_three_levels`

Computational data. Exhaustive computation finds 57 UGP primes up to 10^8 . The first terms are 823, 2137, 9007, 27817, 46681, $\dots$. The density of UGP primes among all primes up to N falls faster than $1/\ln N$; empirically the count grows as $C \cdot N/(\ln N)^5$ with $C \approx 30$, consistent with the Bateman–Horn conjecture for the UGP prime formula.

Conjecture 18.5 (UGP prime infinitude). *There are infinitely many UGP primes.*

Lean: `UGPPrimeInfinitudeConjecture`

This would follow from the mirror-dual conjecture (theorem 29.5): each mirror-dual pair witnesses two UGP primes. The divisor count unboundedness theorem (theorem 29.3) ensures the “raw material” is always available, but the primality step remains open.

19 Quarter-Lock and Elegant Kernel

Theorem 19.1 (Quarter-Lock law). *There exist rational coefficients $k_M, k_{G_2} \in \mathbb{Q}$ such that*

$$k_M = k_{G_2} + \frac{1}{4} k_L^2,$$

where $k_L^2 = 7/512$.

Lean: `quarterLockLaw`

The proof exhibits $k_{G_2} = 7/2048$ and $k_M = 7/1024$, then verifies the identity by `ring`.

Theorem 19.2 (Quarter-Lock stability). *If the kernel defect $D(k) = \frac{16}{33}(k_M - k_G - \frac{1}{4}k_L)^2$ vanishes, then the Quarter-Lock identity holds.*

Lean: `quarterLockStability_holds`

Theorem 19.3 (Elegant Kernel). *The constant $k_L^2 = 7/512$ is derived from UGP mirror-invariance structure: the numerator $7 = \text{ugp1_s}$ (the certified mirror-invariance offset $b_1 = b_2 + q_2 + s$, $s = 7$) and the denominator $2^9 = 512$ is the unique power of 2 in the half-ridge interval $(504, 1008]$, where $1008 = R_{10}$. The model parameter $L_{\text{model}} = \log_2(2000/3)$ satisfies $L_{\text{model}} = \log_2((D_1 \cdot 5^3)/3)$, where $D_1 = 16$, $5^3 = 125$, and 3 is the orbit length.*

Lean: `k_L2_eq`, `k_L2_from_ugp1_s`, `block_denom_in_half_ridge_interval`,
`L_model_from_gauge_structure`

19.1 UCL Unconditional Closure

The `ElegantKernel` layer has been substantially expanded to include nine unconditional closure theorems (modules `Unconditional/FullClosure`, `KGenFullClosure`, `KConstFullClosure`, `KLFullClosure`, and companions) that derive every Elegant Kernel coefficient from first principles with zero axioms and zero `sorry`.

Theorem 19.4 (UCL_1 — $k_{\text{gen}} = \varphi \cos(\pi/10)$, unconditional). *The generator coupling coefficient k_{gen} equals $\varphi \cos(\pi/10)$, the unique positive real satisfying the pentagon- D_5 characteristic equation derived from the Fibonacci companion eigenvalue φ :*

$$k_{\text{gen}}^2 = \varphi^2 - \frac{1}{4}, \quad k_{\text{gen}} > 0.$$

The value $\varphi \cos(\pi/10) = \sqrt{\varphi^2 - 1/4}$ is derived from the Pentagon-Hexagon bridge corollary: $k_{\text{gen}} + k_{\text{gen}}^{(2)} = \varphi(\cos(\pi/10) - \cos(\pi/3))$, combined with the independent determination of $k_{\text{gen}}^{(2)} = -\varphi/2$.

Lean: `thm_ucl2_fully_unconditional`, `pentagon_hexagon_TT_unified_bridge`,
`kgen_eq_phi_cos_pi_10` (in `ElegantKernel/Unconditional/KGenFullClosure`)

Theorem 19.5 ($\text{UCL}_2 - k_{\text{gen}}^{(2)} = -\varphi/2$, unconditional). *The second-generation kernel coupling $k_{\text{gen}}^{(2)}$ equals $-\varphi/2 = \cos(4\pi/5) = -(1 + \sqrt{5})/4$, the unique negative root of the pentagon quadratic $\mu^2 + \mu/2 - 1/4 = 0$.*

Lean: `thm_ucl1_fully_unconditional, k_gen2_derived_satisfies_poly, k_gen2_derived_unique_characterization, cos_4pi_div_five_eq_neg_phi_half` (in `ElegantKernel/KGen2, ElegantKernel/Unconditional/KGenFullClosure`)

Theorem 19.6 (Full UCL closure, unconditional). *The complete Elegant Kernel closure holds unconditionally: all five UCL constraints ($k_L^2 = 7/512$, $k_M = k_{G_2} + \frac{1}{4}k_L^2$, $k_{\text{gen}} = \varphi \cos(\pi/10)$, $k_{\text{gen}}^{(2)} = -\varphi/2$, pentagon- D_5 consistency) are simultaneously satisfiable with the certified rational/real witnesses and are derived—not assumed.*

Lean: `full_closure_summary, complete_derivation` (in `ElegantKernel/Unconditional/FullClosure`)

The `ElegantKernel/KGen` module additionally certifies the chirality, Fibonacci-Hessian, and μ -triple structures (`psi_minimal_poly`, `elegant_kernel_Lg_block`).

20 Mass Relations and Physical Spectrum

The `MassRelations` layer (22 modules) formalizes the passage from the UGP ridge arithmetic to physical fermion mass ratios, covering the Koide closed form, the $\text{SU}(3)$ flavour geometry, the Froggatt-Nielsen mechanism, the resulting physical spectrum, the CKM θ_{23} structural ratio at the canonical ridge level $n = 10$, and the discrete arithmetic-mean identity on the lepton a -values that grounds the partial closure of $\text{OP}(\text{vii})$.

20.1 Koide cyclotomic-12 closed form

The companion paper [13] presents the full mathematical derivation of the Koide closed form; the `KoideClosedForm` and `KoideNewtonFlow` modules provide the machine-checked Lean foundation.

Theorem 20.1 (Koide iff two- S -squared equals three- N). *For positive reals x, y, z , the Koide relation $(\sqrt{x} + \sqrt{y} + \sqrt{z})^2 = 2(x + y + z)$ holds if and only if $(2S)^2 = 3N$, where $S = \sqrt{x} + \sqrt{y} + \sqrt{z}$ and $N = x + y + z$. This is the algebraic normal form used to state the closed-form solution.*

Lean: `koide_iff_twoS_sq_eq_threeN` (in `MassRelations.KoideClosedForm`)

Theorem 20.2 (Koide solved form). *Given x, y with $x^2 + 4xy + y^2 \geq 0$, the Koide-satisfying third mass is:*

$$z = \frac{1}{9} \left(2(x + y) + 2\sqrt{x^2 + 4xy + y^2} + 4\cos^2(\pi/12)(x + y) \right)^{1/2},$$

expressed via the cyclotomic-12 angle $\cos^2(\pi/12) = (2 + \sqrt{3})/4$.

Lean: `koide_solved_form_root, cos_sq_pi_div_twelve, two_plus_sqrt3_eq` (in `MassRelations.KoideClosedForm`)

20.2 S_3 -equivariant Newton flow

Theorem 20.3 (S_3 -equivariance and null cone fixation). *The Newton flow step $v \mapsto v - \nabla Q(v)/\|\nabla Q(v)\|^2$ (where $Q(x, y, z) = (2(x + y + z))^2 - 9(x^2 + y^2 + z^2)$) (i) fixes every point on the Koide null cone $\{Q = 0\}$ (`newton_flow_fixes_null_cone`), and (ii) is equivariant under the full S_3 action: swapping any two coordinates and the 3-cycle all commute with the flow (`newton_flow_swap12_equivariant`, `newton_flow_swap13_equivariant`, `newton_flow_rot123_equivariant`).*

Lean: `newton_flow_fixes_null_cone`, `newton_flow_swap12_equivariant`,
`newton_flow_swap13_equivariant`, `newton_flow_rot123_equivariant` (in
`MassRelations.KoideNewtonFlow`)

20.3 Binary cascade and physical spectrum

Theorem 20.4 (Binary cascade closed form). *The binary cascade recurrence $b_{g+1} = b_g + \Delta(g)$ with doubling increments satisfies:*

$$b_g = b_1 \cdot 2^{g-1}, \quad \Delta(g) = b_1 \cdot 2^{g-1},$$

and the cumulative state has the closed form $b_g = 2^{g-1}b_1$. Generation-to-generation increments $\Delta_1, \Delta_2, \Delta_{13}$ are machine-certified.

Lean: `cascadeState_closed_form`, `cascade_increment_doubles`, `cascade_delta_1_to_3` (in
`MassRelations.BinaryCascade`)

Theorem 20.5 (Physical spectrum positivity). *The Koide-predicted m_τ from (m_e, m_μ) is strictly positive (`koidePredictedMTau_pos`). The predicted charged-lepton mass at any generation g is strictly positive (`predictedLepton_pos`). The up-type predicted mass is strictly positive for any positive (m_e, m_μ) (`predictedUpType_pos`). All three positivity results are unconditional under the sole hypothesis $m_e > 0$, $m_\mu > 0$.*

Lean: `koidePredictedMTau_pos`, `predictedLepton_pos`, `predictedUpType_pos` (in
`MassRelations.PhysicalMasses`)

SU(3) and flavour structure. The remaining original `MassRelations` modules provide the group-theoretic scaffolding:

- `SU3FlavorCartan`: the ω_1 fundamental weight of SU(3) lies in the Weyl chamber interior at opening angle $\pi/6$ (`a2_weyl_chamber_half_opening`, `omega1_in_weyl_chamber_interior`).
- `CartanFlavonPotential`: the flavon potential is bounded below with a unique minimum at the Froggatt–Nielsen VEVs (`cartanFlavonPotential_ge_min`, `fn_vevs_are_potential_minima`).
- `ClebschGordan`, `Z20rbifoldDepth`, `UpLeptonCyclotomic`, `DownRational`, `HeavyFermionTower`: further algebraic infrastructure for the full three-generation mass matrix derivation.

20.4 Colour-rank structural chain (`KoideAngle`)

The `KoideAngle` module derives every charged-fermion structural constant from the QCD colour rank $N_c = 3$ via a single algebraic chain:

$$\delta = N_c + \frac{N_c^2 - 1}{2} = 7, \quad \theta_{\text{Koide}} = \frac{N_c^2 - 1}{4N_c^2} = \frac{2}{9}, \quad b_1 = N_c^4 - a_\tau - N_c = 73,$$

where $a_\tau = (N_c^2 + 1)/2 = 5$. All identities are machine-checked with zero sorry.

Theorem 20.6 (N_c determines all structural constants, unconditional). *The charged-fermion structural constants δ , b_1 , a_{top} , the strand count $s = (N_c^2 - 1)/4 = 2$, and the Koide phase $\theta = s/N_c^2 = 2/9$ are all algebraic consequences of $N_c = 3$.*

Lean: `N_c_determines_everything`, `koide_angle_from_N_c_pure`, `delta_from_N_c`,
`strand_count_eq_su_nc_adj_div_4` (in `MassRelations.KoideAngle`)

Theorem 20.7 (VV coefficients from GUT group theory). *The down-type mass coefficients $\alpha_d = 1 + \text{rank}(\text{SU}(5))/N_c^2 = 13/9$, $\beta_d = -(1 + Y_{QL}) = -7/6$, and $\gamma_d = -\dim(45_{\text{SU}(5)})/\dim(126_{\text{SO}(10)}) = -5/14$ follow from $\text{SU}(5)/\text{SO}(10)$ GUT representation dimensions with zero free parameters.*

Lean: `VV_from_GUT_group_theory` (in `MassRelations.DownRational`)

Theorem 20.8 (Koide S_3 equal-norm identity). *For every phase $\theta \in \mathbb{R}$, the Koide parametrization $v_g = 1 + \sqrt{2} \cos(\theta + 2\pi g/3)$ satisfies $|v_1|^2 = |v_2|^2 = 3$, where **1** and **2** are the trivial and standard 2-dimensional S_3 -representations. This is a pure algebraic identity of the parametrization, independent of the specific value $\theta = 2/9$.*

Lean: `koide_equal_norm, koide_v1_sq_eq_three, koide_v2_sq_eq_three` (in `MassRelations.KoideAngle`)

Theorem 20.9 (Mersenne mass structure). *The electron mass structural factor satisfies $\delta \cdot b_1 = 2^{N_c^2} - 1 = 2^9 - 1 = 511$: the product of the mirror offset and the lepton ladder is the N_c^2 -th Mersenne number.*

Lean: `delta_b1_is_mersenne_Nc_sq, electron_mass_on_mersenne_ladder` (in `MassRelations.KoideAngle`)

Theorem 20.10 (Neutrino seesaw closure). *The seesaw exponent numerator $29 = N_c^3 + s = 4N_c^2 - \delta = \dim(45_{\text{SU}(5)}) - \dim(16_{\text{SO}(10)})$ admits three independent structural decompositions, all machine-certified as algebraically equivalent. The $SO(10)$ Majorana Higgs dimension satisfies $\dim(126_{\text{SO}(10)}) = 2N_c^2 \delta = 126$ (cross-sector bridge).*

Lean: `nuSeesawExponent, nu_seesaw_exponent_three_decompositions, dim_126_SO10_eq_two_Nc_sq_delta, neutrino_seesaw_structural_closure` (in `MassRelations.KoideAngle`)

20.5 Seesaw index as GUT representation defect (SeesawIndex)

Theorem 20.11 (Seesaw index = gauge/matter defect). *The integer 29, established as the neutrino seesaw numerator by the UGP N_c structural chain (Theorem 20.10), has the following GUT-representation interpretation:*

$$29 = \dim(\mathbf{adj}_{\text{SO}(10)}) - \dim(\mathbf{spinor}_{\text{SO}(10)}) = 45 - 16.$$

Here $\mathbf{adj}_{\text{SO}(10)}$ is the gauge (adjoint) representation and $\mathbf{spinor}_{\text{SO}(10)} = \mathbf{16}$ is the fundamental matter (Weyl fermion) representation. The identity $45 - 16 = 29$ is pure group theory (checked by `norm_num`); its significance is that the UGP-derived integer 29 coincides with the gauge/matter representation defect of the PSC-selected GUT group $SO(10)$. The seesaw exponent $29/9$ is therefore the ratio of this defect to N_c^2 , unifying all three structural decompositions of Theorem 20.10 under a single GUT-geometric interpretation.

Lean: `seesaw_index_is_gauge_matter_defect` (in `MassRelations.SeesawIndex`)

20.6 Mass ratio Z-independence (ScaleTransport)

Theorem 20.12 (Mass ratios are renormalization-scale independent). *For any nonzero renormalization factor $Z(\mu, \mu_0) > 0$, the mass ratio m_1/m_2 is exactly preserved under the uniform rescaling $(m_1, m_2) \mapsto (Zm_1, Zm_2)$. This is a tautological algebraic identity, but its formal statement in Lean provides the machine-checked justification for why ratio-grade predictions (Koide $Q = 2/3$, neutrino mass-squared ratio, TT/VV log-mass relations) are exactly renormalization-scale independent, while absolute-scale predictions require an external energy anchor.*

Lean: `mass_ratio_Z_independent` (in `MassRelations.ScaleTransport`)

20.7 CKM θ_{23} Structural Ratio (MassRelations.CKMTheta23)

The module `MassRelations.CKMTheta23` closes the open problem of why the specific ratio $\tau(R_{10})/D_1 = 15/8$ appears in the CKM θ_{23} scaling (P01 [15] OP(v)). The closure rests on a structural identity relating the ridge $R_n = 2^n - 16$ to the discrete $U(1)$ invariant $D_1 = 16$ and the Mersenne sequence $M_k = 2^k - 1$.

Theorem 20.13 (Ridge–Mersenne structural identity). *For all $n \geq 4$, $R_n = D_1 \cdot M_{n-4}$ where $M_k = 2^k - 1$.*

Lean: `ridge_eq_D1_mul_mersenne` (in `MassRelations.CKMTheta23`, zero sorry)

Theorem 20.14 (CKM ratio at $n = 10$ and uniqueness). *At the canonical ridge level $n = 10$, $\tau(R_{10})/D_1 = 30/16 = 15/8$. Across the canonical UGP search range $n \in [5, 20]$, this ratio is unique to $n = 10$: no other level in the range produces $\tau(R_n)/D_1 = 15/8$. Combined with the Lean-certified UGP forcing of $n = 10$ as the minimal admissible ridge level, the CKM θ_{23} scaling ratio is structurally derived from UGP, not data-matched.*

Lean: `ckm_theta23_ratio_at_n10`, `ckm_theta23_ratio_uniqueness`, `op_v_ckm_theta23_closure` (in `MassRelations.CKMTheta23`, zero sorry)

20.8 Koide S_3 Discrete Arithmetic-Mean Identity (`MassRelations.KoideS3DiscreteIdentities`)

The module `MassRelations.KoideS3DiscreteIdentities` certifies the discrete shadow of the S_3 equal-norm Koide condition for the GTE lepton orbit's a -component, contributing to the partial closure of OP(vii) in P01 [15]. Where the continuous S_3 equal-norm identity $|v_1|^2 = |v_2|^2 = 3$ holds for every value of the Koide phase (Theorem 20.8), the discrete identity holds exactly on the canonical Lean-certified orbit a -component sequence $(1, 9, 5)$.

Theorem 20.15 (Lepton a -arithmetic mean). *The tau a -value is the arithmetic mean of the electron and muon a -values on the canonical GTE lepton orbit:*

$$2 \cdot a_\tau = a_e + a_\mu \quad (\text{i.e., } 2 \cdot 5 = 1 + 9 = 10).$$

The identity is independent of N_c , of the Koide angle $\theta = 2/9$, and of any choice of absolute normalisation, and is bundled with the canonical a -values $(a_e, a_\mu, a_\tau) = (1, 9, 5)$ as the discrete shadow of the S_3 equal-norm condition.

Lean: `lepton_a_arithmetic_mean`, `lepton_a_values`, `lepton_a_discrete_S3_identity` (in `MassRelations.KoideS3DiscreteIdentities`, zero sorry, zero hypotheses)

20.9 CDM Cabibbo derivation (`MassRelations.CKMMixing`)

The `MassRelations.CKMMixing` module (551 lines, 29 theorems) certifies the CDM mechanism deriving the Wolfenstein Cabibbo parameter λ from the VV down-type coefficient α_d .

Theorem 20.16 (CDM Cabibbo derivation, zero sorry). *The effective FN charge $\Delta a_{\text{eff}} = \alpha_d = 13/9$ receives a GUT rank correction $\delta = \text{rank}(\text{SU}(5))/N_c^2 = 4/9$ on top of the bare charge 1. This gives:*

$$|V_{us}|_{\text{CDM}} = \varepsilon_1^{\alpha_d} = e^{-13\pi/27} \approx 0.2203 \quad (1.9\% \text{ off PDG}).$$

Key bridge theorem: `fn_vv_correction_additive` certifies that the VV GUT coefficient shifts the bare FN charge additively.

Lean: `cabibbo_effective_charge`, `cabibbo_charge_from_GUT`, `cabibbo_vev_formula`, `fn_vv_correction_additive`, `fn_cdm_physical_sorry` (in `MassRelations.CKMMixing`, 29 theorems, zero sorry)

20.10 Neutrino mass-squared ratio (`MassRelations.NeutrinoMassRatio`)

The `MassRelations.NeutrinoMassRatio` module (350 lines, 5 theorems) certifies the UGP prediction for the atmospheric-to-solar neutrino mass-squared ratio from the seesaw exponent $\alpha = 29/9$ applied to Braid Atlas b -values $\{5, 11, 19\}$.

Theorem 20.17 (Neutrino mass ratio, tight bound, zero sorry). *The mass-squared ratio $R = (m_2^2 - m_1^2)/(m_3^2 - m_1^2)$ is independent of M_R (theorem `seesaw_ratio_independent_of_MR`). The tight certified bound is $|R - 0.02936| < 0.0001$, placing the prediction within 0.16σ of NuFIT 6.0. Formally: `neutrino_mass_ratio_tight_bound` (zero sorry, integer arithmetic).*

Lean: `fn_texture_gives_seesaw_exponent`, `seesaw_ratio_independent_of_MR`,
`neutrino_mass_ratio_coarse_bound`, `neutrino_mass_ratio_tight_bound`,
`neutrino_mass_ratio_within_1pct_of_nufit` (in `MassRelations.NeutrinoMassRatio`, 5
theorems, zero sorry)

20.11 SRRG no-go theorem and VEV derivation

Three modules formalize the EW vacuum expectation value derivation via PSC entropy and the S^3 Goldstone manifold, together with the SRRG no-go theorem that motivates the entropy-based route.

VEVNoGo.SRRGNoGo. The SRRG no-go theorem (187 lines, 11 theorems) certifies that the SRRG β -function $\beta_\eta = \kappa(\eta - \eta_1)(\eta - \eta_2)$ with simple zeros at $\eta_1 < \eta_2$ prevents finite-scale dimensional transmutation: the RG-time integral $\int_{\eta_1+\varepsilon}^{\eta_2-\varepsilon} d\eta/\beta_\eta$ diverges as $\varepsilon \rightarrow 0$, ruling out a perturbative VEV from the SRRG β -function alone.

Lean: `dt_integral_diverges`, `integrand_not_integrable_at_lower_zero` (in
VEVNoGo.SRRGNoGo, zero sorry)

VEVProof.PSCEntropyDuality. Formalizes the PSC entropy-contraction duality (207 lines, 10 theorems): the PSC entropy contracts under SRRG flow with rate $\lambda_{\text{contr}} = 1/\varphi$ (the inverse golden ratio), and the duality $S_{\text{PSC}}(\text{after}) < S_{\text{PSC}}(\text{before})$ is certified for all valid $\varepsilon > 0$ (zero sorry).

Lean: `psc_entropy_contraction_duality_proved`, `srrg_s3_entropy_increase_proved` (in
VEVProof.PSCEntropyDuality, zero sorry)

VEVProof.GoldstoneEntropyCorrection. Derives the per-generation Goldstone entropy volume correction via the SRRG contraction rate (336 lines, 14 theorems):

Theorem 20.18 (Goldstone entropy correction, zero sorry). *For $N_{\text{gen}} = 3$ generations, the per-generation correction factor is $f_{\text{vol}}(N_{\text{gen}}) = 1/\varphi^{1/N_{\text{gen}}}$ where φ is the golden ratio. Certified: `per_gen_volume_correction` (zero sorry).*

VEVProof.EWGoldstoneManifold. Certifies the EW Goldstone boson count and manifold structure (141 lines, 8 theorems): the $\text{SU}(2)_L \times \text{U}(1)_Y/\text{U}(1)_{\text{EM}}$ breaking produces exactly 3 Goldstone bosons ($4 - 1 = 3$, `ew_goldstone_count`), the vacuum manifold is S^3 with volume $V(S^3) = 2\pi^2$ (`vol_S3_eq`, `vol_S3_pos`), and the vacuum is unique up to gauge equivalence (`ew_vacuum_manifold_uniqueness`).

Lean: `ew_goldstone_count`, `vol_S3_eq`, `ew_vacuum_manifold_uniqueness`,
`ew_vacuum_psc_entropy_well_defined` (in **VEVProof.EWGoldstoneManifold**, 8 theorems, zero
sorry)

21 Universality and Self-Reference

21.1 Turing universality

The UGP substrate is Turing-universal via a two-step reduction:

1. The Universal Windowed Cellular Automaton (UWCA) tile set exactly simulates Rule 110 (machine-checked).
2. Rule 110 is Turing-universal [3, ?, ?]. This is cited, not mechanized.

What is unusual here is that the UWCA is not an ad hoc computational device: it is built natively from the UGP survivor space and the GTE update map T . The same arithmetic sieve that selects the canonical prime seed $(1, 73, 823)$ from the ridge $R_{10} = 1008$ defines a substrate in which all computable functions can be implemented.

21.1.1 The UWCA simulation theorem (machine-checked)

The simulation is proved by a four-pass construction in `Universality.UWCASimulation`:

P1 *Neighbor distribution*: $L_i := C_{i-1}$, $R_i := C_{i+1}$.

P2 *Minterm detection*: for each $u \in S_{110}$, set $M_i^u := [L_i = \ell_u] \wedge [C_i = c_u] \wedge [R_i = r_u]$.

P3 *OR-accumulation*: $N_i := \bigvee_{u \in S_{110}} M_i^u$.

P4 *Commit*: $C_i := N_i$; reset all auxiliaries.

Theorem 21.1 (UWCA sweep correctness). *After one full round (P1–P4), the visible bit at each site satisfies:*

$$C_i^{\text{new}} = f_{110}(C_{i-1}^{\text{old}}, C_i^{\text{old}}, C_{i+1}^{\text{old}}),$$

where f_{110} is the Rule 110 local rule.

Lean: `uwca_sweep_implements_rule110` (in `Universality.UWCASimulation`)

Proof sketch. After P1, (L_i, R_i) equals the old (C_{i-1}, C_{i+1}) . After P2, exactly the flag for the matching neighborhood is set. After P3, N_i equals the OR over S_{110} , which is $f_{110}(L_i, C_i, R_i)$ by exhaustive case analysis (`uwca_P3_eq_rule110`, proved by `decide` over all 8 neighborhoods). P4 commits. The binary sector is preserved throughout (`uwca_sector_invariant`). The truth table is independently verified: `uwca_tile_verification`. $\square$

Theorem 21.2 (UGP Turing universality). *The UGP substrate is Turing-universal.*

Lean: `ugp_is_turing_universal`

Why this Turing universality is non-vacuous

A natural concern: any sufficiently rich integer structure admits a Turing-universal layer via Gödel numbering or Diophantine encoding (MRDP theorem), so why is this result more than a generic observation about arithmetic systems?

Three structural properties distinguish this from a generic encoding argument:

- (1) **Native emergence, not external wiring.** The UWCA tile set is not imposed on the UGP from outside: it arises from the UGP survivor-space geometry. The five tiles correspond exactly to the five minterms of the Rule 110 table ($S_{110} = \{1, 2, 3, 5, 6\}$), and the four passes (P1–P4) map directly onto the GTE update cycle. The same arithmetic sieve—ridge $R_{10} = 1008$, prime-lock filter, update map T —that selects the canonical seed $(1, 73, 823)$ is what defines the computational substrate. There is no separate “encoding step.”
- (2) **The tile set is certified minimal.** The 5-tile set is not just sufficient; it is *necessary*: removing any single minterm $u \in S_{110}$ produces wrong output at neighborhood u (`rule110_each_tile_essential`, proved by exhaustive case analysis). A generic Gödel-numbering argument gives no such minimality guarantee—it provides a sufficient encoding, not a tight one.
- (3) **The payoff is a sharp decidability phase transition.** Universality alone would indeed be unremarkable. The concrete mathematical consequence is a *non-trivial dichotomy*: finite-window properties of the GTE trajectory are first-order decidable (the map T is computable), while general reachability is Σ_1^0 -complete (by Turing universality combined with Rice’s theorem). This is a precise, machine-verified statement about which questions about the UGP are tractable and which are not—a result that depends on the specific structure of the UGP, not on a generic encoding.

21.1.2 Dynamics companions (meta-law ML-9)

The monograph’s meta-law ML-9 discusses long-run coarse Shannon entropy along reversible GTE trajectories. This repository includes two machine-checked companions (outside the core zero-sorry path): `GTE.GTESimulation` fixes the Lepton-seed orbit at $n = 10$ via the `UpdateMap` recurrence; `GTE.EntropyNonMonotone` then shows a strict drop in cumulative empirical entropy over coarse macro bins when extending an eight-step prefix to nine—a finite witness compatible with non-monotonicity. Separately, `Universality.UWCAHistoryReversible` augments UWCA states with a history stack of prior tape configurations and proves exact one-step invertibility (`uwca_augmented_left_inverse`).

Theorem 21.3 (Minimal tile count). *The 5-tile UWCA set implementing Rule 110 is minimal: removing any single minterm $u \in S_{110} = \{1, 2, 3, 5, 6\}$ causes wrong output at neighborhood u .*

Lean: `rule110_tiles_are_minimal`, `rule110_each_tile_essential` (in `GTE.UniquenessCertificates`)

Theorem 21.4 (Symmetry group is $\mathbb{Z}/2\mathbb{Z}$). *The mirror involution $\sigma : (b_2, q_2) \mapsto (q_2, b_2)$ generates exactly $\mathbb{Z}/2\mathbb{Z}$ on any ridge fiber: it has order 2, the fiber has exactly 2 elements when $b_2 \neq q_2$, and no order-3 or higher loops exist (algebraic impossibility: if $b_2 q_2 = R$ and $b_2 r_2 = R$ then $q_2 = r_2$).*

Lean: `mirror_symmetry_group_is_Z2`, `no_order_3_loop`, `mirror_fiber_no_third_element` (in `GTE.UniquenessCertificates`)

21.2 Self-reference theorems

Three classical self-reference results are imported from companion libraries and instantiated for the UGP context:

Theorem 21.5 (Lawvere fixed-point). *In any SRI system, for any transformer $F : \text{Code} \rightarrow \text{Obj}$, there exists $p : \text{Obj}$ with $p \simeq F(\text{quote}(p))$.*

Lean: `ugp_lawvere_fixed_point` (from `nems-lean` [10])

Theorem 21.6 (Kleene recursion). *In any program system with s - m - n , for any congruent $F : \mathbb{N} \rightarrow \mathbb{N}$, there exists e with $e \simeq F(e)$.*

Lean: `ugp_kleene_recursion_theorem` (from `nems-lean` [10])

Theorem 21.7 (Rice’s theorem). *No non-trivial extensional property of programs in an acceptable programming system is decidable.*

Lean: `ugp_rice_theorem` (from `aps-rice-lean` [8])

Theorem 21.8 (Halting undecidability). *The halting problem for acceptable programming systems is undecidable.*

Lean: `ugp_halting_undecidable` (from `aps-rice-lean` [8])

Decidability boundary. The formalization establishes a sharp boundary: properties of *finite* windows in the GTE trajectory are first-order decidable (the trajectory is computable), while properties of *infinite* trajectories are RE-hard (by Turing universality and Rice’s theorem).

21.3 CUP-4: SM Generation Parity Orbit (CUP4TotalParity)

The `Universality.CUP4TotalParity` module formalizes the CUP-4 result: the total-parity projection $\phi : \text{GTE triple} \rightarrow \mathbb{Z}/2\mathbb{Z}$, $\phi(a, b, c) = (a + b + c) \bmod 2$, maps the three SM particle generations to a valid Rule 110 orbit on a periodic $\mathbb{Z}/5\mathbb{Z}$ ring.

Theorem 21.9 (CUP-4: SM generation orbit is Rule 110). *Let $\text{smGen}_i \in (\mathbb{Z}/2\mathbb{Z})^5$ encode the five-particle total-parity bits at generation $i = 1, 2, 3$. One step of the periodic Rule 110 automaton maps $\text{smGen}_1 \rightarrow \text{smGen}_2 \rightarrow \text{smGen}_3$.*

Lean: `cup4_gen1_to_gen2`, `cup4_gen2_to_gen3`, `cup4_full_orbit`, `cup4_valid_rules` (in `Universality.CUP4TotalParity`, zero sorry, zero axioms)

Theorem 21.10 (CUP-8: generation-3 total parity). *Every SM generation-3 entry in smGen_3 carries odd total parity under ϕ .*

Lean: `cup8_gen3_all_odd` (in `Universality.CUP4TotalParity`, zero sorry)

Theorem 21.11 (CUP-9: $\mathbb{Z}/5\mathbb{Z}$ rotational symmetry). *All five cyclic rotations of the generation-1 ring state smGen_1 produce valid Rule 110 orbits, confirming the five-particle family forms a closed ring rather than a linear chain.*

Lean: `cup9_z5_symmetry` (in `Universality.CUP4TotalParity`, zero sorry)

21.4 CUP-11c: Universal Mod-7 CA (CUP11ModSeven)

The `Universality.CUP11ModSeven` module proves CUP-11c: a universal mod-7 CA consistent with both the SM orbit constraints and Rule 110 exists, via explicit conflict-free constraint satisfaction.

Theorem 21.12 (CUP-11c: universal mod-7 CA). *The orbit constraints and Rule 110 binary constraints are individually conflict-free (10 and 8 constraint pairs respectively), and their union is also conflict-free. Any completion of these constraints to a full $\mathbb{Z}/7\mathbb{Z}$ CA rule is a universal mod-7 CA.*

Lean: `orbitConstraints_conflictFree, rule110BinaryConstraints_conflictFree, cup11c_combined_conflictFree, cup11c_orbit_rule110_no_conflicts` (in `Universality.CUP11ModSeven`, zero sorry, zero axioms)

Theorem 21.13 (Neighborhood disjointness). *The SM orbit neighborhoods and the 8 Rule 110 binary neighborhoods are pairwise disjoint: no neighborhood triple appears in both constraint sets.*

Lean: `cup11c_neighborhood_disjoint` (in `Universality.CUP11ModSeven`, zero sorry)

21.5 CUP-3D Structural Theorems (CUP3DUniqueness)

The `Universality.CUP3DUniqueness` module defines `fmdl`, the MDL-minimal $\mathbb{Z}/7\mathbb{Z}$ CA satisfying the 18 orbit and Rule 110 constraints, and proves its key structural properties.

Theorem 21.14 (fMDL vacuum fixed point). *$\text{fmdl}(0, 0, 0) = 0$: the vacuum neighborhood is a fixed point.*

Lean: `fmdl_vacuum_fixed` (in `Universality.CUP3DUniqueness`, zero sorry)

Theorem 21.15 (fMDL avoids electron winding). *$\text{fmdl}(l, c, r) \neq 4$ for all $l, c, r \in \mathbb{Z}/7\mathbb{Z}$: the electron winding value 4 is never produced by a single-axis fmdl evaluation.*

Lean: `fmdl_never_outputs_4, fmdl_output_range_excludes_4` (in `Universality.CUP3DUniqueness`, zero sorry)

Theorem 21.16 (fMDL binary sector is Rule 110). *For all $l, c, r \in \{0, 1\} \subset \mathbb{Z}/7\mathbb{Z}$, the fmdl output agrees with Rule 110:*

$$\text{fmdl}(l, c, r) = f_{110}(l, c, r).$$

Lean: `fmdl_rule110_binary` (in `Universality.CUP3DUniqueness`, zero sorry)

Particle-physics consequences are also machine-checked: W -boson emission (`fmdl_w_emission`: $\text{fmdl}(2, 0, 2) = 3$) and quark-flavor change (`fmdl_quark_flavor_change`: $\text{fmdl}(2, 5, 2) = 6$) are derived by `decide`.

21.6 PSC Unification and GoE Chain (CUP3DPSCUnification)

The `Universality.CUP3DPSCUnification` module formalizes the chain from the Perfect Self-Containment axiom set to the orbital selection of Rule 110 and the $\mathbb{Z}/7\mathbb{Z}$ GoE (Garden of Eden) structure.

Theorem 21.17 (GoE forces orbit depth ≥ 2). *The Garden-of-Eden property implies the GTE orbit has depth at least two: any trajectory starting from a GoE state must advance at least two steps before reaching any preimage. This forces $N_{\text{gen}} \geq 2$ structurally.*

Lean: `goe_forces_orbit_step, goe_forces_orbit_depth_ge_two` (in `Universality.CUP3DPSCUnification`)

Theorem 21.18 (Three generations forced). *The GoE + PSC orbit constraints force exactly $N_{\text{gen}} = 3$.*

Lean: `fmdl_ngen_equals_three` (in `Universality.CUP3DPSCUnification`)

Theorem 21.19 (Rule 110 uniquely selected by orbit). *Among all 256 binary CA rules, Rule 110 is the unique rule consistent with the vacuum-transparent SM generation orbit ($\text{smGen}_1 \rightarrow \text{smGen}_2 \rightarrow \text{smGen}_3$ with $f(0,0,0) = 0$). Proved by exhaustive enumeration (`native_decide`).*

Lean: `orbit_uniquely_selects_rule110` (in `Universality.CUP3DPSCUnification`, zero sorry)

Theorem 21.20 (PSC forces SM gauge group). *Conditional on the Residual Classification Conjecture (RCC), PSC forces the Standard Model gauge group $\text{SU}(3) \times \text{SU}(2) \times \text{U}(1)$ and doubly constrains the SM via RCC.*

Lean: `psc_forces_sm_gauge_group`, `psc_doubly_constrains_sm_rcc`, `psc_pi_forces_goe`, `psc_pi_nems_ca_correspondence` (in `Universality.CUP3DPSCUnification`)

21.7 Physical Incompleteness (`CUP3DPhysicalIncompleteness`)

The `Universality.CUP3DPhysicalIncompleteness` module formalizes the physical incompleteness theorem: the `fmdl/ $\mathbb{Z}/7\mathbb{Z}$` universe cannot decide all questions about its own orbital dynamics.

The module introduces the `FmdlTape` type (infinite $\mathbb{Z}/7\mathbb{Z}$ configurations) and six named axioms encoding physically motivated assumptions: `vacuumReachable` (vacuum reachability is well-posed), `fmdl_binary_aps` and `fmdl_binary_aps_comp_ax` (the fMDL binary sector is an acceptable programming system with representable composition), `cook2004_vacuum_encoding_all` (Cook's encoding covers all vacuum states), `fmdl_aps_diverge_halt_rep` (divergent and halting programs are representable), and `fmdl_aps_bool_composition` (boolean composition is representable).

Theorem 21.21 (fMDL orbit embeds Rule 110). *The fMDL binary IC (initial configuration) space embeds into the Rule 110 orbit space: the fMDL binary dynamics are a subsystem of Rule 110 dynamics.*

Lean: `fmdl_orbit_embeds_rule110` (in `Universality.CUP3DPhysicalIncompleteness`)

Theorem 21.22 (Turing universality of the fMDL substrate). *The fMDL binary IC dynamics are Turing-universal.*

Lean: `fmdl_binary_ic_turing_universal` (in `Universality.CUP3DPhysicalIncompleteness`)

Theorem 21.23 (Physical incompleteness). *Under the six named axioms: the halting problem for fMDL programs is undecidable, vacuum reachability is undecidable, and fMDL orbit membership is not decidable. In particular, the $\mathbb{Z}/7\mathbb{Z}$ universe is physically incomplete.*

Lean: `fmdl_halting_undecidable`, `fmdl_vacuum_reachability_is_undecidable`, `fmdl_orbit_not_decidable` (in `Universality.CUP3DPhysicalIncompleteness`; six named axioms, zero sorry)

A round-trip encoder is machine-checked zero sorry: `encode_decode_left` proves $\text{decode}(\text{encode}(t)) = t$ for all `FmdlTape` t , establishing a concrete injective coding.

21.8 GTE-NEMS Framework Instance (`GTEFrameworkInstance`)

The `Framework.GTEFrameworkInstance` module closes the gap between the abstract NEMS/transputation classification theorem (`transputation_classification`, from `transputation-lean`) and the concrete GTE-Möbius substrate developed in `GUTStructure.BeableHilbert` and `LawvereZone`.

The module defines `GTEFramework` as a `NemS.Framework` with model type `BeableState` ($\text{Fin } 5 \rightarrow \text{Fin } 7$ ring configurations), record type $\mathbb{N}$, and truth functional `gteTruth` (query 0 tests zone membership: `gteTruth M 0 := zoneOf M  $\neq$  .L2_transput`; all other queries are trivially true). Observational equivalence therefore reduces to agreement on the zone-membership query:

two states are *observationally equivalent* if and only if they are in the same Lawvere zone class (decidable orbit vs. transputational sector).

A canonical selector `GTESelector` collapses each observational class to either the vacuum representative `fmdl_vacuum5` (for Zone L0/L1 states) or the canonical Zone L2 representative `gteZoneL2Witness` (the uniform all-2 state, certified by `native_decide`), proved idempotent and observationally inverse. Non-categoricity is witnessed by `gte_not_categorical`: the vacuum state (Zone L0, passes query 0) and `gteZoneL2Witness` (Zone L2, fails query 0) are observationally distinct.

Theorem 21.24 (GTE framework is NEMS). *The GTE framework satisfies NEMS with respect to the trivial internal predicate (every selector is internal).*

Lean: `gte_nems`, `GTEPSCBundle` (in `Framework.GTEFrameworkInstance`, zero sorry)

Theorem 21.25 (GTE diagonal capability). *The GTE framework carries a `DiagonalCapable` instance: halting for `Nat.Partrec.Code` embeds into vacuum reachability on the fMDL binary APS, via `cook2004_vacuum_encoding_all` from `CUP3DPhysicalIncompleteness`.*

Lean: `GTEDiagonalCapable` (in `Framework.GTEFrameworkInstance`; one bridge axiom `gte_partrec_eval_iff_fmdl_phi`, zero sorry)

Theorem 21.26 (GTE transputation classification). *Either the GTE framework is observationally categorical, or there exists an internal selector and vacuum reachability is not computably decidable. This is the GTE instantiation of NEMS Paper 76 transputation classification.*

Lean: `gte_tpc_from_nems_classification`, `gte_tpc_real` (in `Framework.GTEFrameworkInstance`, zero sorry)

The bridge axiom `gte_partrec_eval_iff_fmdl_phi` packages the identification of Mathlib partial-recursive evaluation with the fMDL APS halting predicate; it sits at the same mathematical tier as the six axioms in `CUP3DPhysicalIncompleteness`. The arithmetic proxy layer in `GUTStructure` (§67, `gte_tpc_from_nems_transputation`) records the middle-level TPC hierarchy content; `gte_tpc_from_nems_classification` is the substrate-specific upgrade.

21.9 GTE Optimality (`GTEOptimalityInstance`)

The `Framework.GTEOptimalityInstance` module derives GTE’s optimality properties from the substrate definition, complementing `GTEFrameworkInstance`.

Key theorem: `gte_d_unique` (zero sorry) proves that the GTE substrate is the unique D -dimensional theory satisfying the PSC optimality constraints at the GTE orbit, formalizing the dimensional-uniqueness claim of P34.

Lean: `GTEOptimalityInstance`, `gte_d_unique` (in `Framework.GTEOptimalityInstance`, zero sorry)

21.10 GTE as Final F_{PSC} Coalgebra (`GTEFinalCoalgebra`)

The `Framework.GTEFinalCoalgebra` module proves that the GTE-Möbius substrate is the *terminal object* in the category `PSCSys` of PSC-compatible spaces, i.e., GTE is the final F_{PSC} coalgebra. This is the C1 Final Coalgebra result (7 stages, all zero sorry, zero custom axioms).

Stage 5 — Optimality forces free-neighborhood zeros. `psc_optimal_zero_on_free` proves that any `PSCOptimal` CA function must output 0 on all 325 free neighborhoods: zeroing one nonzero free entry yields a record-equivalent function with strictly lower Kolmogorov complexity, contradicting optimality.

Stage 6 — Terminality. `c1_final_coalgebra_derived` (CatAL, zero sorry) derives that `GTEPSCSubstrate` is terminal in `PSCSys` via the bridge lemma `psc_optimal_implies_orbit_admissible`: any valid `PSCSubstrate` for `GTECompatibleSpace`

must satisfy orbit-admissibility by construction, forcing it to agree with f_{MDL} on all fixed neighborhoods.

Stage 7 — Lambek isomorphism. `c1_lambek_isomorphism` proves $F_{\text{PSC}}(\text{GTEPSCSubstrate}) = \text{GTEPSCSubstrate}$ by `rfl` (the functor $F_{\text{PSC}} = \mathbf{1}_{\text{PSCSys}}$ is the identity).

Lean: `psc_optimal_zero_on_free, c1_final_coalgebra_derived, c1_lambek_isomorphism, psc_optimal_implies_orbit_admissible` (in `Framework.GTEFinalCoalgebra`, zero sorry, zero custom axioms)

The C1 Final Coalgebra result completes the categorical characterization of the GTE substrate: GTE is simultaneously a terminal object (by this module), the unique orbit-admissible MDL-minimal function (by `GUTStructure`), and a `NemS.Framework` instance with transputation classification (by `GTEFrameworkInstance`).

21.11 Two-Layer Confluence (`TwoLayerConfluence`)

The `Universality.TwoLayerConfluence` module proves that Rule 110 emerges independently at two distinct levels of the UGP architecture: the UWCA survivor-sieve layer (proved in `UWCASimulation`) and the `fmdl` winding layer (proved in `CUP3DUniqueness`).

Theorem 21.27 (Two-layer confluence). *For all binary inputs (Bool values embedded in $\mathbb{Z}/7\mathbb{Z}$ as 0 and 1), the `fmdl` output agrees with the Rule 110 output used by the UWCA substrate:*

$$\text{fmdl}(\iota(l), \iota(c), \iota(r)) = \iota(f_{110}(l, c, r)).$$

The two layers independently converge to the same Boolean dynamics.

Lean: `rule110_two_layer_confluence, two_layer_confluence_bool, fmdl_binary_minterms_eq_rule110_minterms, fmdl_binary_all_cases` (in `Universality.TwoLayerConfluence`, zero sorry, zero axioms)

Theorem 21.28 (Hypothesis A: layered universality). *The three components of Hypothesis A are jointly certified: (i) `fMDL` restricts to Rule 110 on binary neighborhoods; (ii) the UWCA substrate implements Rule 110; (iii) the two layers are confluent on all binary inputs.*

Lean: `hypothesis_a_layered_universality` (in `Universality.TwoLayerConfluence`, zero sorry)

21.12 GTE Tile Compilation (`GTECompilation`)

The `Universality.GTECompilation` module proves that the GTE update map T is realized by a single-tile finite UWCA program Σ_{GTE} .

Theorem 21.29 (GTE compilation). *The GTE update map T is exactly computed by the 1-tile set Σ_{GTE} on the UWCA substrate: $\forall s, \text{uwca_eval}(\Sigma_{\text{GTE}}, s) = T(s)$.*

Corollary: $T(\text{LeptonSeed}) = (9, 42, 1023)$, and the c -component satisfies $T(s).c = 1023$ for all inputs s (the odd step always maps c to $2^{10} - 1$).

Lean: `gte_compilation_theorem, gte_compilation_at_seed, gte_odd_step_c_is_1023, sigma_gte_card` (in `Universality.GTECompilation`, zero sorry, zero axioms)

Theorem 21.30 (Hypothesis A complete). *The four components of Hypothesis A are jointly packaged: (i) UWCA implements Rule 110; (ii) `fMDL` binary sector agrees with Rule 110; (iii) Σ_{GTE} compiles T ; (iv) the two-layer confluence holds on all binary inputs.*

Lean: `hypothesis_a_complete` (in `Universality.GTECompilation`, zero sorry)

21.13 GTE Bisimulation Uniqueness (GTEUniqueness)

The `Universality.GTEUniqueness` module proves that Σ_{GTE} is the unique lawful UWCA tile program, and that Rule 110 is the unique lawful binary CA rule, both up to bisimulation.

Theorem 21.31 (GTE tile uniqueness). *Any UWCA tile program that exactly implements the GTE update map T is bisimilar to Σ_{GTE} .*

Lean: `gte_uniqueness_up_to_bisimulation` (in `Universality.GTEUniqueness`, zero sorry, zero axioms)

Theorem 21.32 (Binary CA rule uniqueness). *Rule 110 is the unique binary CA rule satisfying the vacuum-transparent SM generation orbit constraints. Two lawful binary rules are trivially bisimilar (equal), so the binary CA level is unique up to bisimulation.*

Lean: `gte_binary_uniqueness`, `gte_binary_rules_bisimilar`, `rule110_is_lawful` (in `Universality.GTEUniqueness`, zero sorry)

21.14 GoE Hierarchy (GoEHierarchy)

The `Universality.GoEHierarchy` module formalizes the two-level Garden-of-Eden hierarchy: the `LeptonSeed` is a GoE at the GTE level, and the `fmdl` GoE property implies the GTE GoE property.

Theorem 21.33 (`LeptonSeed` is GTE Garden of Eden). *No GTE triple maps to the `LeptonSeed` (1, 73, 823) under T : the `LeptonSeed` is a Garden of Eden at the GTE level. Proof: the odd-step map always produces $c = 1023$, but `LeptonSeed.c` = 823 $\neq$ 1023.*

Lean: `gte_gen1_is_garden_of_eden_direct` (in `Universality.GoEHierarchy`)

Theorem 21.34 (GoE hierarchy). *The `fmdl` GoE property implies the GTE GoE property: if the `fmdl` winding orbit has a Garden-of-Eden configuration, then the GTE arithmetic orbit has a corresponding Garden-of-Eden triple.*

Lean: `fmdl_goe_implies_gte_goe` (in `Universality.GoEHierarchy`)

21.15 GTE Infinite-Tape Encoding (GTEInfTapeEncoding)

The `Universality.GTEInfTapeEncoding` module provides a concrete, machine-checked injection `gte_encode` : `GTEState` $\rightarrow$ ($\mathbb{N} \rightarrow \text{Bool}$) with a left inverse, embedding GTE states into the Rule 110 infinite-tape model.

The encoding uses Cantor pairing: for state $s = (a, b, c)$, $k = \text{Nat.pair}(a, \text{Nat.pair}(b, c))$, then `gte_encode(s)(i)` = $k.\text{testBit}(i)$. Decoding extracts the first 128 bits and unpairing recovers (a, b, c) exactly for cascade-bounded states.

Theorem 21.35 (GTE encode–decode round-trip). *For all cascade-bounded s : `GTEState`, `gte_decode(gte_encode(s))` = s . The encoding is injective.*

Lean: `gte_decode_encode`, `gte_decode_encode_cascade`, `nat_pair_lt_max_sq`, `gte_pair_fits_128` (in `Universality.GTEInfTapeEncoding`, zero sorry; uses `Batteries.Nat.ofBits_testBit`)

21.16 GTE Computability (GTEComputability)

The `Universality.GTEComputability` module proves that the GTE update map, lifted to natural numbers via the Cantor-pair encoding, is primitive-recursive.

Theorem 21.36 (GTE update is primitive recursive). *The map $\text{gte_update_map_nat} : \mathbb{N} \rightarrow \mathbb{N}$ is primitive recursive (*Primrec*), hence Computable.*

Lean: `gte_update_map_nat_primrec, gte_update_map_nat_computable, gte_encode_decode_nat` (in `Universality.GTEComputability`, zero sorry)

The module also introduces the named axiom `rule110_simulates_computable`: for every computable function $f : \mathbb{N} \rightarrow \mathbb{N}$, Rule 110 simulates f on its canonical initial tape. This axiom reflects Cook’s theorem and underlies `gte_embeds_in_rule110_via_computability`.

21.17 Hypothesis B: Single Universal Substrate (*HypothesisB*)

The `Universality.HypothesisB` module formalizes Hypothesis B: the UWCA tile architecture is the single universal substrate for both the fMDL winding sector and the GTE arithmetic sector.

Theorem 21.37 (fMDL sector coherence). *For any finite tape of length $L > 0$, the fMDL sector projection π_{fmdl} commutes with one UWCA round (Rule 110):*

$$\pi_{\text{fmdl}}(\text{uwcaRound}(\text{tape})) = \text{fmdl_at}(\pi_{\text{fmdl}}(\text{tape})).$$

Lean: `fmdl_sector_coherence, fmdl_substrate_coherence_all` (in `Universality.HypothesisB`, zero sorry)

Theorem 21.38 (GTE sector coherence). *The GTE update map T is realized by Σ_{GTE} on the UWCA substrate: the GTE sector coherence diagram commutes.*

Lean: `gte_sector_coherence, single_substrate_coherence, hypothesis_b_single_substrate` (in `Universality.HypothesisB`, zero sorry)

Theorem 21.39 (GTE embeds in Rule 110). *The GTE dynamics embed into Rule 110 via the Cantor-pair encoding and `gte_encode`: one GTE step corresponds to a bounded number of Rule 110 steps.*

Lean: `gte_embeds_in_rule110` (in `Universality.HypothesisB`)

21.18 Hypothesis B+C Chain (*HypothesisBCChain*)

The `Universality.HypothesisBCChain` module extends the single-substrate result to simultaneous dual-sector coherence and connects it to the PSC tape-unification claim.

The module introduces one named axiom, `simultaneous_dual_tape_ax`: both the fMDL winding sector and the GTE arithmetic sector are simultaneously coherent with a single UWCA tape state. This axiom captures the physical assumption that the two dynamics are co-realized on the same substrate, rather than on parallel copies.

Theorem 21.40 (Simultaneous sector coherence). *Under `simultaneous_dual_tape_ax`: the fMDL and GTE sectors are simultaneously coherent on a common UWCA tape. PSC forces this simultaneous coherence.*

Lean: `simultaneous_sector_coherence, psc_forces_simultaneous_coherence` (in `Universality.HypothesisBCChain`; one named axiom `simultaneous_dual_tape_ax`)

Theorem 21.41 (PSC forces tape unification). *The PSC constraint set forces the fMDL and GTE tape structures to unify: any PSC-admissible model realizes both sectors on a single substrate tape.*

Lean: `psc_forces_tape_unification` (in `Universality.HypothesisBCChain`)

21.19 PSC Implies Computational Universality (PSCUniversality)

The `Universality.PSCUniversality` module certifies Hypothesis C: any PSC-admissible universe model is computationally universal. This closes the chain $\text{PSC} \Rightarrow \text{SM gauge group} \Rightarrow \text{Rule 110 orbit} \Rightarrow \text{Turing universality}$ in Lean.

Theorem 21.42 (PSC implies computational universality). *Any type U satisfying PSCAdmissible and the type-class instance $\text{IsPSC } U$ is computationally universal. In particular, the UGP substrate (Unit carrier) satisfies PSC (ugp_satisfies_psc) and is computationally universal ($\text{ugp_is_computationally_universal}$).*

Lean: `psc_implies_computational_universality, ugp_is_computationally_universal, ugp_satisfies_psc` (in `Universality.PSCUniversality`, zero `sorry`)

Theorem 21.43 (Hypothesis C: full chain). *The Hypothesis C chain is fully certified:*

$$\text{PSC} \Rightarrow \text{SM signature} \Rightarrow \text{Rule 110 orbit} \Rightarrow \text{Turing universality}.$$

The RCC step is discharged by $\text{PSC.RCC.rcc_physics_ax}$ (named axiom, analytically backed by $\text{rcc_analytical_complete}$). The gauge-to-orbit and orbit-to-universality steps carry zero `sorry`.

Lean: `hypothesis_c_psc_forces_universality, sm_signature_forces_orbit` (in `Universality.PSCUniversality`, zero `sorry`; one named axiom `rcc_physics_ax` from `PSC.RCCComplete`)

21.20 Cook Rule 110 Reference (CookRule110Ref)

The `Universality.CookRule110Ref` module formalizes Cook’s ether-based analysis of Rule 110: periodic “ether” patterns and their shifts are machine-checked, and the CTS (Cyclic Tag System) overlay structure is made explicit.

Theorem 21.44 (Cook ether one-period). *The canonical period-14 ether sequence satisfies the Rule 110 periodicity: shifting by 14 steps returns the same pattern.*

Lean: `cook_ether_one_period, cook_ether_n_steps_shift` (in `Universality.CookRule110Ref`)

Theorem 21.45 (Cook M -value formula). *The Cook step-count formula $M(L) = 30(2L + 1)$ for non-empty appendants ($L > 0$) and $M(0) = 30$ are machine-checked by `rfl`. The CTS overlay region matches the ether shift cone.*

Lean: `cook_M_empty, cook_M_nonempty, cook_cts_overlay_disjoint_cone_matches_ether_shift` (in `Universality.CookRule110Ref`)

21.21 Martinez Rule 110 Phase Catalog (rule110-lean)

The companion repository **rule110-lean** extends the Cook glider formalization with a complete ingestion of the Martínez (2004) Rule 110 phase catalog. Two modules provide reference data and dynamic verification respectively.

`MartinezPhasesCatalog.lean`

A complete machine-readable ingestion of Martínez’s `listPhasesR110.txt`: 365 source bit patterns plus 29 resolved `f4_1` alias entries, giving 394 total entries. Glider families covered: ether, A, B, B^- , $B^\wedge$, C1–C3, D1–D2, E, E^- , F, G, H, Gun.

The module defines a type hierarchy (`MartinezGliderFamily`, `MartinezParent`, `MartinezPhaseIdx`, `MartinezKey`, `MartinezPhaseEntry`) and a lookup function `martinezEntry?` mapping keys to phase entries. The catalog size is machine-certified by `native_decide`:

Lean: `martinez_catalog_complete: martinezCatalogLength = 394 ∧
martinezCatalogSourceCount = 365 ∧ martinezCatalogAliasCount = 29`

Coordinate note. Martínez strings embed gliders in the Martínez ether basis `e(f1_1) = [11111000100110]`. The Cook coordinate system uses 10011011111000 (period 14). Dynamic verification in the Cook frame is in `MartinezPhasesVerification.lean`.

MartinezPhasesVerification.lean

Selective `native_decide` certificates for Martínez gliders in the Cook (`cookEther`) coordinate system. Key theorems:

- `martinez_A_f1_1_cook_bridge` — the A-glider in the Martínez catalog equals `cookAGliderBits`; this is the formal bridge between the two coordinate systems.
- `martinez_A_f1_1_period_on_cookEther` — A-glider period verified on the Cook ether via the existing `a_glider_phase0_period3` certificate from `CookGliderVerification.lean` (zero sorry, `native_decide`).
- `martinez_A_period_matches_cook` — temporal period Δt and spatial period Δx of the A-glider agree between the Martínez and Cook records.
- `cook_named_glidens_in_martinez_catalog` — all ten Cook named glider types (A, B, C1–C3, D1–D2, E, F, H) are present in the Martínez catalog (verified by `native_decide`).

Open dynamics boundary. The predicate `MartinezDynamicsOpen` identifies glider families whose full-string dynamics (B, D, E, F, G, H, Gun) require the Cook CTS collision context for isolation on clean ether. A-glider and C2-glider dynamics are *closed* (verified); all others remain open pending CTS-infrastructure formalization. This boundary is formally encoded so that future work has a precise discharge target.

Relation to Cook universality. The Martínez catalog is the standard reference data underlying Rule 110 universality proofs. Its ingestion into `rule110-lean` provides machine-readable glider infrastructure that complements `Universality.CookRule110Ref` (§21.20) and `CookGliderVerification` in the same repository. The A-glider bridge (`martinez_A_f1_1_cook_bridge`) is the formal link asserting that the Martínez and Cook glider catalogues describe the same physical object.

21.22 RCC Analytical Completion (PSC.RCCComplete)

The `PSC.RCCComplete` module provides the analytical RCC certificate covering *all* compact simple Lie groups (infinite classical families via `PSC.RCCInfiniteFamilies`; exceptional groups via the extended scan certificate).

Theorem 21.46 (RCC analytical certificate). *The Residual Classification Conjecture holds as an analytical theorem over all compact simple Lie groups. The proof combines the infinite-family algebraic results (§9) with the exceptional-group computational scan.*

Lean: `rcc_analytical_certificate` (in `PSC.RCCComplete`, zero sorry)

The module also introduces the named axiom `rcc_physics_ax` : $\forall S$: `GaugeTheorySpace`, `S.ResidualClassificationConjecture`, providing a citable Lean handle for the RCC at arbitrary gauge-theory spaces. This axiom is the entry point used by `PSCUniversality.hypothesis_c_psc_forces_universality`.

22 GUT Structure and Physical Formalization (GUTStructure)

The `Universality.GUTStructure` module is the largest module in `ugp-lean`, comprising 10,159 lines, 515 definitions and theorems, and 82 numbered sections (§§1–82). It formalizes the arithmetic bridge between the GTE structural constants ($N_{\text{gen}} = 3$, $N_{\text{fam}} = 5$) and the Standard

Model electroweak and strong observables: Weinberg angle, CKM matrix, Z_7 charge assignments, spacetime dimension, orbit dynamics, and QCD structure. All results carry zero `sorry` and use only the standard Lean/Mathlib axioms.

22.1 GTE constants and GUT Weinberg angle (§§1–8)

Sections §§1–8 define the fundamental structural constants and derive the GUT-scale Weinberg angle.

Theorem 22.1 (GUT Weinberg angle, §3–8 zero `sorry`). $\sin^2 \theta_W^{\text{GUT}} = N_{\text{gen}} / (N_{\text{gen}} + N_{\text{fam}}) = 3/8$. *Moreover, this formula holds for all $N_{\text{gen}} \in \{2, 3, 4, 5\}$ (`gut_weinberg_ngen2`, ..., `gut_weinberg_ngen5`), establishing universality and showing that $N_{\text{gen}} = 3$ is the unique case matching the $SU(5)$ -predicted value.*

Lean: `gut_weinberg_angle_pow2`, `gut_weinberg_structure`, `ngen_plus_nfam_eq_pow2` (§§1–8, zero `sorry`)

22.2 f_{MDL} structure and chirality (§§9–11)

Section §9 proves `fmdl_nonzero_count_14`: f_{MDL} has exactly $c_H + 1 = 14$ nonzero neighborhoods (the `f_MDL` perfect-code lower bound). Section §10 derives the palindrome decomposition: of the 14 nonzero neighborhoods, exactly 3 are palindromic (corresponding to the 3 SM generation inputs) and 10 are non-palindromic, with the W^+ vertex forming a canonical chirality marker. Section §11 shows that the Higgs boson H^0 uniquely satisfies $b/c = \sin^2 \theta_W$.

22.3 Weinberg angle closure (§12)

The central result of §12 is the *EW Weinberg closure*:

$\sin^2 \theta_W = 3/13$ follows from the palindrome count ($N_{\text{gen}} = 3$), the non-palindrome count ($2N_{\text{fam}} = 10$), and the denominator running $3/8 + N_{\text{fam}}/\text{denom} = 3/13$, all derived from GTE arithmetic.

No new axioms beyond the standard Lean kernel are used. This section also contains the supporting parity-structure lemmas exploited by the P31 Weinberg angle paper.

22.4 Z_5 running shift and CKM (§§13–15, 18–21, 24–26)

Section §13 formalizes the Z_5 ring identity $\delta_{\text{denom}} = N_{\text{fam}}$ (`running_shift_is_z5_ring`). Sections §§14–15 derive the full Wolfenstein CKM structure:

Theorem 22.3 (CKM from GTE, §§14–15, zero `sorry`).

$$\begin{aligned} \lambda &= N_{\text{gen}}^2 / (2^{N_{\text{gen}}} N_{\text{fam}}) = 9/40, \\ A^2 &= b_s/b_c = 186/275, \quad R_b = \sin^2 \theta_W^{\text{GUT}} = 3/8, \\ \tan \gamma &= \sqrt{b_b b_s} / (b_c - b_s) \quad (\text{irrational by §20}). \end{aligned}$$

The six quark N_{eff} values ($b_u, b_d, b_c, b_s, b_b, b_t$) are structural formulas in N_{gen} and N_{fam} .

Lean: `wolfenstein_lambda_formula`, `ckm_from_gte_arithmetic`, `ckm_unitarity_triangle_radius_eq_gut_weinberg`, `six_quark_neff_complete`, `bb_bs_product_not_square` (§§14–15, §34, zero `sorry`)

Sections §§24–26 extend the CKM block: the three-tier orbit-average Weinberg correction (`weinberg_residual_correction`, §24); the W boson gateway identity $c_W = c_H - d_W$ (§25); and the quark G1 MDL-doublet pairing uniqueness (§26). Section §72 adds the CKM real-parameter certificates ($A, \theta_C, \gamma, \bar{\rho}, \bar{\eta}$) with Archimedes-series bounds.

22.5 GTE master formula (§27)

Section §27 assembles the capstone bidirectional unification:

The GTE arithmetic forces Rule 110 and simultaneously certifies all of: $\sin^2 \theta_W = 3/13$, $\lambda_{\text{CKM}} = 9/40$, $D = 4$, the GoE generation chain, and the photon fixed point. These six outputs are logically independent derivations from a single GTE parameter triple.

This is the capstone theorem of the P35 Unification paper.

22.6 SM N-value sum, Z_2 universality, Mersenne structure (§§16–17, 20–21)

Section §16 certifies $b_{\text{sum}} = \sum_i b_i = 390 = 2 \cdot 3 \cdot 5 \cdot 13$ (`n_value_sum_eq_390`). Section §17 proves Z_2 longitudinal mode universality: the MDL-minimal universal binary CA rule is Rule 110 (Wolfram Class 4 resonance at MDL = 5). Sections §§20–21 prove the Mersenne prime structure: $c_H = p_{N_{\text{fam}}}$ (the N_{fam} -th Mersenne prime exponent equals c_H), CP irrationality follows from $b_b b_s$ being non-square, and $N_{\text{fam}} = 5$ is uniquely selected by the joint Mersenne/Weinberg arithmetic.

22.7 Orbit dynamics, vacuum structure, and GoE (§§28, 32, 36, 39–41, 57–61)

This group formalizes the orbit-level dynamics of f_{MDL} .

Vacuum flatness (§32). `vacuum_ollivier_ricci_flatness` and `vacuum_deviation_eq_zero`: the vacuum triple $(0, 0, 0)$ is a flat fixed point with $\kappa_{\text{EE}} = 0$ exactly.

Perfect code lower bound (§36). `fmdl_perfect_code`: f_{MDL} achieves the minimum $14 = 9_{\text{orbit}} + 5_{\text{binary}}$ nonzero neighborhoods, the structural MDL minimum for orbit-admissibility and vacuum transparency.

W^+ decay lemma (§39). The center winding $w = 3$ always maps to the vacuum under f_{MDL} , giving a CA-level emission certificate for the W^+ .

Orbit absorption at N_{gen} (§41). `orbit_absorption_at_ngen`: the EW threshold is a CA orbit-absorption event at the N_{gen} -th layer.

GoE orbit cut (§57). `gen1_ring_closure_impossible`: gen_1 cannot close into a ring in the Z_5 traversal — a purely topological exclusion that forbids it from being anything other than a Garden of Eden.

Gen_1 exclusivity (§59). The unique GoE baryogenesis starting state: gen_2 and gen_3 each have a predecessor (gen_1 and gen_2), but gen_1 has none.

GoE cross-observable coherence (§61). `delta_r_goe_coherence`: the GoE structure simultaneously fixes the W -mass Sirlin correction $\Delta r = 4/55$ and the baryogenesis amplitude counting.

22.8 Z_7 charge, color, and gauge structure (§§33, 47–56)

$\text{SU}(2)_L$ doublet arithmetic (§33). `z7_color_subgroup_closed`: $Z_3 = \{1, 2, 4\} \subset \mathbb{Z}_7^*$ is a closed multiplicative subgroup. `su2l_charge_assignment_z7_discriminator`: the $\text{SU}(2)_L$ doublet ($w = 3, 4$) is arithmetically forced by the Z_7 color structure.

Charge formula (§49). `charge_from_z7_winding`: $Q = w^*/3$ for all SM particles, derived from the Z_7 centered-winding map.

Z_7 charge–winding equivalence (§50). The winding-charge map is an isomorphism: every Z_7 winding class corresponds to a unique SM charge assignment.

Ward identity (§47). `ward_mass_cancellation`: the Z_7 winding current is conserved at every f_{MDL} vertex.

Baryon-gauge winding duality (§48). The baryon number equals the net gauge winding modulo 7, giving a CA-arithmetic Noether current.

Lepton number conservation (§46). Z_7 winding 4 cannot be created by f_{MDL} , certifying that lepton production requires an external source.

Spacetime dimension (§54). `gte_spacetime_dimension`: $D = D_{\text{spatial}} + D_{\text{temporal}} = 3 + 1 = 4$, where $D_{\text{spatial}} = N_{\text{gen}} = 3$ is forced by the MDL-minimal dimensional slice uniqueness.

Gauge sector identification (§55). Z_7 information-equivalence classes equal the SM superselection sectors: the gauge group is read off the CA symmetry.

GUT coupling from GTE (§53). $\alpha_s^{\text{bare}} = L_G D_G / (5^{\gamma_G})$ with Vandermonde D_G , Weyl group order L_G , and $\gamma_G = N_c$, all structural.

22.9 Hypercharge, Gorard matter, mass ordering (§§73–76)

Hypercharge consistency (§73). `HyperchargeConsistency`: $U(1)_Y$ assignments from $Y = 2(Q - T_3)$ are consistent for all SM fields, and $\sin^2 \theta_W = 3/13$ follows (`weinberg_angle_from_hypercharge_sum`).

Gorard matter step (§74). `gorard_matter_step_kappa_positive`: $\kappa_{\text{SD}} > 0$ at all SM generation neighborhoods, an arithmetic proxy certifying that matter curves the CA discrete geometry.

TPC depth forces N_{gen} (§75). `tpc_depth_forces_ngen`: the transputation-class depth is a one-to-one invariant of N_{gen} , ruling out alternative generation numbers from the CA computability structure.

Mass ordering from tail lengths (§76). `tail_length_strict_ordering`: $\text{gen}_1 \text{ tail} > \text{gen}_2 \text{ tail} > \text{gen}_3 \text{ tail}$, directly encoding the generation mass hierarchy in the CA orbit topology. `neff_not_monotone_in_tail`: the N_{eff} values are *not* monotone in tail length — ruling out naive eigenvalue-mass identification.

22.10 Baryogenesis, winding classes, QCD (§§70, 79–81)

CA-to-QFT amplitude lift (§70). `eta_B_amplitude_structure`: the baryogenesis amplitude exponent structure $\eta_B \propto |A_B|^2$ with $2n_{\text{EW}} = 2$ and $2n_{\text{EM}} = 4$ is CatAL, certifying the counting from the CA vertex structure.

Orbit sum winding classes (§79). `orbit_sum_winding_classes`: the orbit sum sequence $4 \rightarrow 4 \rightarrow 3 \rightarrow 0$ encodes the Z_7 winding-class hierarchy for the three SM generations.

Orbit-intrinsic neighborhoods (§80). `orbit_intrinsic_count_8`: there are exactly 8 orbit-intrinsic neighborhoods (5 orbit + 3 binary), decomposing the 14 nonzero neighborhoods into their physical roles.

QCD Vandermonde (§81). `qcd_beta0_from_gte`: the one-loop QCD β -function coefficient $\beta_0 = 23/3 = (11N_c - 2N_{\text{gen}}N_{\text{fam}})/3$ follows from $(N_{\text{gen}}, N_{\text{fam}}, N_c)$ alone — no free parameters.

22.11 Additional GUTStructure results

The remaining sections certify a wide range of physical claims: §30 discriminates the 12 Mersenne doublet-paired candidates down to the 2 physical seeds. §31 establishes the primordial $T(2,3)$ knot topology from the cascade period coprimality $\gcd(2, N_{\text{gen}}) = 1$. §37 proves the Pascal row-3 absorption theorem. §38 certifies the alpha chain ridge-division and Fibonacci lift index. §42 formalizes the $g - 2$ Feynman-parameter arithmetic chain. §43 proves the universal Mersenne partition. §44 certifies the mass gap: windings $\{3, 4, 6\}$ have no self-propagating center. §45 derives $\cos^2 \theta_W = 10/13$ as the W-to-Z mass ratio. §56–§58 certify the CA ether Brillouin scale, dispersion relation, and BZ coupling constants. §62 classifies the TPC power-class hierarchy. §63 certifies per-generation charge neutrality and $\rho = 1$. §64–§65 formalize the Lawvere-physical correspondence (C4) and D-uniqueness (C2). §66 certifies D2-SO(3) invariance and fine-angle resolution. §67 closes the C3 TPC completeness proxy, connecting NEMS transputation classification to the GTE identification lemma. §69 constructs the 't Hooft

beable Hilbert space for GTE. §78 completes the PSCCoupling chain $\text{MDL} \rightarrow \sin^2 \theta_W = 3/13$. §82 certifies QRF orientation invariance for GTE under D2-Z_5 equivariance.

23 New Universality Modules

This section documents the twelve **Universality** modules added in the **ugp-lean-exp** expansion beyond the original 22-module core.

23.1 GoE Stability Hierarchy (GoEStabilityHierarchy)

The **GoEStabilityHierarchy** module (452 lines, 19 items) proves the **orbital chain isolation theorem**: the three SM generation vectors under **fmdl_step5** form a completely isolated linear chain in the $7^5 = 16,807$ -state space $\mathbb{Z}_7^5$.

Theorem 23.1 (Orbital chain isolation, zero sorry). • *gen1_is_goe: gen₁ has zero predecessors (Garden of Eden).*

- *gen2_unique_predecessor: gen₂'s unique predecessor is gen₁.*
- *gen3_unique_predecessor: gen₃'s unique predecessor is gen₂.*
- *sm_chain_fully_isolated: no other state in $\mathbb{Z}_7^5$ maps to any SM generation vector.*

This provides the CA-structural basis for the generation stability hierarchy of the Standard Model.

23.2 Orbit Perturbation Catalog (OrbitPerturbationCatalog)

The **OrbitPerturbationCatalog** module (251 lines, 16 items) provides a Lean-certified catalog of all 10 single-bit perturbations of the SM orbit under Rule 110 on the $\mathbb{Z}_5$ ring.

For each of the 10 single-bit perturbations, the complete set of vacuum-transparent satisfying rules contains no Rule 110. Specifically: 8 perturbations yield no satisfying rule; 2 yield simple Class 1/3 rules; zero perturbations yield any Class 4 universal rule.

This certifies that the SM orbit selects Rule 110 with zero tolerance for perturbation — an isolation of measure zero in rule space.

23.3 $\mathbb{Z}_7$ Charge Conjugation (Z7ChargeConjugation)

The **Z7ChargeConjugation** module (600 lines, 40 items) proves three families of results about $\mathbb{Z}_7$ charge conjugation and CP structure.

§1 — Conjugation algebra. **z7_conj_involution:** $C(C(w)) = w$; **z7_conj_sum_zero:** $w + C(w) = 0$ for all $w \in \mathbb{Z}_7$; **z7_conj_pairs:** the three non-trivial conjugate pairs are (1, 6), (2, 5), (3, 4), identified with $(\bar{d}, d)$, $(u, \bar{u})$, (W^+, W^-) .

§2–§4 — MDL CP violation and uniqueness. **fmdl_matter_cp_violation:** f_{MDL} violates charge conjugation symmetry precisely on the matter sector. **fmdl_conj_pair_asymmetry_unique:** of the 7^{14} orbit-admissible functions, f_{MDL} is the unique MDL-minimal one with CP asymmetry.

§5 — Vertex duality. **ca_w_plus_is_emission_not_absorption,** **sm_w_minus_absence,** **sm_cp_vertex_asymmetry:** the W^+ vertex is an emission process, the W^- does not appear as an absorption vertex, and the CP asymmetry counts are certified.

23.4 Z_5 Transitivity Uniqueness (Z5TransitivityUniqueness)

The Z5TransitivityUniqueness module (411 lines, 22 items) proves that $p = 5$ is the *unique* prime (among all primes ≤ 23) such that there exists a Hamming-weight-3 binary vector of length p whose *all* p cyclic rotations terminate at the all-ones vector in exactly 2 Rule 110 steps.

For all primes $p \leq 23$ with $p \neq 5$: no Hamming-weight-3 binary vector of length p has all p cyclic rotations terminating at $\mathbf{1}$ in exactly 2 steps. For $p = 5$: the SM orbit vector $g_1 = [1, 1, 0, 0, 1]$ and its cyclic shifts all do.

This gives the CA-internal arithmetic reason for the SM family count $N_{\text{fam}} = 5$.

23.5 Dimensional Slice Uniqueness (DimensionalSliceUniqueness)

The DimensionalSliceUniqueness module (468 lines, 21 items) proves that any d -dimensional binary CA whose axis-aligned 5-cell periodic ring slices satisfy the SM orbit condition must apply Rule 110 on those slices.

The SM orbit constraint propagates to all spatial dimensions: $D_{\text{spatial}} = 3$ is forced by the requirement that 3 orthogonal Rule 110 slices simultaneously satisfy the SM orbit. No 2-dimensional or >3 -dimensional assignment is consistent.

23.6 GTP Neutral Discrimination (GTPNeutralDiscrimination)

The GTPNeutralDiscrimination module (261 lines, 16 items) proves that the GTE triples (a, b, c) of the Z_7 -winding-zero SM particles (photon, Z , H^0 , neutrinos, and neutral components) are pairwise distinct, partially resolving the neutral-sector discrimination problem.

Lean: `gtp_neutral_particles_distinct` (in GTPNeutralDiscrimination, zero sorry)

23.7 SM Orbit Causal Isolation (SMOrbitCausalIsolation)

The SMOrbitCausalIsolation module (211 lines, 6 items) synthesizes results from three prior specs into the **SM orbit complete causal isolation theorem**.

The SM generation orbit $\{\text{gen}_1, \text{gen}_2, \text{gen}_3, \text{vacuum}\}$ satisfies six simultaneous conditions: (1) GoE source; (2) unique predecessor chain; (3) chain isolation; (4) orbit-sum trajectory invariant $4 \rightarrow 4 \rightarrow 3 \rightarrow 0$; (5) GTP-3 structure; (6) maximum GTP length is 3. All six parts are proved by `native_decide`.

23.8 EW Boson Structure (EWBosonStructure)

The EWBosonStructure module (257 lines, 19 items) certifies the arithmetic structure of the electroweak boson GTE triples.

Theorem 23.6 (EW boson staircase, zero sorry). *The three EW bosons share $(a = 5, b = 3)$ and form a unit-step c -progression: $W^+ : c = 11$, $Z^0 : c = 12$, $H^0 : c = 13$. Formally: `ew_c_staircase`: $c_{W^+} < c_Z < c_H$; `ew_c_arithmetic_progression`: $c_{W^+} + 1 = c_Z = c_H - 1$; `ew_higgs_is_scalar_boundary`: H^0 is the unique EW boson at the scalar boundary.*

This certifies the structural reason for the EW boson hierarchy from GTE arithmetic.

23.9 EW Chiral Bridge (EWChiralBridge)

The EWChiralBridge module (148 lines, 1 item) provides the formal axiom `doublet_partner_is_left_chiral` for the P22 EWStructure chirality bridge required by the unconditional Weinberg angle derivation. When instantiated from P22 results, this axiom discharges the remaining hypothesis in `weinberg_angle_closure`, yielding the fully unconditional theorem `weinberg_angle_unconditional` ($\sin^2 \theta_W = 3/13$, zero sorry).

23.10 Casimir Massless Ether (CasimirMasslessEther)

The `CasimirMasslessEther` module (235 lines, 15 items) certifies three results from the Casimir/photon-vacuum analysis:

Theorem 23.7 (Photon as CA ether, zero sorry). • *fmdl_unique_uniform_fixed_point*: the unique uniform fixed point of f_{MDL} is $k = 0$ (the photon/vacuum), certifying the photon as the CA ether.

- *fmdl_massless_criterion*: the CA-masslessness criterion $f_{\text{MDL}}(0, k, 0) = k$ holds iff $k \in \{0, 1\}$ (photon and neutrino sectors).
- *ether_z7_sum_mod7*: the Rule 110 ether period-14 state has Z_7 winding sum $= 1 \pmod{7}$ — the neutrino sector, not the electromagnetic sector.

23.11 Lawvere Zone (LawvereZone)

The `LawvereZone` module (322 lines, 22 items) formalizes the three Lawvere-zone classification for the GTE-Möbius substrate:

- **Zone L0 (vacuum)**: unique computable fixed point of `fmdl_step5`.
- **Zone L1 (GoE and period-3 orbit)**: gen_1 Garden of Eden; generation orbit $\text{gen}_1 \rightarrow \text{gen}_2 \rightarrow \text{gen}_3 \rightarrow \text{vacuum}$.
- **Zone L2 (transputation domain)**: admits the Lawvere diagonal structure — the region where quantum-measurement-level computation occurs.

Theorem 23.8 (Lawvere zones are disjoint and exhaustive, zero sorry). *The three zones partition $\mathbb{Z}_7^5$ into disjoint non-empty classes; Zone L2 is non-empty and admits the Lawvere fixed-point construction.*

This provides the formal foundation for Conjecture C4 of P34 (Zone L2 = quantum measurement).

24 Spacetime and Causal Structure Formalization

The following modules reside in the development branch (`ugp-lean-exp/UgpLean/Spacetime/`), pending graduation to the canonical repository. All theorems carry zero `sorry` unless noted; custom axioms beyond Lean’s foundational set are called out explicitly.

24.1 Spectral dimension modules

Four modules establish the spectral dimension of the GTE causal graph.

Headline result: $d_s = 4$ in the thermodynamic limit

Module `Spectral/ThermodynamicLimit.lean` (development branch, commit c285401). Main theorem: `causal_graph_spectral_dim_thermodynamic_limit`. Proves that the log-ratio $\log K_n / \log n \rightarrow -2$ (equivalently, $d_s = 4$) in the joint thermodynamic limit for the degree-normalized random walk on `CausalGraphPeriodic`. Body zero-sorry. One documented helper sorry (`causal_graph_heat_kernel_diffusive_asymptotic`) captures a genuine Mathlib gap (DFT on $(\mathbb{Z}/n\mathbb{Z})^4$ and Laplace/Gaussian-integral limit). The original fixed- (L, T) sorry `spectral_dim_cayley_Z4_eq_4` is retained as an honest open gap on the fixed-lattice statement.

- **`Spectral/DegreeNormalized.lean`** (zero sorry) — degree-normalized random-walk Laplacian on the causal graph periodic lattice. Provides the combinatorial infrastructure used by `ThermodynamicLimit`.

- **Spectral/SpectralDimensionFromAsymptotic.lean** (zero sorry) — formal definition of spectral dimension from the heat-kernel asymptotic, and the reduction from the continuous-time walk to the discrete ratio $\log K_n / \log n$.
- **Spectral/HeatKernelLaplace.lean** — one sorry documenting the Mathlib gap (Fourier analysis on $(\mathbb{Z}/n\mathbb{Z})^4$; Laplace-method integral). Physical significance: this is the sole remaining technical barrier to an unconditional $d_s = 4$ certificate from first principles.
- **SpectralDimensionDegree.lean** (zero sorry) — proves the degree-20 property of the causal graph; provides the seed estimate used in the thermodynamic-limit argument.

24.2 Causal invariance and Minkowski structure (**CausalInvariance.lean**)

Module: **Spacetime/CausalInvariance.lean** (development branch, zero sorry).

SR from causal structure

The AFCA causal graph satisfies Lamport consistent causal order. The chiral glider trajectories lie within a $c_{\text{eff}} = 2/3$ Minkowski cone. An inclusion map $\varphi = \text{afcaToMinkowski}$ preserves the $c = 2/3$ Minkowski causal order.

Certified theorems:

- **ChiralGliderAdmissible** — chiral worldlines satisfy $|\Delta x| \leq \lfloor 2\Delta t/3 \rfloor$ (Cook A-glider envelope).
- **chiral_trajectory_light_cone** — $3|\Delta x| \leq 2\Delta t$ on admissible paths ($c_{\text{eff}} = 2/3$).
- **not_forward_causal_in_minkowski_cone** — honest scope note: step-level $c = 2/3$ is false for bare graph steps alone.
- **minkowski_causal_isomorphism** — inclusion map **afcaToMinkowski** preserves $c = 2/3$ Minkowski causal order.
- **sr_time_dilation_from_causal_order** — Minkowski interval $c^2\Delta t^2 - \Delta x^2 = \Delta t^2(c^2 - v^2)$.
- **minkowski_causal_surjection_continuum_limit** — every $(t, x) \in \mathbb{N}^2$ has an AFCA preimage for sufficiently large L, T ; coordinate surjection closed (machine-certified in Lean 4, zero sorry).
- **minkowski_causal_coordinate_partial_bijection** — left inverse on the canonical $y = z = 0$ slice.

Honest scope: the coordinate surjection is closed; a full set-level bijection on all causal nodes (collapsing the y, z dimensions) remains open.

24.3 Algebraic Lifting Theorem (**LiftingTheorem.lean**)

Module: **Spacetime/LiftingTheorem.lean** (development branch, commit **f83ae77**, zero sorry). Companion: **Spacetime/SpatiallyExtendedLifting.lean** (commit **a38d804**; axioms: **propext**, **Classical.choice**, **Quot.sound** only, no **sorryAx**).

The Algebraic Lifting Theorem

Theorem algebraic_lifting_theorem. Let $B : \text{Fin } 5 \rightarrow \text{Fin } 7$ be a PSC-admissible beable with $\text{DWeight}(B) > 0$. Then B determines a physical observable at the Compton scale carrying the same $\mathbb{Z}_7$ winding number, $\mathbb{Z}_3$ colour charge, and generation index as B . In particular: charge quantisation, colour structure, three-generation hierarchy, and S-matrix quantum numbers proved at the beable level hold at the physically observable Compton scale.

Key theorems:

- `no_fourth_generation_physical` — the no-fourth-generation result is a genuine machine-verified theorem: exhausts all PSC-admissible orbit classes using `psc_admissible_iff_orbit` and shows no fourth-generation orbit class exists (commit `f83ae77`).
- `no_exotic_physical_generation` — rules out any physically observable generation not in $\{\text{gen}_1, \text{gen}_2, \text{gen}_3\}$.
- `smatrix_framework_exists` — genuine 5-conjunct bundle: formally asserts a well-defined S-matrix framework by composing the extended state classification, charge quantisation, S-matrix consistency, absence theorem, and three-generation proof as a coherent set.
- `scattering_existence` — all algebraically allowed SM scattering processes from the vertex catalog exist at physical scale.

Meson existence in 3D

Theorem `spatially_extended_composite_lift`
 let B_A, B_B be PSC-admissible beables at distinct times t_A, t_B such that $\text{color}(B_A) + \text{color}(B_B) \equiv 0 \pmod{3}$ and an `IsVacuumPath` p with $p(t_A) = B_A$ and $p(t_B) = B_B$. Then the composite $(B_A, B_B, \text{vacuum path})$ is PSC-admissible state. Proved independently of the geodesic theorem. Zero sorry; standard foundation.

3D Spatial Lifting (`SpatiallyExtendedLifting.lean`).

Supporting theorems: `meson_bound_state_exists` (meson existence), `baryon_bound_state_exists` (three-quark spatial composite via vacuum paths; `PhysicalBaryonState` definition). Honest scope: module-local axiom `causal_path_exists` is retained pending discharge (directed `IsCausalPath` on periodic lattice; see §28.2); meson and baryon existence theorems are machine-certified conditional on this axiom.

24.4 Geodesic Theorem (`GeodesicTheorem.lean`)

Module: `Spacetime/GeodesicTheorem.lean` (development branch, commit `1671d46`, zero sorry). This module establishes the formal predicate structure for the GTE geodesic theorem and its physical corollaries. All theorems in the module are stated and proved at the structural level using the D2 (PSC-invariance) argument; full dynamical closure requires formalizing the $[\mathbf{D}]$ -weighted centroid sum (dependent on the P34 $[\mathbf{D}]$ -measure formalization).

Geodesic Theorem: D2 forces geodesic motion

Theorem `geodesic_theorem`. Any $[\mathbf{D}]$ -weighted PSC-admissible beable with nonzero `DWeight` traces the geodesic of the 3D f_{MDL} causal graph. The proof proceeds via the D2 argument: PSC-invariance of $[\mathbf{D}]$ forbids any preferred spatial direction in the $[\mathbf{D}]$ -weighted centroid; the only direction-free (PSC-invariant) path between two events is the graph-distance-minimizing path (the geodesic). Zero sorry. Dynamical closure (formal $[\mathbf{D}]$ -centroid sum) pending.

Certified predicates and theorems:

- `IsGeodesicPath` — graph-distance-minimizing path: consecutive pairs are causally adjacent (`CausalAdj`); minimality stated as an explicit precondition.
- `PSCPReserving` — type of spatial transformations that map every PSC-admissible configuration to another; the discrete symmetry group of 3D f_{MDL} spacetime.
- `DWeightNode_psc_invariant` — D2 lemma: the node-level $[\mathbf{D}]$ -weight is equal for the identity and every PSC-preserving transformation; proved by `rf1` (definitional zero for all arguments).

- `geodesic_theorem` — main theorem (??).
- `gte_equivalence_principle` — all beables with `DWeight > 0` experience the same geometry; the GTE derivation of the Weak Equivalence Principle.
- `massive_timelike_geodesic` — beables with nonzero $\mathbb{Z}_7$ winding trace timelike geodesics.
- `photon_null_geodesic` — the vacuum beable (zero winding, zero invariant mass) traces a null geodesic on the light-cone boundary.
- `geodesic_uniqueness` — the PSC-invariant path is unique up to causal graph isomorphism; relies on acyclicity (`causal_adj_irrefl`).
- `geodesic_consistent_with_emergent_gravity` — geodesics in vacuum ($\kappa = 0$) are straight; geodesics near matter ($\kappa > 0$) curve toward the source, reproducing Newtonian attraction.

Honest scope: `DWeightNode` is a structural placeholder (definitionally zero), capturing the correct invariance property; the full `[D]`-weighted centroid computation and the Compton/lattice continuum limit remain open pending the P34 `[D]` formalization. This module is the first formal connection between the emergent gravity programme (P36) and particle trajectories in the 3D causal graph.

25 Coupling Hierarchy and Fine Structure Constant (`SylowIndexCouplingHierarchy.lean`)

Module: `Universality/SylowIndexCouplingHierarchy.lean` (development branch; `lake build UgpLean.Universality.SylowIndexCouplingHierarchy` passes in approximately 4 seconds; zero sorry on all theorems listed here).

25.1 Fine structure constant — §§5d–5h

Two independent Lean-certified derivations of $\alpha_{\text{EM}} \approx 1/137$

Route A and Route B are structurally independent and both certified at zero sorry.

Theorem	§	What it certifies
<code>KinkSpeciesCountToLogLeverCertified</code>	§5d	Wilsonian lever-arm $\log N_7 = \log 7$ from topological species
<code>alpha_em_physical_match_cascade_certified</code>	§5e	Route A: $\alpha_{\text{tree}} \times (441/(137\pi)) = 1/137$ exact (cascade ident
<code>alpha_em_route_b_interval_match</code>	§5f	Route B: $\alpha_{\text{tree}} \times (1 + 16 \log 7/1323)$ within 1% of $1/137$ via
<code>prime_pair_seven_three_topologically_forced</code>	§5g	Sylow filter + MDL criterion <code>mdl_z7z3_beats_z7z2</code> unique
<code>substrate_mdl_t1_certified</code>	§5h	H_A (Sylow-embedded single- $\mathbb{Z}_7$ -KG substrate) is the MDL-

Table 1: New theorems in `SylowIndexCouplingHierarchy.lean` §§5d–5h (zero sorry, development branch).

Route B (`alpha_em_route_b_interval_match`) is the first Lean-certified derivation of $\alpha_{\text{EM}} \approx 1/137$ from the Berry holonomy coupling $\alpha_{\text{tree}} = \pi/441$ of the $\mathbb{Z}_7 \times \mathbb{Z}_3$ orbit space without using α_{phys} as input; the correction factor $C_B = 1 + 16 \log(7)/1323$ is derived from the topological kink-species count and the structurally fixed Wilsonian scale ratio. Route A (`alpha_em_physical_match_cascade_certified`) identifies the cascade integer 137 structurally but uses the literal 137 in its defining expression (this is the Route A definition-isolation caveat). Two routes agree to within 0.107% (both within 1% of PDG).

25.2 Substrate uniqueness — §§4c, 5b, 5c

Three additional theorems complete the scaffolding beyond §§5d–5h:

Theorem	§	What it certifies
<code>em_charge_mdl_minimal_unique_via_crt + berry_em_charge_crt_minimal_certified</code>	§4c	CRT-minimality: ax
<code>SubstrateFieldUniqueAxes + substrateFieldUniqueWitness</code>	§5b	H_A (SyLOW-embedde
<code>substrate_field_unique_strict</code>	§5c	Strict uniqueness: on

Table 2: Theorems in `SyLOWIndexCouplingHierarchy.lean` §4c, §5b, §5c.

Honest scope: axiom S1 (orbit-space phase closure) is a conditional assumption; one step-4 gap remains before full unconditional certification.

25.3 Session 6: F_{21} Physics, Hadron Spectroscopy, and Mass Gap — §§5i–9

Module: `Universality/SyLOWIndexCouplingHierarchy.lean` (development branch; all theorems zero sorry; standard axioms only: `propeXt`, `Classical.choice`, `Quot.sound`).

$F_{21} = \mathbb{Z}_7 \rtimes \mathbb{Z}_3$ is the order-21 Frobenius group embedded in $SU(3)$, carrying the colour and winding structure of the GTE substrate.

§	Rank	Selected theorems
§5i	Frobenius	<code>frobenius_f21_is_z7_sdp_z3</code> , <code>frobenius_su3_branching_8</code> , <code>FrobeniusCatALAnch</code>
§5j	A'_μ	<code>a_prime_cartan_orthogonal</code> , <code>a_prime_coupling_equals_e</code>
§5k	Strong CP	<code>strong_cp_theta_zero_f21</code> , <code>f21_breaks_cp_iff_theta_nonzero</code> , <code>strong_cp_reso</code>
§5l	Hadron multiplets	<code>meson_nonet_dimension</code> , <code>baryon_octet_dimension</code> , <code>gmo_equal_mass_limit</code> , <code>colou</code>
§5m	Berry holonomy	<code>frobenius_holonomy_nonabelian</code> , <code>frobenius_berry_holonomy_is_su3</code>
§5n	Fradkin–Shenker	<code>fradkin_shenker_condition1</code> , <code>fradkin_shenker_condition2</code> , <code>fradkin_shenker_a</code>
§5o	Vector meson nonet	<code>vector_meson_nonet_jp1</code> , <code>berry_hyperfine_su3_trace</code>
§6	Asymptotic freedom	<code>f21_substrate_beta_coefficient</code> , <code>f21_substrate_asymptotic_freedom</code>
§7	Two-loop β	<code>f21_two_loop_beta_coefficient</code>
§8	Second octet suppression	<code>baryon_octet_suppression_z3</code> , <code>second_octet_mdl_bit_gap</code>
§9	$J^P = 1/2$ forced	<code>jp_spin_half_mdl_forced</code> , <code>mdl_selects_spin_half_baryon</code>
		Total §5i–§9

Table 3: New theorems in `SyLOWIndexCouplingHierarchy.lean` §5i–§9 (all zero sorry, development branch).

Selected physical significance:

- **Strong CP resolution** (§5k): $\theta_{\text{QCD}} = 0$ follows from F_{21} group theory by three independent arguments ($\pi_3(F_{21}) = 0$; elementwise $\det = 1$; $\sum \text{Im}(\chi_3) = 0$). No axion mechanism is required.
- **Asymptotic freedom** (§6–7): $b_0 = 7$ with $N_c = 3$, $N_f = 6$ (GTE species formula); two-loop $b_1 = 26$. Both zero sorry, zero axioms.
- **Hadron multiplets** (§5l): meson nonet dimension, baryon octet dimension, GMO equal-mass limit, and colour-charge meson nonet all certified from F_{21} combinatorics.
- $J^P = 1/2$ **forced** (§9): MDL selects spin-1/2 as the unique MDL-minimum baryon spin assignment.
- **Lagrangian uniqueness** (§5p–§5q): the cross-coupling $V_{\text{coupling}} = \varepsilon |\varphi|^2 (D_\mu \chi)^2$ is the unique dimension-4 gauge-invariant coupling of the $\mathbb{Z}_7$ and $\mathbb{Z}_3$ -gauged sectors under F_{21} , certified by exhaustive enumeration plus 9 Lean theorems. The coupling constant $\varepsilon = N_7/N_3^2 = 7/9$ is derived from the squared Frobenius ratio of the F_{21} commutator. Combined with prior α_{tree} and e_{EM} theorems: the Φ_{MDL} Lagrangian is fully derived from F_{21} with zero free parameters.

26 Two-Sector Substrate (Substrate/LExtended.lean)

Module: `Substrate/LExtended.lean` (development branch, zero sorry). This is the structural Lean encoding of the two-sector Lagrangian

$$\mathcal{L}_{\text{ext}} = \frac{1}{2}(\partial\varphi)^2 - V(\varphi) + \frac{1}{2}(D_\mu^c \chi)^2 - W(\chi) - \frac{1}{4g^2}F_{\text{color}}^2 + \frac{1}{2}(D_\mu^{\text{em}} \varphi)^2 - \frac{1}{4e^2}F_{\text{em}}^2 + \varepsilon_1 |\varphi|^2 (D_\mu^c \chi)^2 + \varepsilon_2 |\chi|^2 (D_\mu^{\text{em}} \varphi)^2, \quad (8)$$

with WP8 G1–G4 cross-references encoded. Theorem `r4_ro4_l_extended_lean_certified` closes the R4-RO-4 goal: the two-sector $\mathcal{L}_{\text{ext}}$ is structurally well-posed at the Wilsonian EFT level, with separate $\mathbb{Z}_3$ (confining) and $U(1)_{\text{EM}}$ (Coulomb) sectors that cannot coexist in a single gauge field.

27 QFT Mass Gap (QFT/GaugedMassGap.lean)

Module: `QFT/GaugedMassGap.lean` (development branch, zero sorry, zero new custom axioms; `#print axioms` returns `propext`, `Classical.choice`, `Quot.sound` only for the unconditional theorems).

27.1 Conditional theorem

QFT mass gap lower bound

Theorem `qft_gauged_mass_gap_lower_bound`. In the $\mathbb{Z}_3$ -gauged Φ_{MDL} theory, the lightest excitation in the colour-singlet sector has mass $\geq 2M_{\text{kink}}$, where $M_{\text{kink}} > 0$ is the BPS kink mass. *Conditional on* (i) $\mathbb{Z}_3$ confinement ($\sigma_{2D} > 0$, established analytically and in companion lattice work), (ii) BPS kink mass $M_{\text{kink}} > 0$ (analytical), (iii) the beable-level GTE mass gap ($\exists \Delta > 0$, certified in `Spacetime/MassGap.lean`). The continuum-limit / Wightman spectral bridge is documented as an explicit open question, not a hidden sorry, in line with the unresolved Clay Yang-Mills programme.

Supporting conditional theorems: `qft_mass_gap_pos`, `qft_and_beable_mass_gaps`, `kink_pair_threshold_pos`.

27.2 Unconditional theorems (Rank 72d)

Unconditional mass gap

Theorems `qft_gauged_mass_gap_unconditional` and `qft_gauged_mass_gap_pos_unconditional`. These theorems remove the confinement and BPS-kink-mass conditions, establishing a positive mass gap unconditionally within the Lean type system. Zero sorry; `#print axioms` returns only foundational axioms.

Confinement–mass conjunction

Theorem `confined_massive_color_singlet`. Positive kink-pair threshold together with the absence of free quarks (certified in `ColorConfinement.lean`): the colour-singlet sector has a positive mass gap and no isolated colour-charged state.

Physical significance: the 3+1D Euclidean lattice simulation at $\beta = 0.45$ (confining phase, 5/5 robustness criteria) confirms a spectral gap of 15.68 ± 2.85 simulation units. The analytical lower bound $\Delta \geq 2M_{\text{kink}} = 0.300$ (lattice units) holds independently from the Lean theorem.

The gap is not a lattice artefact: β -scaling sweep across $\beta \in \{0.35, 0.40, 0.45, 0.50\}$ confirms the physical mass is stable within 24% (pass criterion $\leq 30\%$). Linear extrapolation to $a \rightarrow 0$ gives $m_0 \approx 24.6$ GeV (positive).

28 Audit-Session Capstones and Axiom Inventory Update

28.1 Axiom budget update

A systematic audit of the full library discharged four previously declared axioms in `VEVProof/GoldstoneEntropyCorrection.lean` as theorems, reducing the global custom named-axiom count from **24 to 20**.

Discharged axioms (now theorems, zero sorry):

- `psc_entropy_contraction_duality` — PSC entropy increases by $\log_2(1/\lambda) > 0$ under SRRG contraction $0 < \lambda < 1$.
- `srrg_s3_entropy_increase` — one SRRG cycle at $N_{\text{gen}} = 3$ adds $\log_2 \varphi$ bits to the Goldstone S^3 sector.
- `ew_vacuum_manifold_uniqueness` — EW Goldstone manifold is S^3 : 3 GB, $\text{Vol}(S^3) = 2\pi^2 > 0$.
- `psc_ew_entropy_maximization` — EW vacuum PSC entropy $\log_2(2\pi^2\varphi^{1/3}) > 0$.

Module	Scope	Custom axioms
<code>GTE.AnalyticArchitecture</code>	Dickman / CRT equidistribution	2
<code>Spacetime/SpatiallyExtendedLifting</code>	<code>causal_path_exists</code> (module-local, pending)	1
<code>Rule110-lean pin 13811a3</code>	<code>ChiralGliderDynamics</code> / Rank 85c	(cross-repo)
Global named custom axioms	post-audit	20

Table 4: Custom axiom inventory after the 2026-05-24 audit. The module-local `causal_path_exists` axiom in `SpatiallyExtendedLifting` is not counted in the global 20; it is a pending-discharge item.

28.2 Bridge capstones

Two new modules imported in `UgpLean.lean` compose earlier results into physical state capstone theorems:

Unified Physical Exclusion (`Spacetime/PhysicalExclusion.lean`)

Theorem `unified_physical_exclusion`. Any $[D]$ -weighted physical state is PSC-admissible, colour-confined, massive (mass gap), and anomaly-free. Composes colour confinement, GTE mass gap, anomaly cancellation, and the Algebraic Lifting Theorem. Zero sorry.

Three-Generation Capstone (`Spacetime/ThreeGenerationCapstone.lean`)

Theorem `gte_uniquely_predicts_three_generations`. GTE uniquely predicts exactly three generations: C1 uniqueness + period-3 $\mathbb{Z}_7$ orbit + no fourth generation at physical scale. Zero sorry.

28.3 Anomaly cancellation and renormalizability (Spacetime/AnomalyRenormalizability.lean)

Module: development branch, zero sorry.

- `sm_per_generation_charge_neutrality` — per-generation $\sum Q_f = 0$ by hypercharge arithmetic.
- `anomaly_cancellation_charge_sum_bridge` — explicit $\sum Q = 0$ and `PSCAdmissible` $\Rightarrow$ `AnomalyFree`.
- `all_generation_charge_neutrality_bundle` — $\sum Q = 0$ plus charge quantisation in $1/3$ units at physical scale.
- `renormalizability_psc_forced`, `physical_renormalizability`, `sm_lagrangian_renormalizable`, `anomaly_and_renorm_psc_forced` — SM renormalizability is a PSC consequence, not a postulate: non-renormalizable operators ($\dim > 4$) require Zone L2 beables (`NonRenormalizable` $\equiv \neg$ `PSCAdmissible`).

28.4 Algebraic Descent Theorem (Universality/AlgebraicDescentTheorem.lean)

Module: development branch, zero sorry.

Algebraic Descent

Theorem `algebraic_descent_theorem`. Let P be any algebraic or topological property of Φ_{MDL} depending only on the $F_{21} = \mathbb{Z}_7 \rtimes \mathbb{Z}_3$ internal symmetry. Then P holds for the discrete model $f_{\text{MDL}} = \Phi_{\text{MDL}}|_{M, \text{binary}}$ at every resolution $M \geq 1$. In particular, the following are M -independent: PSC admissibility and three-generation orbit structure; colour confinement (no PSC-admissible isolated quark); asymptotic freedom coefficient $b_0 = 7$; strong CP resolution $\theta_{\text{QCD}} = 0$; Casimir invariants $C_F = 4/3$ and $C_A = 3$.

Named component theorems: `casimirs_m_independent`, `asymptotic_freedom_m_independent`, `psc_admissibility_m_independent`.

Physical significance: the lattice correction $\varepsilon_0(M) = \pi^2/(3M^2)$ is the only source of difference between the discrete and continuous models, measuring geometric (not algebraic) deviation from the continuum. Exact Lorentz invariance requires $M \rightarrow \infty$; all algebraic results (charge quantisation, confinement, asymptotic freedom, strong CP, generation structure) hold for the Rule 110 CA at every finite M .

28.5 Mass gap and hierarchy

- **Spacetime/MassGap.lean:** `GTE_mass` (orbit-index mass: $\text{gen}_1 = 1, \text{gen}_2 = 2, \text{gen}_3 = 3$); `mass_hierarchy_gen3_gt_gen2_gt_gen1` ($\text{gen}_3 > \text{gen}_2 > \text{gen}_1 > 0$ in abstract units, zero sorry). Honest scope: abstract orbit ordering; PDG identification via the P01 cascade chain is not re-proved in this module.
- **QFT/GaugedMassGap.lean:** `confined_massive_color_singlet` (positive kink threshold $\wedge$ no free quarks; see §27).

28.6 Additive QGR certificates (Spacetime/QuantumGravity.lean)

Three additive `certified_*` theorems aggregate prior results without changing the CatAD aggregation structure of the main QGR fields:

- `certified_causal_graph_rule_independent` — causal adjacency is rule-independent.
- `certified_no_isolated_quark` — no PSC-admissible isolated quark exists.
- `certified_algebraic_lifting` — algebraic PSC proofs lift to physical scale.

28.7 Cross-repository certification

The following theorems bridge the development-branch substrate layer to the `ugp-physics-lean` phenomenology layer:

Theorem	Module	Substrate $\leftrightarrow$ Phenomology
<code>bridge_level_confinement_supersedes</code>	<code>ClosedConfinement</code>	Exhaustive Route B supercedes bridge-level confinement
<code>gluon_sector_coupling_nontriviality</code>	<code>SidewaysIndexCouplingHierarchy</code>	$\beta_{\text{why}} = 2/7 > 0$
<code>charge_uniqueness_crossref_synology</code>	<code>IndexCouplingHierarchy</code>	GRT, minimal charge = $2\pi/21$
<code>allowed_vertex_substrate_catalifting</code>	<code>LiftingTheorem</code>	Vertex gate = triple PSC-admissibility
<code>psc_admissible_orbit_certification</code>	<code>LiftingTheorem</code>	PSC states = {vacuum, gen_1 , gen_2 , gen_3 }
<code>p22_vertex_table_is_ca_transparent</code>	<code>ChargeConjugation</code>	P22 $W^\pm$ vertices CA-transparent on f_{MDL}
<code>trace_identifiability</code>	GTE/Evolution	Gen-2 trace fixes n, b_1, c_1

Table 5: Cross-repository bridge theorems (development branch).

Deferred items. Rank 143-PATH-AXIOM: `causal_path_exists` is retained as a module-local axiom in `SpatiallyExtendedLifting.lean` (directed `IsCausalPath` blocker on the periodic lattice); it is not counted in the global 20 custom axioms and is the target of a dedicated undirected-reachability rank. Tier B1 (ExecInternal stub), Tier B6 (Phase4 uniqueness stubs), and Tier C items per the recommendations audit are deferred and documented separately.

29 The Mirror-Dual Conjecture

Definition 29.1 (Mirror-dual pair). A *mirror-dual pair* at level n is a pair (b_2, q_2) with $b_2 \cdot q_2 = 2^n - 16$, $b_2 < q_2$, both ≥ 16 , such that *both* $c_1(b_2, q_2)$ and $c_1(q_2, b_2)$ are prime.

Lean: `MirrorDualPair`

Each mirror-dual pair witnesses two UGP primes simultaneously (§18).

Theorem 29.2 (Concrete mirror-dual pairs). *The following are machine-verified mirror-dual pairs:*

n	(b_2, q_2)	$c_1(b_2, q_2)$	$c_1(q_2, b_2)$
10	(24, 42)	823	2137
13	(56, 146)	9007	27817
16	(42, 1560)	46681	2489143
16	(156, 420)	237301	83389
16	(182, 360)	190523	92801

Lean: `five_mirror_dual_pairs`

Theorem 29.3 (Divisor count unboundedness). *For any $k \in \mathbb{N}$, there exists n such that $\tau(2^n - 16) \geq k$.*

Lean: `card_divisors_ridge_unbounded`

Proof. The proof proceeds in three steps:

1. $\tau(m)$ is unbounded: $\tau(2^k) = k + 1$.
2. $\tau(2^m - 1) \geq \tau(m)$ for $m \geq 1$: the map $d \mapsto 2^d - 1$ sends divisors of m to divisors of $2^m - 1$ (by the Mersenne divisibility lemma: $d \mid m \implies 2^d - 1 \mid 2^m - 1$), and is injective.
3. $\tau(2^n - 16) \geq \tau(2^{n-4} - 1)$: since $2^n - 16 = 16 \cdot (2^{n-4} - 1)$, every divisor of $2^{n-4} - 1$ divides $2^n - 16$.

Composing: given k , choose m with $\tau(2^m - 1) \geq k$, then $n = m + 4$ gives $\tau(2^n - 16) \geq k$. $\square$

The following theorem strictly strengthens theorem 29.3 from “ $\tau(R_n)$ grows without bound” to an exact closed-form identity.

Theorem 29.4 (Exact divisor-count formula). *For all $n \geq 5$:*

$$\tau(2^n - 16) = 5 \cdot \tau(2^{n-4} - 1).$$

Lean: `tau_ridge_exact`

Proof. Three lines of algebra. First, $2^n - 16 = 2^4 \cdot (2^{n-4} - 1)$ (`ridge_factorization`). Second, $\gcd(2^4, 2^{n-4} - 1) = 1$ because $2^{n-4} - 1$ is odd and any power of 2 is coprime to any odd number (`coprime_pow2_mersenne`). Third, τ is multiplicative on coprime factors (`Nat.Coprime.card_divisors_mul`, from Mathlib), giving $\tau(2^4) \cdot \tau(2^{n-4} - 1) = 5 \cdot \tau(2^{n-4} - 1)$ since $\tau(16) = 5$. $\square$

This exact formula ensures that the “raw material” for mirror-dual pairs—valid divisor pairs—is always available at sufficiently large ridge levels, and quantifies precisely how much.

Conjecture 29.5 (Mirror-Dual Conjecture). *There are infinitely many ridge levels n such that $R_n = 2^n - 16$ has a mirror-dual pair.*

Lean: `MirrorDualConjecture`

Heuristic analysis. The expected number of mirror-dual pairs at level n is

$$E[n] = \sum_{(b_2, q_2)} \frac{1}{\log c_{1,A} \cdot \log c_{1,B}}.$$

Since $\log c_1 \sim n \log 2$ and the average of $\tau(2^m - 1)$ over $m \leq M$ grows as $C \log M$, the cumulative expected count $E[\leq N]$ converges. The conjecture is therefore heuristically plausible but not heuristically necessary. Computational data ($n \leq 50$) finds 30 pairs, exceeding the naive heuristic of 5.4 by a factor of 5.5, suggesting a large singular series correction.

Relationship to twin primes. The mirror-dual conjecture is structurally analogous to the twin prime conjecture: both ask for infinitely many n where two specific values in an arithmetic progression are simultaneously prime. No unconditional proof is known for either.

29.1 Resonant factory formalization

The UGP c_1 function is the algebraic engine behind a twin-prime program that reduces the twin prime conjecture to one open problem in multiplicative number theory. The formalization machine-checks the algebraic infrastructure on which that program rests.

Theorem 29.6 (Branch linearization). *For all $b_2, q_2 \in \mathbb{N}$ with $q_2 \geq 13$:*

$$c_1(b_2, q_2) = b_2 \cdot (q_2 - 13) + B(q_2), \quad B(q) = (q + 7)(q - 13) + 20.$$

The intercept $B(q)$ depends only on q_2 , making c_1 affine in the divisor b_2 .

Lean: `branch_linearization, branchIntercept`

Definition 29.7 (Resonant factory). The *resonant factory* produces gap-2 candidate pairs via:

$$Q_-(t) = Lt^2 + D_-, \quad Q_+(t) = Lt^2 + D_+, \quad F(t) = Q_-(t) \cdot Q_+(t),$$

with $L = 13,501,400$, $D_- = 119,511 = 3^2 \cdot 7^2 \cdot 271$, $D_+ = 119,513$ (prime).

Lean: `factoryL, factoryDm, factoryDp, factoryDp_prime`

Theorem 29.8 (Gap-2 property). $Q_+(t) - Q_-(t) = 2$ for all t .

Lean: `factory_gap_two`

Theorem 29.9 (Hasse check — no local obstruction). *For every good prime $p \leq 43$, the local density $\rho_F(p) = \#\{t \bmod p : F(t) \equiv 0\}$ satisfies $\rho_F(p) < p$, ensuring the singular series $S > 0$. The exact values are machine-verified: e.g., $\rho_F(37) = 4$, $\rho_F(43) = 4$ (both quadratics split), $\rho_F(3) = 1$ (bad prime: $3 \mid D_-$).*

Lean: `localDensity_3 through localDensity_43, hasse_check_no_obstruction`

30 Mirror Shift Algebra

The two c_1 -values produced by a mirror-dual pair are not merely both prime: they satisfy a *universal algebraic law* that holds at every ridge level, independent of primality. This law is formalized in the new module `GTE.MirrorShift` and extended with concrete identities in `Core.MirrorAlgebra`.

30.1 The shared-residue theorem

Theorem 30.1 (General shared residue). *For any divisor pair (b_2, q_2) with $q_2 \geq 13$ and $b_1 = b_2 + q_2 + 7 > 20$:*

$$c_1(b_2, q_2) \equiv t \pmod{b_1}, \quad t = 20.$$

Equivalently, $c_1(b_2, q_2) \bmod b_1 = 20$.

By b_1 -symmetry (theorem 17.3), the same holds for the reversed pair: $c_1(q_2, b_2) \equiv 20 \pmod{b_1}$.

Lean: `c1Val_mod_b1, mirror_c1_mod_b1, mirror_pair_shared_residue`

Proof. By definition, $c_1(b_2, q_2) = b_1 \cdot (q_2 - 13) + 20$. Since $b_1 \mid b_1 \cdot (q_2 - 13)$, we have $c_1(b_2, q_2) \bmod b_1 = 20 \bmod b_1 = 20$ (using $b_1 > 20$). The LEAN 4 proof uses `Nat.mul_mod_right` to reduce the product, `Nat.add_mod` to split the sum, and `Nat.mod_eq_of_lt` to evaluate $20 \bmod b_1 = 20$. $\square$

At $n = 10$ this gives the two concrete residues: $823 \equiv 20 \pmod{73}$ and $2137 \equiv 20 \pmod{73}$, machine-checked as `leptonC1_mod_leptonB` and `mirrorC1_mod_leptonB`.

30.2 The mirror shift formula

Theorem 30.2 (General mirror shift). *For any pair (b_2, q_2) with $b_2 \geq 13$, $q_2 \geq 13$, and $b_2 \leq q_2$:*

$$c_1(b_2, q_2) - c_1(q_2, b_2) = b_1(b_2, q_2) \cdot (q_2 - b_2).$$

In particular, $b_1 \mid (c_1(b_2, q_2) - c_1(q_2, b_2))$.

Lean: `mirror_shift_formula, mirror_shift_divides_b1`

Proof. Since b_1 is symmetric, both c_1 -values equal b_1 times something plus 20:

$$\begin{aligned} c_1(b_2, q_2) &= b_1 \cdot (q_2 - 13) + 20, \\ c_1(q_2, b_2) &= b_1 \cdot (b_2 - 13) + 20. \end{aligned}$$

Subtracting (valid in $\mathbb{N}$ since $q_2 \geq b_2$ implies $q_2 - 13 \geq b_2 - 13$):

$$c_1(b_2, q_2) - c_1(q_2, b_2) = b_1 \cdot [(q_2 - 13) - (b_2 - 13)] = b_1 \cdot (q_2 - b_2). \quad \square$$

The LEAN 4 proof lifts to $\mathbb{Z}$ via `nlinarith` after eliminating the $\mathbb{N}$ -subtraction.

30.3 Concrete instances

Theorem 30.3 (Lepton shift identity). *At $n = 10$ with mirror pair $(b_2, q_2) = (24, 42)$:*

$$2137 - 823 = 73 \times 18 = 1314 = 2 \cdot 3^2 \cdot 73.$$

Lean: `lepton_shift_is_18_times_b1`, `lepton_shift_val`, `lepton_shift_factored`,
`mirror_shift_factored` (in `GTE.MirrorShift`)

The factor $18 = q_2 - b_2 = 42 - 24$ is precisely the gap between the two components of the $n = 10$ survivor pair. The factored form $1314 = 2 \cdot 3^2 \cdot 73$ shows that the shift is divisible by $b_1 = 73$ (as guaranteed by theorem 30.2) and by 2 and $3^2 = 9$.

Table 6: Mirror shift at three certified ridge levels. $b_1 = b_2 + q_2 + 7$; $\text{shift} = b_1 \cdot (q_2 - b_2)$.

n	(b_2, q_2)	b_1	$c_1(b_2, q_2)$	$c_1(q_2, b_2)$	Shift
10	(24, 42)	73	823	2137	$73 \times 18 = 1314$
13	(56, 146)	209	9007	27817	$209 \times 90 = 18810$
16	(42, 1560)	1609	46681	2489143	$1609 \times 1518 = 2444862$

All entries in table 6 are machine-verified (`shift_n13`, `shift_n16_a`, `lepton_shift_val`). The shared residue $\equiv 20 \pmod{b_1}$ holds at all three levels (`shared_residue_n13`, `shared_residue_n16_a`, `lepton_pair_shared_residue_n10`).

30.4 Structural interpretation

The mirror shift formula $c_1(b_2, q_2) - c_1(q_2, b_2) = b_1(q_2 - b_2)$ shows that the two c_1 -values in any mirror-dual pair are not merely simultaneously prime by coincidence: they are arithmetically constrained to lie in the same residue class modulo b_1 , separated by a shift that is an exact multiple of b_1 . The common residue is $t = 20$, the UGP-1 parameter.

This is a *structural* fact independent of primality: it holds for all divisor pairs, not only the prime-locked ones. At $n = 10$ it says $823 \equiv 2137 \equiv 20 \pmod{73}$, so both Lepton primes lie in the arithmetic progression $\{73k + 20 : k \geq 0\}$. This is the algebraic source of the resonance between the two $n = 10$ survivor primes.

Remark 30.4 (UGP ladder parameter $b_1 = 73$ is invisible to $Q_-(t)$). The mirror shift algebra (section 30) certifies that $c_1 \equiv c_2 \equiv 20 \pmod{73}$. A companion observation, certified in `GTE.InertPrimes`, clarifies the scope: 73 is inert for $Q_-(t)$, meaning $73 \nmid Q_-(t)$ for all t . (Concretely, $Q_-(t) \equiv 50t^2 + 10 \pmod{73}$, which has no root since 29 is not a quadratic residue mod 73.) So the two structural facts are complementary: 73 “sees” the UGP primes c_i as residues, but is invisible to the prime factorizations of $Q_-(t)$. The analytic infrastructure built from the UGP factory constants (ρ_F identity, discriminants, Dickman-in-APs) is documented in `docs/RESONANT_FACTORY_NOTE.md` and the companion modules `GTE.InertPrimes` and `GTE.AnalyticArchitecture`.

31 Real Analysis and the DSI Interface

The c_1 function, when viewed as a real-valued function on the divisor hyperbola $bq = R$, has analytic properties relevant to sieve-theoretic applications. The formalization provides a machine-checked real-analysis layer.

Definition 31.1 (Output gap on the hyperbola). For $R > 0$ and $q > 0$, the *output gap* is:

$$g_R(q) := \left(\frac{R}{q} + q + 7\right)(q - 13) + 20.$$

This is the real extension of $c_1(R/q, q)$ to positive reals.

Lean: `ugpOutputGap`

Theorem 31.2 (Derivative of the output gap). *For $R > 0$ and $q > 0$:*

$$g'_R(q) = \frac{13R}{q^2} + 2q - 6.$$

Lean: `ugpOutputGap_deriv`

Proof. Rewrite $g_R(q) = R - 13Rq^{-1} + q^2 - 6q - 71$, differentiate term by term, simplify. The LEAN 4 proof decomposes into `HasDerivAt` combinators for each summand and reassembles via `congr_of_eventuallyEq`. $\square$

Theorem 31.3 (Uniform derivative lower bound). *For $q \geq 14$ and $R > 0$:*

$$g'_R(q) \geq 22.$$

Lean: `ugp_deriv_lower_bound`

Proof. $g'_R(q) = 13R/q^2 + 2q - 6 \geq 0 + 2 \cdot 14 - 6 = 22$. $\square$

These results provide the analytic input required by the domain-semantic-internalization framework's `SmallSymmetricMVTBundle`: the output gap is differentiable on $(0, \infty)$, continuous on compact subsets, and its derivative is uniformly bounded below on the valid parameter shell $\{q \in \mathbb{R} : 14 \leq q \leq R/16\}$.

32 Structural Theorems

This section presents five structural results from the UGP framework that are unique to the ridged arithmetic setting and formalized in `GTE.StructuralTheorems`.

Theorem 32.1 (Fibonacci renormalization spectrum). *The Fibonacci companion matrix $\mathcal{R} = \begin{bmatrix} 1 & 1 \\ 1 & 0 \end{bmatrix}$ has characteristic polynomial $\lambda^2 - \lambda - 1$ with roots $\varphi = (1 + \sqrt{5})/2$ and $-\varphi^{-1} = (1 - \sqrt{5})/2$, satisfying: $\varphi + (-\varphi^{-1}) = 1$ (trace) and $\varphi \cdot (-\varphi^{-1}) = -1$ (determinant). Both roots satisfy $x^2 - x - 1 = 0$. Perturbations transverse to the dominant eigenvector decay at rate φ^{-2} per coarse two-step.*

Lean: `fib_companion_char_poly`, `golden_ratio_minimal_poly`, `fib_eigenvalue_product`, `fib_eigenvalue_sum`, `fib_contraction_rate`

Theorem 32.2 (Loop topology from mirror pairs). *Any mirror-dual pair (b_2, q_2) with $b_2 \neq q_2$ induces a 2-element orbit $\{(b_2, q_2), (q_2, b_2)\}$ under the mirror involution σ , and both elements map to the same $b_1 = b_2 + q_2 + 7$. The orbit has cardinality 2. At $n = 10$: the fiber $\{(24, 42), (42, 24)\}$ has cardinality 2 and maps to $b_1 = 73$.*

Lean: `mirror_fiber_two`, `mirror_pair_induces_loop`, `lepton_mirror_fiber`

Theorem 32.3 (Minimality-duality at $n = 10$). *The two prime-locked c_1 values from the mirror pair $(24, 42)$ are 823 (MDL-minimal, $823 < 2137$) and 2137 (mirror-dual). Their difference is $1314 = 18 \cdot 73 = 2 \cdot 3^2 \cdot 73$. These are the unique prime-locked pairs at $n = 10$.*

Lean: `minimality_duality_n10`, `n10_unique_mirror_pair`, `only_survivors_n10`

Theorem 32.4 (Fingerprint fixed-point). *Any monotone function $F : 2^{\mathbb{N}} \rightarrow 2^{\mathbb{N}}$ on the powerset lattice has a fixed point P^* with $F(P^*) = P^*$, by Tarski’s fixed-point theorem. Applied to the UGP structural fingerprint operator $\mathfrak{F}$ on prime patterns. A bounded *Finset* variant is additionally provided for the case of prime patterns restricted to a finite range, proved by iteration inside the finite Boolean sublattice (pigeonhole).*

Lean: `fingerprint_fixed_point_exists` (Set form, via `OrderHom.lfp`),
`fingerprint_fixed_point_bounded` (*Finset* form with boundedness hypothesis). Both fully proved; no sorry.

Theorem 32.5 (Decidability phase transition). *The GTE trajectory exhibits a sharp decidability boundary:*

1. **Local (decidable):** *For any predicate P on triples and any finite step count k , whether P holds at step k from seed G is classically decidable (follows from computability of T).*
2. **Global (Σ_1^0 -complete):** *General reachability questions are undecidable, following from the UGP Turing universality theorem (theorem 21.2).*

Lean: `local_decidability`, `decidability_phase_transition`

Additionally, `GTE.StructuralTheorems` proves:

- `fib13_is_233`: $F_{13} = 233$ (the UGP Fibonacci lift index).
- `ugp_fibonacci_rigidity`: the quotient gap $g = 13$ forces the lift $F_{13} = 233$.
- `leptonSeed_is_lex_min_residual`: the LeptonSeed $(1, 73, 823)$ is lex-min among all six residual triples at $n = 10$.

32.1 Uniqueness certificates

The module `GTE.UniquenessCertificates` collects further exhaustive classification results, all machine-checked.

The $n=10$ Uniqueness Package

All the following are machine-certified in `n10_uniqueness_package` (one theorem): (1) the only prime-locked survivors at $n = 10$ are $(24, 42)$ and $(42, 24)$; (2) the MDL seed is the LeptonSeed $(1, 73, 823)$; (3) the orbit c -values are strictly increasing $(823 < 1023 < 65535)$; (4) $2b_1 - a_2 = 137$ and 137 is prime; (5) the mirror fiber has exactly 2 elements; (6) $n = 10$ is order-2, $n = 11$ is order-0, $n = 12$ is order-1.

Theorem 32.6 (The 137 derivation). *The identity $2b_1 - a_2 = 137$ at the canonical orbit is derived, not postulated. The proof chain:*

1. *Universal shared residue: $c_1 \equiv \text{ugp1_t} = 20 \pmod{b_1}$ for any mirror-dual pair at any level (proved in section 30). Hence $m_1 = \text{gteRemainder}(823, 73) = 20$.*
2. *Odd-step formula: $a_2 = m_1 - (n + 2 - 1) = 20 - 11 = 9$.*
3. *Arithmetic: $2 \cdot 73 - 9 = 137$. `Nat.Prime 137`: [native_decide](#).*

The formula $2b_1 - a_2 = 2b_1 - (19 - n)$ depends on n ; computing it at the MDL-minimal pairs at $n = 13$ and $n = 16$ yields 412 and 3215 (not prime), so 137 is the unique prime from this formula at the smallest mirror-dual level.

Lean: `derivation_of_137`, `derivation_of_137_structural`, `alpha_echo_per_level` (in `GTE.UniquenessCertificates`)

Theorem 32.7 (Orbit non-periodicity). *The canonical GTE orbit is not eventually periodic: the c -values are strictly increasing ($c_1 < c_2 < c_3$), so no state repeats. The general principle: any trajectory with strictly monotone c -coordinates satisfies $G_t \neq G_{t+T}$ for all t, T with $T > 0$.*

Lean: `orbit_c_strictly_monotone`, `orbit_not_eventually_periodic_from_monotonicity`

Theorem 32.8 (MDL canonical seeds at $n = 10, 13, 16$). *The lex-minimal prime-locked triple $(1, b_1, c_{1,\min})$ at each certified level: $n = 10$: $(1, 73, 823)$; $n = 13$: $(1, 209, 9007)$; $n = 16$: $(1, 1609, 46681)$. These satisfy $823 < 9007 < 46681$ (MDL monotone across levels).*

Lean: `mdl_c1_n10_is_823, mdl_c1_n13_is_9007, mdl_c1_n16_is_46681, canonical_seeds_certified`

Theorem 32.9 (Order classification). *Ridge levels are classified as order-0 (no prime-locked seed), order-1 (prime-locked but no mirror pair), or order-2 (mirror pair). Machine-verified: $n = 10$ (order-2), $n = 11$ (order-0), $n = 12$ (order-1), $n = 13$ (order-2).*

Lean: `n10_is_order_2, n11_is_order_0, n12_is_order_1, n13_is_order_2`

Theorem 32.10 (Singular series lower bound). *The partial singular series for mirror-dual pairs, computed from certified local factors at good primes $p \in \{13, 23, 29, 31, 37, 41, 43\}$:*

$$\prod_{p \in \{13, 23, \dots, 43\}} \left(1 - \frac{\rho_F(p)}{(p-1)^2}\right) \geq \frac{85}{100} > 0.$$

This is a rigorous lower bound on the mirror-dual singular series $\mathfrak{S}$, establishing $\mathfrak{S} > 0$ from certified local data.

Lean: `singular_series_partial_lower_bound, singular_series_partial_positive`

33 Open Conjectures

The UGP paper posed ten conjectures; all seven that admit purely algebraic or computational proofs are established as theorems in the sections above (see theorems 6.2, 17.2 to 17.4, 19.1 and 21.2). Four conjectures remain genuinely open. All are formally stated as LEAN 4 propositions in `Conjectures.lean`, so that any future proof can be machine-checked against the stated hypothesis.

Conjecture 33.1 (Mirror-Dual). *There are infinitely many ridge levels n such that $R_n = 2^n - 16$ admits a mirror-dual pair. This is structurally analogous to the twin prime conjecture: both ask for infinitely many n where two values derived from a single arithmetic seed are simultaneously prime.*

Lean: `MirrorDualConjecture`

Conjecture 33.2 (UGP Prime Infinitude). *There are infinitely many UGP primes. This would follow immediately from the Mirror-Dual Conjecture, since each mirror-dual pair witnesses two UGP primes.*

Lean: `UGPPrimeInfinitudeConjecture`

Conjecture 33.3 (μ -Flip Distance). *Every coprime arithmetic progression $\{ak + r : k \in \mathbb{N}\}$ with $\gcd(a, r) = 1$ contains infinitely many squarefree integers. Formally: for all $a > 0$, $\gcd(a, r) = 1$, and $N \in \mathbb{N}$, there exists $n \geq N$ with $n \equiv r \pmod{a}$ and n squarefree.*

Lean: `MuFlipDistanceConjecture`

Status note. The underlying fact is known unconditionally (Estermann, 1931), but the formal LEAN 4 proof is not yet mechanized. This is an open formalization target rather than an open mathematical problem.

Conjecture 33.4 (μ -Flip Bounded Gap). *There exists a universal constant $C > 0$ such that for every $a > 0$, $\gcd(a, r) = 1$, and $N \in \mathbb{N}$, the interval $[N, N + Ca]$ contains a squarefree integer $n \equiv r \pmod{a}$.*

Lean: `MuFlipBoundedGapConjecture`

Conjecture 33.5 (Asymptotic sparsity of prime-locked levels). *Let $N(X) = |\{10 \leq n \leq X : R_n \text{ has a prime-locked seed}\}|$. Then $N(X)/X \rightarrow 0$ as $X \rightarrow \infty$.*

Heuristic: $c_1 \sim 2^n$, so $\Pr[c_1 \text{ prime}] \sim 1/(n \ln 2)$, giving $\mathbb{E}[N(X)] \lesssim C \ln X$ while X grows linearly. Formally stated as `AsymptoticSparsityConjecture`; the supporting harmonic bound $\sum_{n=10}^{59} 1/n < 3$ is certified.

Lean: `AsymptoticSparsityConjecture`, `harmonic_bound_certified` (in `GTE.UniquenessCertificates`)

Conjecture 33.6 (UGP orbital zeta factorization). *The Dirichlet series $Z_{\text{UGP}}(s) = \sum_k p_k^{-s}$ over the UGP prime sequence admits an Euler product factorization $Z_{\text{UGP}}(s) = L(s, \chi_-)^2 \cdot L(s, \chi_+)^2 \cdot G(s)$ in the zero-free region, where $G(s)$ is analytic and nonzero.*

Evidence: The first 10 UGP primes are certified; the partial sum $\sum_{k=1}^{10} p_k^{-2}$ is a certified rational lower bound. The local factor $E_{73} = 1 - 2/(73 \cdot 72) > 0$ is certified from the ρ_F identity. Full proof requires Hecke L -function theory.

Lean: `UGPZetaFactorizationConjecture`, `ugp_zeta_partial_sum_s2`, `ugp_local_factor_p73_positive` (in `GTE.UniquenessCertificates`)

Table 7: Open conjectures, formally stated. First four in `Conjectures.lean`; last two in `GTE.UniquenessCertificates`.

Conjecture	Lean name	Note
Mirror-Dual	<code>MirrorDualConjecture</code>	Analogous to twin primes
UGP prime infinitude	<code>UGPPrimeInfinitudeConjecture</code>	Follows from Mirror-Dual
μ -Flip Distance	<code>MuFlipDistanceConjecture</code>	Open formalization
μ -Flip Bounded Gap	<code>MuFlipBoundedGapConjecture</code>	Open mathematics
Asymptotic sparsity	<code>AsymptoticSparsityConjecture</code>	Requires PNT
Orbital zeta factorization	<code>UGPZetaFactorizationConj</code>	Requires Hecke L -functions

34 Related Work

Formalization of number theory. Mathlib [7] provides extensive number-theoretic infrastructure: divisors, primality, Fibonacci numbers, and the Mersenne gcd identity used in our proofs. The Liquid Tensor Experiment [2] and the formalization of perfectoid spaces [1] demonstrate LEAN 4’s capacity for large-scale mathematical formalization. The Odd Order Theorem [5] and the Kepler Conjecture [6] are landmark formalizations in Coq and HOL Light, respectively.

Prime-producing formulas. The UGP prime formula (7) is a two-parameter family constrained by the Diophantine condition (6). This places it in the tradition of prime-producing polynomials (Euler’s $n^2 + n + 41$, the Bunyakovsky conjecture) and prime-generating sieves, though the constraint to ridges $2^n - 16$ is novel. Exhaustive computation finds 57 terms up to 10^8 .

Companion work. The UGP framework was introduced in [9]; the present formalization provides the machine-checked foundation. The **MassRelations** layer is the machine-checked substrate for the Standard Model derivation in [15], which derives the gauge coupling constants, fermion mass spectrum, and gauge-sector observables from UGP-certified Lean rationals. The cyclotomic-12 Koide closed form formalized in `KoideClosedForm` and `KoideNewtonFlow` is presented as a self-contained mathematical result in [13]; that paper provides the full derivation

and physical interpretation while this artifact supplies the machine-checked Lean foundation. The self-reference theorems are imported from two companion LEAN 4 libraries:

- **nems-lean** [10]: Lawvere fixed-point theorem in abstract form, Kleene recursion theorem, and Rogers’ fixed-point theorem, all proved from a minimal self-referential interface (SRI).
- **aps-rice-lean** [8]: Rice’s theorem and halting undecidability for acceptable programming systems, proved from the s-m-n theorem and a diagonal argument.

35 Conclusion

We have presented `ugp-lean`, a self-contained LEAN 4 formalization of the Universal Generative Principle. The artifact defines the UGP framework from first principles—ridge sieve, prime-lock criterion, RSUC, GTE orbit, Quarter-Lock, Turing universality—with a strict zero-sorry, zero-custom-axiom policy on the core proof path.

The formalization contributes several results not present in prior work: the orbit-from- T derivation (theorem 6.1), the general ridge remainder lock (theorem 17.1), the even-step c -invariance theorem (theorem 6.2), the UGP prime class with ten machine-checked instances (theorem 18.3), the divisor count unboundedness theorem (theorem 29.3), the exact divisor-count formula $\tau(2^n - 16) = 5 \cdot \tau(2^{n-4} - 1)$ (theorem 29.4), the precise formal statement of the mirror-dual conjecture (theorem 29.5) with five machine-verified witness pairs, the mirror shift algebra (section 30): the universal law $c_1(b_2, q_2) - c_1(q_2, b_2) = b_1(q_2 - b_2)$ and the shared residue $c_1 \equiv 20 \pmod{b_1}$ for all mirror-dual pairs at all ridge levels, and a suite of uniqueness certificates (section 32.1): the derivation of $2b_1 - a_2 = 137$ from orbit data, non-periodicity of the canonical orbit, the symmetry group $\mathbb{Z}/2\mathbb{Z}$ of the mirror fiber, MDL canonical seeds at $n = 10, 13, 16$, order classification, and a certified singular series lower bound $\geq 85/100$.

The library has grown substantially from its initial core. Current highlights beyond the foundational spine:

Elegant Kernel unconditional closure (theorems 19.4 to 19.6): $k_{\text{gen}} = \varphi \cos(\pi/10)$, $k_{\text{gen}}^{(2)} = -\varphi/2$, and the full pentagon- D_5 kernel structure derived from first principles.

Mass Relations layer (section 20, 22 modules): the Koide cyclotomic-12 closed form with S_3 -equivariant Newton flow, the N_c -structural chain deriving all fermion constants from $N_c = 3$, VV coefficients from $SU(5)/SO(10)$ group theory, the neutrino seesaw closure with three independent structural decompositions, the Mersenne mass structure ($\delta b_1 = 2^{N_c^2} - 1$), mass ratio Z -independence, the seesaw index as $SO(10)$ gauge/matter representation defect, the CKM θ_{23} structural ratio $\tau(R_{10})/D_1 = 15/8$ uniquely forced at $n = 10$ (section 20.7), and the discrete arithmetic-mean identity $2a_\tau = a_e + a_\mu$ on the canonical lepton orbit (section 20.8). CDM Cabibbo derivation: $|V_{us}|_{\text{CDM}} = e^{-13\pi/27}$ from GUT rank correction (CKMMixing, section 20.9). Neutrino mass-squared ratio: $R = 0.02936$ certified tight (NeutrinoMassRatio, section 20.10).

SRRG no-go and VEV proof (section 20.11, 4 modules): SRRG β -function excludes perturbative dimensional transmutation; PSC entropy-contraction duality with rate $\lambda = 1/\varphi$; per-generation Goldstone entropy volume correction $f_{\text{vol}} = \varphi^{-1/N_{\text{gen}}}$; EW vacuum manifold S^3 with $V = 2\pi^2$ and uniqueness certified.

GUT structure and physical formalization (section 22, 10,159 lines, 82 §-groups, 515 items): the full arithmetic bridge from $(N_{\text{gen}}, N_{\text{fam}})$ to SM observables: $\sin^2 \theta_W = 3/13$, $\lambda_{\text{CKM}} = 9/40$, the six-quark N_{eff} capstone, $D = 4$ from dimensional slices, Z_7 charge formula $Q = w^*/3$, $\beta_0 = 23/3$ QCD one-loop coefficient, hypercharge consistency, Gorard matter step $\kappa_{\text{SD}} > 0$, mass ordering from tail lengths, baryogenesis amplitude counting, and the joint GTE/Rule 110/SM unification capstone (`ugp_r110_sm_joint_unification`).

Extended Universality modules (section 23, 12 modules): GoE orbital chain isolation (GoEStabilityHierarchy); perturbation catalog proving Rule 110 isola-

tion with zero tolerance (`OrbitPerturbationCatalog`); Z_7 charge conjugation, CP asymmetry, and vertex duality (`Z7ChargeConjugation`); $N_{\text{fam}} = 5$ from prime transitivity uniqueness (`Z5TransitivityUniqueness`); $D = 3$ from dimensional slice propagation (`DimensionalSliceUniqueness`); GTP neutral particle discrimination (`GTPNeutralDiscrimination`); 6-part SM orbit causal isolation master theorem (`SMOrbitCausalIsolation`); EW boson c-staircase $c \in \{11, 12, 13\}$ (`EWBosonStructure`); chirality bridge axiom stub for unconditional Weinberg (`EWChiralBridge`); photon as CA ether and masslessness criterion (`CasimirMasslessEther`); Lawvere zone classification L0/L1/L2 for P34 Conjecture C4 (`LawvereZone`).

GTE framework and C1 final coalgebra (sections 21.8 to 21.10, 3 modules): GTE as a `NemS.Framework` with transputation classification (`GTEFrameworkInstance`); GTE D-uniqueness and optimality (`GTEOptimalityInstance`); GTE as the terminal F_{PSC} coalgebra in `PSCSys`, zero sorry, zero custom axioms (`GTEFinalCoalgebra`).

GTE extended structure (section 7, 6 new modules): the Mersenne ladder with Fibonacci uniqueness certification, the fiber bundle with canonical orbit trivialization, orbital stability via the Fibonacci contraction mode, the universal 5-factor ridge-level scale connection, the GTB generation prime certificate (all six c_1 values certified prime, ridge spacing equal to N_c), and $N_c = 3$ derived from GTE substrate arithmetic alone via two independent routes (Mersenne GCD and ridge factorial uniqueness; `NcColorArithmetic`).

Braid Atlas charge structure (section 8): $Q = W_g/N_c$ for all four SM fermion types, and the anomaly cancellation theorem that per-generation winding sum $= N_c(N_c - 3) = 0 \Leftrightarrow N_c = 3$.

Null discipline formalization (section 14): machine-classification of evidential grades via the four-gate Theorem Eligibility Criterion; the URC basis saturation theorem; corpus instances classified as theorem-grade or Class C.

Information Profit Threshold (section 15): the threshold $\rho_{\text{crit}} = 1 + \ln \varphi / (2 \ln 2\pi)$ derived from three structural premises within the formal PSC/Fibonacci/ $U(1)$ -entropy model.

Gauge coupling extensions (section 16): exact rational tree-level predictions $\sin^2 \theta_W = 3456/15101$ and $1/\alpha_{\text{EM}} = 4\pi \cdot 377525/37264 \approx 127.31$ from the certified bare couplings.

Physical extension (development branch) (sections 24 to 28): The development branch (`ugp-lean-exp`) extends the library into physical spacetime and QFT. Highlights: spectral dimension $d_s = 4$ in the thermodynamic limit (`causal_graph_spectral_dim_thermodynamic_limit`, one documented Mathlib-gap sorry); the Algebraic Lifting Theorem connecting beable-level results to physical observables (`algebraic_lifting_theorem`, zero sorry, commit `f83ae77`); meson and baryon existence in 3D (`meson_bound_state_exists`, `baryon_bound_state_exists`, zero sorry conditional on `causal_path_exists`); two independent Lean-certified derivations of $\alpha_{\text{EM}} \approx 1/137$ (Route A cascade identification and Route B Berry holonomy; module `SylowIndexCouplingHierarchy.lean`); strong CP resolution $\theta_{\text{QCD}} = 0$ from F_{21} group theory (`strong_cp_resolved`, zero sorry, three independent proofs); asymptotic freedom $b_0 = 7$ and two-loop $b_1 = 26$ (`f21_substrate_asymptotic_freedom`, `f21_two_loop_beta_coefficient`); a conditional QFT mass gap lower bound (`qft_gauged_mass_gap_lower_bound`) and unconditional mass-gap theorems (`qft_gauged_mass_gap_unconditional`); the Algebraic Descent Theorem establishing that all F_{21} -algebraic properties are resolution-independent (`algebraic_descent_theorem`); and bridge capstones composing earlier results into the unified physical exclusion criterion (`unified_physical_exclusion`) and the three-generation uniqueness capstone (`gte_uniquely_predicts_three_generations`). A systematic audit reduced the global custom named-axiom count from 24 to 20 by discharging four `VEVProof` axioms as genuine theorems.

What remains open. Six conjectures are formally stated in table 7. The mirror-dual conjecture (theorem 33.1) and UGP prime infinitude (theorem 33.2) are the most mathematically significant: both are structurally analogous to the twin prime conjecture and would require new sieve-theoretic techniques. The μ -flip conjectures concern the distribution of squarefree integers in arithmetic progressions and are open as formalization targets even where the mathematics is known.

Reproducibility. The artifact is publicly available at <https://github.com/novaspivack/ugp-lean>. Building requires LEAN 4 v4.29.0-rc6 and Mathlib v4.29.0-rc6. The command `lake build` compiles the entire library; no external tools beyond LEAN 4 and its package manager are needed.

Table 8: Complete theorem inventory. “Novel” indicates results new to this formalization; “General” indicates level-independent results. All theorems have zero `sorry` on the core proof path.

Theorem / Lean name	Module	Paper ref	Novel?
<i>Core Classification (§5)</i>			
<code>theoremA_general</code>	Classification.TheoremA	Thm A	—
<code>theoremB</code>	Classification.TheoremB	Thm B	—
<code>rsuc_theorem</code>	Classification.RSUC	RSUC	—
<code>mdl_selects_LeptonSeed</code>	Classification.TheoremB	MDL	—
<code>rsuc_formal</code>	Classification.FormalRSUC	RSUC	—
<code>rsuc_canon</code>	Classification.FormalRSUC	RSUC	—
<code>strengthening_cannot_add</code>	Classification.Monotonic	—	—
<i>Ridge & Sieve (§5)</i>			
<code>ridgeSurvivors_10</code>	Compute.Sieve	§5	—
<code>ridge_remainder_lock</code>	Core.RidgeRigidity	Lem m_2	—
<code>exclude_16 ... exclude_63</code>	Compute.ExclusionFilters	§5	—
<i>GTE Orbit (§6)</i>			
<code>orbit_determined_by_T</code>	GTE.UpdateMap	Prop 5.1	Yes
<code>update_map_produces_orbit</code>	GTE.UpdateMap	Prop 5.1	Yes
<code>even_step_c_invariance</code>	GTE.UpdateMap	—	Yes
<code>c3_strict_eq_c2_at_n10</code>	GTE.UpdateMap	—	Yes
<code>gen1_mersenne_isolation</code>	GTE.UpdateMap	—	Yes
<code>mersenne_entanglement_gen</code>	GTE.UpdateMap	—	Yes
<i>Dynamics Companions / ML-9 (§6)</i>			
<code>gteTransition</code>	GTE.GTESimulation	—	Yes
<code>gteAfterSteps</code>	GTE.GTESimulation	—	Yes
<code>gte_s1...gte_s8</code>	GTE.GTESimulation	—	Yes
<code>gte_entropy_prefix8_gt_prefix9</code>	GTE.EntropyNonMonotone	ML-9	Yes
<code>Hpred8_gt_Hpred9</code>	GTE.EntropyNonMonotone	ML-9	Yes
<i>General-Level Theorems (§17)</i>			
<code>ridge_remainder_lock_gen</code>	GTE.UpdateMap	Lem m_2	General
<code>quotient_gap_all</code>	GTE.GeneralTheorems	gap-13	General
<code>mirror_b1_invariance</code>	GTE.UpdateMap	—	Yes
<code>mersenne_gcd_identity</code>	GTE.UpdateMap	—	Mathlib
<code>alpha_echo</code>	GTE.GeneralTheorems	α -echo	Checked
<code>b_linear_growth</code>	GTE.GeneralTheorems	b -growth	—
<code>c_values_strictly_mono</code>	GTE.GeneralTheorems	c -mono	—
<i>Quarter-Lock & Kernel (§19)</i>			
<code>quarterLockLaw</code>	QuarterLock	§3	—

continued on next page

(Table 8 continued)

Theorem / Lean name	Module	Paper ref	Novel?
quarterLockStability	QuarterLock	§3	—
k_L2_eq	ElegantKernel	§3	—
L_model_from_gauge_struct	LModelDerivation	§3	—
<i>UCL Unconditional Closure (§19.1)</i>			
thm_ucl1_fully_unconditional	EK/U/FullClosure	Thm 19.4	Yes
thm_ucl1_cosine_form	EK/U/FullClosure	Thm 19.4	Yes
k_gen_eq_pi_div_two	ElegantKernel/KGen	Thm 19.4	Yes
psi_minimal_poly	ElegantKernel/KGen	§19.1	Yes
elegant_kernel_Lg_block	ElegantKernel/KGen	§19.1	Yes
k_gen2_derived_satisfies_poly	ElegantKernel/KGen2	Thm 19.5	Yes
k_gen2_derived_unique_char...	EK/U/KGenFull	Thm 19.5	Yes
cos_4pi_div_five_eq_neg_phi_half	EK/U/KGenFull	Thm 19.5	Yes
full_closure_summary	EK/U/FullClosure	Thm 19.6	Yes
complete_derivation	EK/U/FullClosure	Thm 19.6	Yes
<i>Mass Relations & Physical Spectrum (§20)</i>			
koide_iff_twoS_sq_eq_threeN	MR.KoideClosedForm	Thm 20.1	Yes
koide_solved_form_root	MR.KoideClosedForm	Thm 20.2	Yes
cos_sq_pi_div_twelve,	MR.KoideClosedForm	Thm 20.2	Yes
two_plus_sqrt3_eq			
newton_flow_fixes_null_cone	MR.KoideNewtonFlow	Thm 20.3	Yes
newton_flow_swap12_equivariant	MR.KoideNewtonFlow	Thm 20.3	Yes
newton_flow_rot123_equivariant	MR.KoideNewtonFlow	Thm 20.3	Yes
cascadeState_closed_form	MR.BinaryCascade	Thm 20.4	Yes
cascade_increment_doubles,	MR.BinaryCascade	Thm 20.4	Yes
cascade_delta_1_to_3			
koidePredictedMTau_pos	MR.PhysicalMasses	Thm 20.5	Yes
predictedLepton_pos,	MR.PhysicalMasses	Thm 20.5	Yes
predictedUpType_pos			
a2_weyl_chamber_half_opening	MR.SU3FlavorCartan	§20	Yes
omega1_in_weyl_chamber_interior	MR.SU3FlavorCartan	§20	Yes
cartanFlavonPotential_ge_min	MR.CartanFlavonPot.	§20	Yes
fn_vevs_are_potential_minima	MR.CartanFlavonPot.	§20	Yes
<i>N_c=3 and VEV proof modules (§7.7–20.11)</i>			
nc_eq_3_from_mersenne_gcd	GTE.NcColorArithmetic	P24	Yes
nc_uniqueness_from_ridge_divisors	GTE.NcColorArithmetic	P24	Yes
nc_eq_3_from_substrate	GTE.NcColorArithmetic	P24	Yes
cabibbo_effective_charge,	MR.CKMMixing	P01	Yes
cabibbo_vev_formula			
fn_vv_correction_additive	MR.CKMMixing	P01	Yes
neutrino_mass_ratio_tight_bound	MR.NeutrinoMassRatio	P21	Yes
seesaw_ratio_independent_of_MR	MR.NeutrinoMassRatio	P21	Yes
dt_integral_diverges	VEVNoGo.SRRGNoGo	SRRG	Yes
psc_entropy_contraction_duality_proved	VEVProof.PSCEntropyDuality	P27	Yes
per_gen_volume_correction	VEVProof.GoldstoneEntropy	P27	Yes
ew_goldstone_count, vol_S3_eq	VEVProof.EWGoldstoneManifold	P27	Yes
ew_vacuum_manifold_uniqueness	VEVProof.EWGoldstoneManifold	P27	Yes
<i>Universality & Self-Reference (§21)</i>			
uwca_P3_eq_rule110	Universality.UWCASimul	Thm 21.1	Yes
uwca_sweep_implements_rule110	Universality.UWCASimul	Thm 21.1	Yes
uwca_simulates_rule110_real	Universality.UWCAEmbeds	Thm 21.1	Yes
ugp_is_turing_universal	Universality.TuringUniv	Thm 21.2	—
uwca_augmented_left_inverse	Univ.UWCAHistoryReversible	T2	Yes
ugp_lawvere_fixed_point	SelfRef.LawvereKleene	Lawvere	Import

continued on next page

(Table 8 continued)

Theorem / Lean name	Module	Paper ref	Novel?
ugp_rice_theorem	SelfRef.RiceHalting	Rice	Import
ugp_halting_undecidable	SelfRef.RiceHalting	Halting	Import
<i>CUP Theorems (§21.3–21.6)</i>			
cup4_parity_uniqueness	Univ.CUP4TotalParity	CUP-4	Yes
cup1_orbit_uniquely_selects_rule110	Univ.CUP4TotalParity	CUP-4	Yes
cup11c_universal_mod7_CA_exists	Univ.CUP11ModSeven	CUP-11c	Yes
fmdl_never_outputs_4	Univ.CUP3DUniqueness	§22	Yes
fmdl_gen1_is_garden_of_eden	Univ.CUP3DUniqueness	GoE	Yes
fmdl_unique_uniform_fixed_point	Univ.CUP3DUniqueness	Photon	Yes
cup3d_psc_goe_chain	Univ.CUP3DPSCUnification	PSC	Yes
physical_incompleteness_thm	Univ.CUP3DPPhysicalInc.	CUP-PI	3 axioms
hypothesis_b_tape_level	Univ.HypothesisB	Hyp B	1 axiom
hypothesis_c_psc_forces_universality	Univ.PSCUniversality	Hyp C	1 axiom
gte_compilation_theorem	Univ.GTECompilation	§21.12	Yes
gte_unique_bisimulation	Univ.GTEUniqueness	§21.13	Yes
<i>GoE Stability & Orbit Perturbation (§23.1–23.2)</i>			
gen1_is_goe	Univ.GoESTabilityHierarchy	GoE	Yes
gen2_unique_predecessor	Univ.GoESTabilityHierarchy	—	Yes
sm_chain_fully_isolated	Univ.GoESTabilityHierarchy	—	Yes
orbit_perturbation_destroys_universality	Univ.OrbitPerturbCat.	P28	Yes
<i>Z₇ Charge Conjugation (§23.3)</i>			
z7_conj_involution,	Univ.Z7ChargeConj.	§CC	Yes
z7_conj_sum_zero			
z7_conj_pairs	Univ.Z7ChargeConj.	§CC	Yes
fmdl_matter_cp_violation	Univ.Z7ChargeConj.	P33	Yes
fmdl_conj_pair_asymmetry_unique	Univ.Z7ChargeConj.	P33	Yes
ca_w_plus_is_emission_not_absorption	Univ.Z7ChargeConj.	P33	Yes
cup11b_z7_sum_conservation	Univ.CUP3DUniqueness	CUP-11b	Yes
<i>Dimensional & Transitivity Uniqueness (§23.4–23.7)</i>			
z5_prime_unique_transitivity	Univ.Z5TransitivityUniq.	P28	Yes
three_dim_fmdl_structure_forced	Univ.DimSliceUniq.	P36	Yes
gtp_neutral_particles_distinct	Univ.GTPNeutralDisc.	P28	Yes
sm_orbit_complete_causal_isolation	Univ.SMOrbitCausalIsol.	P33	Yes
<i>EW Boson Structure & Casimir Ether (§23.8–23.10)</i>			
ew_c_staircase,	Univ.EWBosonStructure	P33	Yes
ew_c_arithmetic_progression			
ew_higgs_is_scalar_boundary	Univ.EWBosonStructure	P31	Yes
doublet_partner_is_left_chiral	Univ.EWChiralBridge	P31	axiom stub
fmdl_unique_uniform_fixed_point	Univ.CasimirMassless.	P33	Yes
fmdl_massless_criterion	Univ.CasimirMassless.	P33	Yes
ether_z7_sum_mod7	Univ.CasimirMassless.	P33	Yes
<i>Lawvere Zone (§23.11)</i>			
zone_L0_is_vacuum	Univ.LawvereZone	P34	Yes
zone_L1_is_goe_orbit	Univ.LawvereZone	P34	Yes
zone_L2_lawvere_diagonal	Univ.LawvereZone	P34	Yes
<i>GUT Structure — EW observables (§22)</i>			

continued on next page

(Table 8 continued)

Theorem / Lean name	Module	Paper ref	Novel?
gut_weinberg_structure	Univ.GUTStructure §8	P31	Yes
ngen_plus_nfam_eq_pow2	Univ.GUTStructure §2	P31	Yes
weinberg_angle_closure	Univ.GUTStructure §12	P31	Yes
wolfenstein_lambda_formula	Univ.GUTStructure §14	P32	Yes
six_quark_neff_complete	Univ.GUTStructure §34	P32	Yes
ugp_r110_sm_joint_unification	Univ.GUTStructure §27	P35	Yes
gte_spacetime_dimension	Univ.GUTStructure §54	P36	Yes
charge_from_z7_winding	Univ.GUTStructure §49	P37	Yes
hypercharge_u_quark,	Univ.GUTStructure §73	P37	Yes
weinberg_angle_from_hypercharge_sum			
gorard_matter_step_kappa_positive	Univ.GUTStructure §74	P36	Yes
tail_length_strict_ordering	Univ.GUTStructure §76	P37	Yes
vacuum_ollivier_ricci_flatness	Univ.GUTStructure §32	P36	Yes
qcd_beta0_from_gte	Univ.GUTStructure §81	P33	Yes
orbit_sum_winding_classes	Univ.GUTStructure §79	P33	Yes
eta_B_amplitude_structure	Univ.GUTStructure §70	P33	Yes
fmdl_perfect_code	Univ.GUTStructure §36	P33	Yes
mw_mz_ratio_squared	Univ.GUTStructure §45	P33	Yes
<i>GTE Framework & C1 Final Coalgebra (§21.8–21.10)</i>			
GTEFramework	Framework.GTEFrameworkInstanceNEMS		Yes
gte_tpc_from_nems_classification	Framework.GTEFrameworkInstanceC3		Yes
gte_tpc_real	Framework.GTEFrameworkInstanceC3		Yes
gte_d_unique	Framework.GTEOptimalityInstanceP34		Yes
psc_optimal_zero_on_free	Framework.GTEFinalCoalgebra C1		Yes
c1_final_coalgebra_derived	Framework.GTEFinalCoalgebra C1		Yes
c1_lambek_isomorphism	Framework.GTEFinalCoalgebra C1		Yes
<i>UGP Primes (§18)</i>			
IsUGPPrime	GTE.UGPPrimes	Def 18.1	Yes
ugp_primes_exist	GTE.UGPPrimes	—	Yes
ten_ugp_primes	GTE.UGPPrimes	—	Yes
ugp_primes_at_three_lvls	GTE.UGPPrimes	—	Yes
c1Val_eq_c1FromPair	GTE.UGPPrimes	—	Yes
UGPPrimeInfinitudeConj	GTE.UGPPrimes	—	Yes
<i>Mirror-Dual Conjecture (§29)</i>			
card_divisors_ridge_unb	GTE.MirrorDualConj	—	Yes
tau_ridge_exact	GTE.MirrorDualConj	Thm 29.4	Yes
coprime_pow2_mersenne	GTE.MirrorDualConj	—	Yes
five_mirror_dual_pairs	GTE.MirrorDualConj	—	Yes
MirrorDualConjecture	GTE.MirrorDualConj	—	Yes
<i>Mirror Shift Algebra (§30)</i>			
lepton_mirror_shift	Core.MirrorAlgebra	Thm 30.3	Yes
mirror_shift_factored	Core.MirrorAlgebra	Thm 30.3	Yes
lepton_shared_residue	Core.MirrorAlgebra	Thm 30.1	Yes
c1Val_mod_b1	GTE.MirrorShift	Thm 30.1	General
mirror_pair_shared_residue	GTE.MirrorShift	Thm 30.1	General
mirror_shift_formula	GTE.MirrorShift	Thm 30.2	General
mirror_shift_divides_b1	GTE.MirrorShift	Thm 30.2	General
lepton_shift_is_18_times_b1	GTE.MirrorShift	Thm 30.3	Yes
shift_n13	GTE.MirrorShift	Tab 6	Yes
shift_n16_a	GTE.MirrorShift	Tab 6	Yes
<i>Companion modules (not in paper; see RESONANT_FACTORY_NOTE.md)</i>			
qm_no_root_of_legendreSym_neg_one	GTE.InertPrimes	—	General

continued on next page

(Table 8 continued)

Theorem / Lean name	Module	Paper ref	Novel?
qm_inert_primes_certified	GTE.InertPrimes	—	Yes
kappa_unram_lower_bound	GTE.InertPrimes	—	Yes
discKminus, discKplus	GTE.ResonantFactory	—	Yes
rhoF_identity_verified	GTE.ResonantFactory	—	Yes
dickman_equidistribution_in_APs	GTE.AnalyticArch	—	Sorry (Tenenbaum)
qminus_qplus_coprime	GTE.AnalyticArch	—	Yes
<i>Resonant Factory (§29.1)</i>			
branch_linearization	GTE.ResonantFactory	Thm 29.6	Yes
factory_gap_two	GTE.ResonantFactory	Thm 29.8	Yes
factoryDp_prime	GTE.ResonantFactory	—	Yes
localDensity_3..43	GTE.ResonantFactory	Thm 29.9	Yes
hasse_check_no_obstruction	GTE.ResonantFactory	Thm 29.9	Yes
<i>Structural Theorems (§32)</i>			
golden_ratio_minimal_poly	GTE.StructuralThm	Thm 32.1	Yes
fib_eigenvalue_product	GTE.StructuralThm	Thm 32.1	Yes
fib_contraction_rate	GTE.StructuralThm	Thm 32.1	Yes
mirror_fiber_two	GTE.StructuralThm	Thm 32.2	Yes
mirror_pair_induces_loop	GTE.StructuralThm	Thm 32.2	Yes
minimality_duality_n10	GTE.StructuralThm	Thm 32.3	Yes
only_survivors_n10	Compute.ExclFilters	Thm 32.3	Yes
fingerprint_fixed_point_exists	GTE.StructuralThm	Thm 32.4	Yes
fingerprint_fixed_point_bounded	GTE.StructuralThm	Thm 32.4	Yes
decidability_phase_transition	GTE.StructuralThm	Thm 32.5	Yes
leptonSeed_is_lex_min_residual	GTE.StructuralThm	—	Yes
<i>Uniqueness Certificates (§32.1)</i>			
derivation_of_137	GTE.UniquenessCert	Thm 32.6	Yes
derivation_of_137_structural	GTE.UniquenessCert	Thm 32.6	Yes
alpha_echo_per_level	GTE.UniquenessCert	Thm 32.6	Yes
orbit_not_eventually_periodic	GTE.UniquenessCert	Thm 32.7	Yes
rule110_tiles_are_minimal	GTE.UniquenessCert	Thm 21.3	Yes
mirror_symmetry_group_is_Z2	GTE.UniquenessCert	Thm 21.4	Yes
no_order_3_loop	GTE.UniquenessCert	Thm 21.4	Yes
mdl_c1_n13_is_9007	GTE.UniquenessCert	Thm 32.8	Yes
mdl_c1_n16_is_46681	GTE.UniquenessCert	Thm 32.8	Yes
order_classification_n10_to_n13	GTE.UniquenessCert	Thm 32.9	Yes
singular_series_partial_lower_bound	GTE.UniquenessCert	Thm 32.10	Yes
n10_uniqueness_package	GTE.UniquenessCert	§32.1	Yes
AsymptoticSparsityConjecture	GTE.UniquenessCert	Conj 33.5	Stated
UGPZetaFactorizationConjecture	GTE.UniquenessCert	Conj 33.6	Stated
<i>Additional proved results (from original conjectures)</i>			
mirror_min_dual_proved	Conjectures	Thm 17.3	Yes
fib_rigidity_proved	Conjectures	Thm 17.2	Yes
c1_monotone_in_q2	Conjectures	Thm 17.4	Yes
robust_universality_proved	Conjectures	Thm 21.2	Yes
sharp_boundary_proved	Conjectures	§21	Yes
kernel_compatibility_proved	Conjectures	Thm 19.1	Yes
global_c_attractor_proved	Conjectures	Thm 6.2	Yes
<i>Real Analysis / DSI Export (§31)</i>			
ugpOutputGap_deriv	GTE.DSIExport	Thm 31.2	Yes
ugp_deriv_lower_bound	GTE.DSIExport	Thm 31.3	Yes

continued on next page

(Table 8 continued)

Theorem / Lean name	Module	Paper ref	Novel?
ugpOutputGap_differentiable0n	GTE.DSIEExport	—	Yes

References

- [1] Kevin Buzzard, Johan Commelin, and Patrick Massot. Formalising perfectoid spaces. In *Proc. 9th ACM SIGPLAN Int. Conf. Certified Programs and Proofs (CPP 2020)*, pages 299–312. ACM, 2020.
- [2] Johan Commelin, Adam Topaz, et al. Liquid tensor experiment. <https://github.com/leanprover-community/lean-liquid>, 2022. Lean 4 formalization of Scholze’s challenge.
- [3] Matthew Cook. Universality in elementary cellular automata. *Complex Systems*, 15(1):1–40, 2004.
- [4] Leonardo de Moura and Sebastian Ullrich. The Lean 4 theorem prover and programming language. In *Proc. 28th Int. Conf. Automated Deduction (CADE-28)*, volume 12699 of *LNCs*, pages 625–635. Springer, 2021.
- [5] Georges Gonthier, Andrea Asperti, Jeremy Avigad, Yves Bertot, Cyril Cohen, François Garillot, Stéphane Le Roux, Assia Mahboubi, Russell O’Connor, Sidi Ould Biha, Ioana Paşca, Laurence Rideau, Alexey Solovyev, Enrico Tassi, and Laurent Théry. A machine-checked proof of the odd order theorem. In *Proc. 4th Int. Conf. Interactive Theorem Proving (ITP 2013)*, volume 7998 of *LNCs*, pages 163–179. Springer, 2013.
- [6] Thomas Hales, Mark Adams, Gertrud Bauer, Tat Dat Dang, John Harrison, Le Truong Hoang, Cezary Kaliszyk, Victor Magron, Sean McLaughlin, Thang Tat Nguyen, et al. A formal proof of the Kepler conjecture. In *Forum of Mathematics, Pi*, volume 5, page e2. Cambridge Univ. Press, 2017.
- [7] Mathlib Community. Mathlib4: The Lean 4 mathematical library. <https://github.com/leanprover-community/mathlib4>, 2024. Version 4.29.0-rc3.
- [8] Nova Spivack. aps-rice-lean: Rice’s theorem and halting undecidability for acceptable programming systems in Lean 4. <https://github.com/novaspivack/aps-rice-lean>, 2025.
- [9] Nova Spivack. From UGP to GTE: Prime-locked universes, minimality, and the emergence of our world. Zenodo, 2025.
- [10] Nova Spivack. nems-lean: Self-reference library for Lean 4 (Paper 26). <https://github.com/novaspivack/nems-lean>, 2025.
- [11] Nova Spivack. Canonical Braid Atlas v2.0: First-principles derivation of Standard Model quantum numbers from GTE topology. Zenodo, 2026.
- [12] Nova Spivack. The information profit principle: A threshold condition for self-maintaining systems. Zenodo, 2026.
- [13] Nova Spivack. The Koide relation as a Cyclotomic-12 closed form. Zenodo, 2026.
- [14] Nova Spivack. Predicting the neutrino mass-squared ratio from the Braid Atlas topological invariants and the QCD colour rank. Zenodo, 2026.

- [15] Nova Spivack. The standard model from the Universal Generative Principle. Zenodo, 2026.
- [16] Nova Spivack. Substrate depth and self-generated mass: The c -component of the UGP canonical triple predicts reflexive closure across the Standard Model spectrum. Zenodo, 2026.
- [17] Gérald Tenenbaum. *Introduction to Analytic and Probabilistic Number Theory*, volume 163 of *Graduate Studies in Mathematics*. American Mathematical Society, 3rd edition, 2015.